

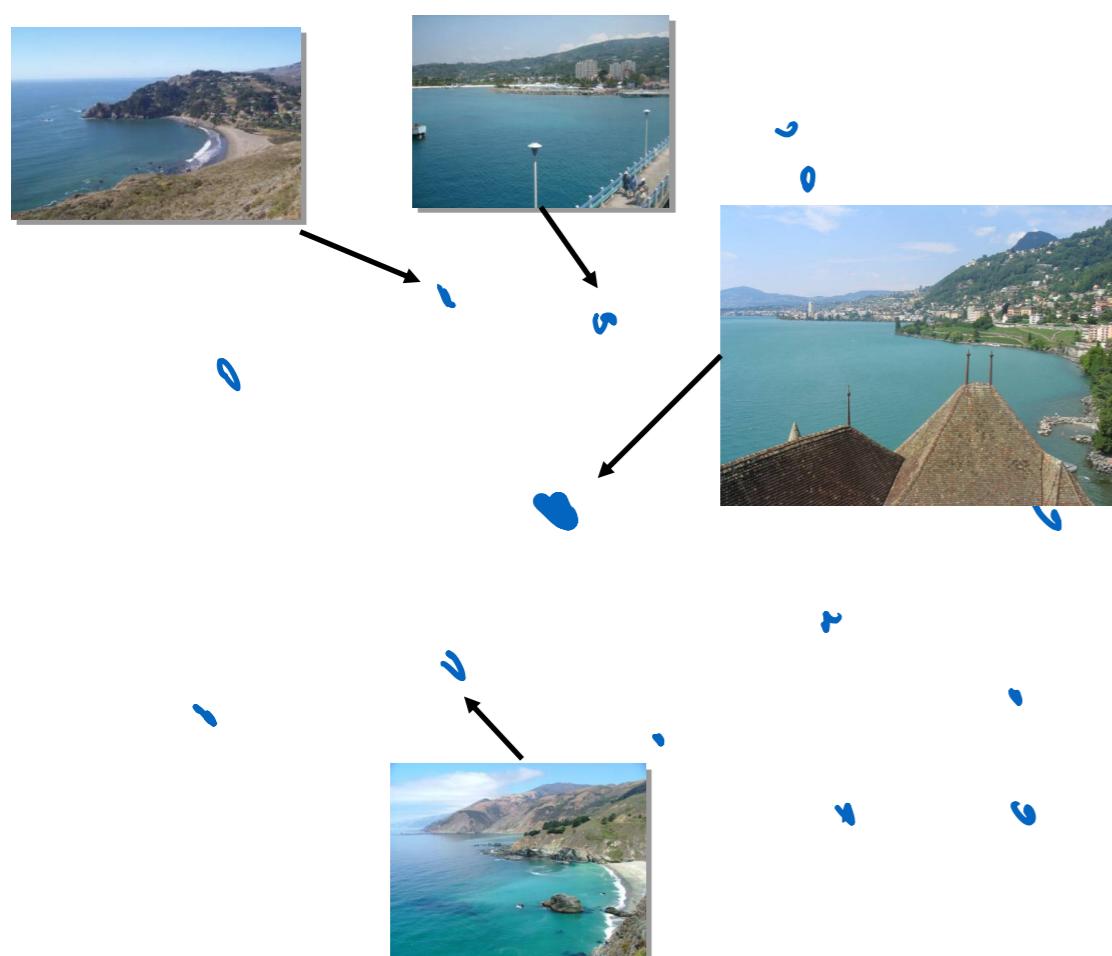
Convolutional Network for Image Synthesis

Jun-Yan Zhu

16-726 Learning-based Image Synthesis, Spring 2021

many slides from Alyosha Efros, Phillip Isola, Richard Zhang, James Hays, and Andrea Vedaldi, Jitendra Malik.

Review (data-driven graphics)



Review (data-driven graphics)



Nearest neighbor methods:

1. Stored examples
2. Calculate distance between two examples
3. Voting (label transfer): image blending/averaging

Visual similarity via labels



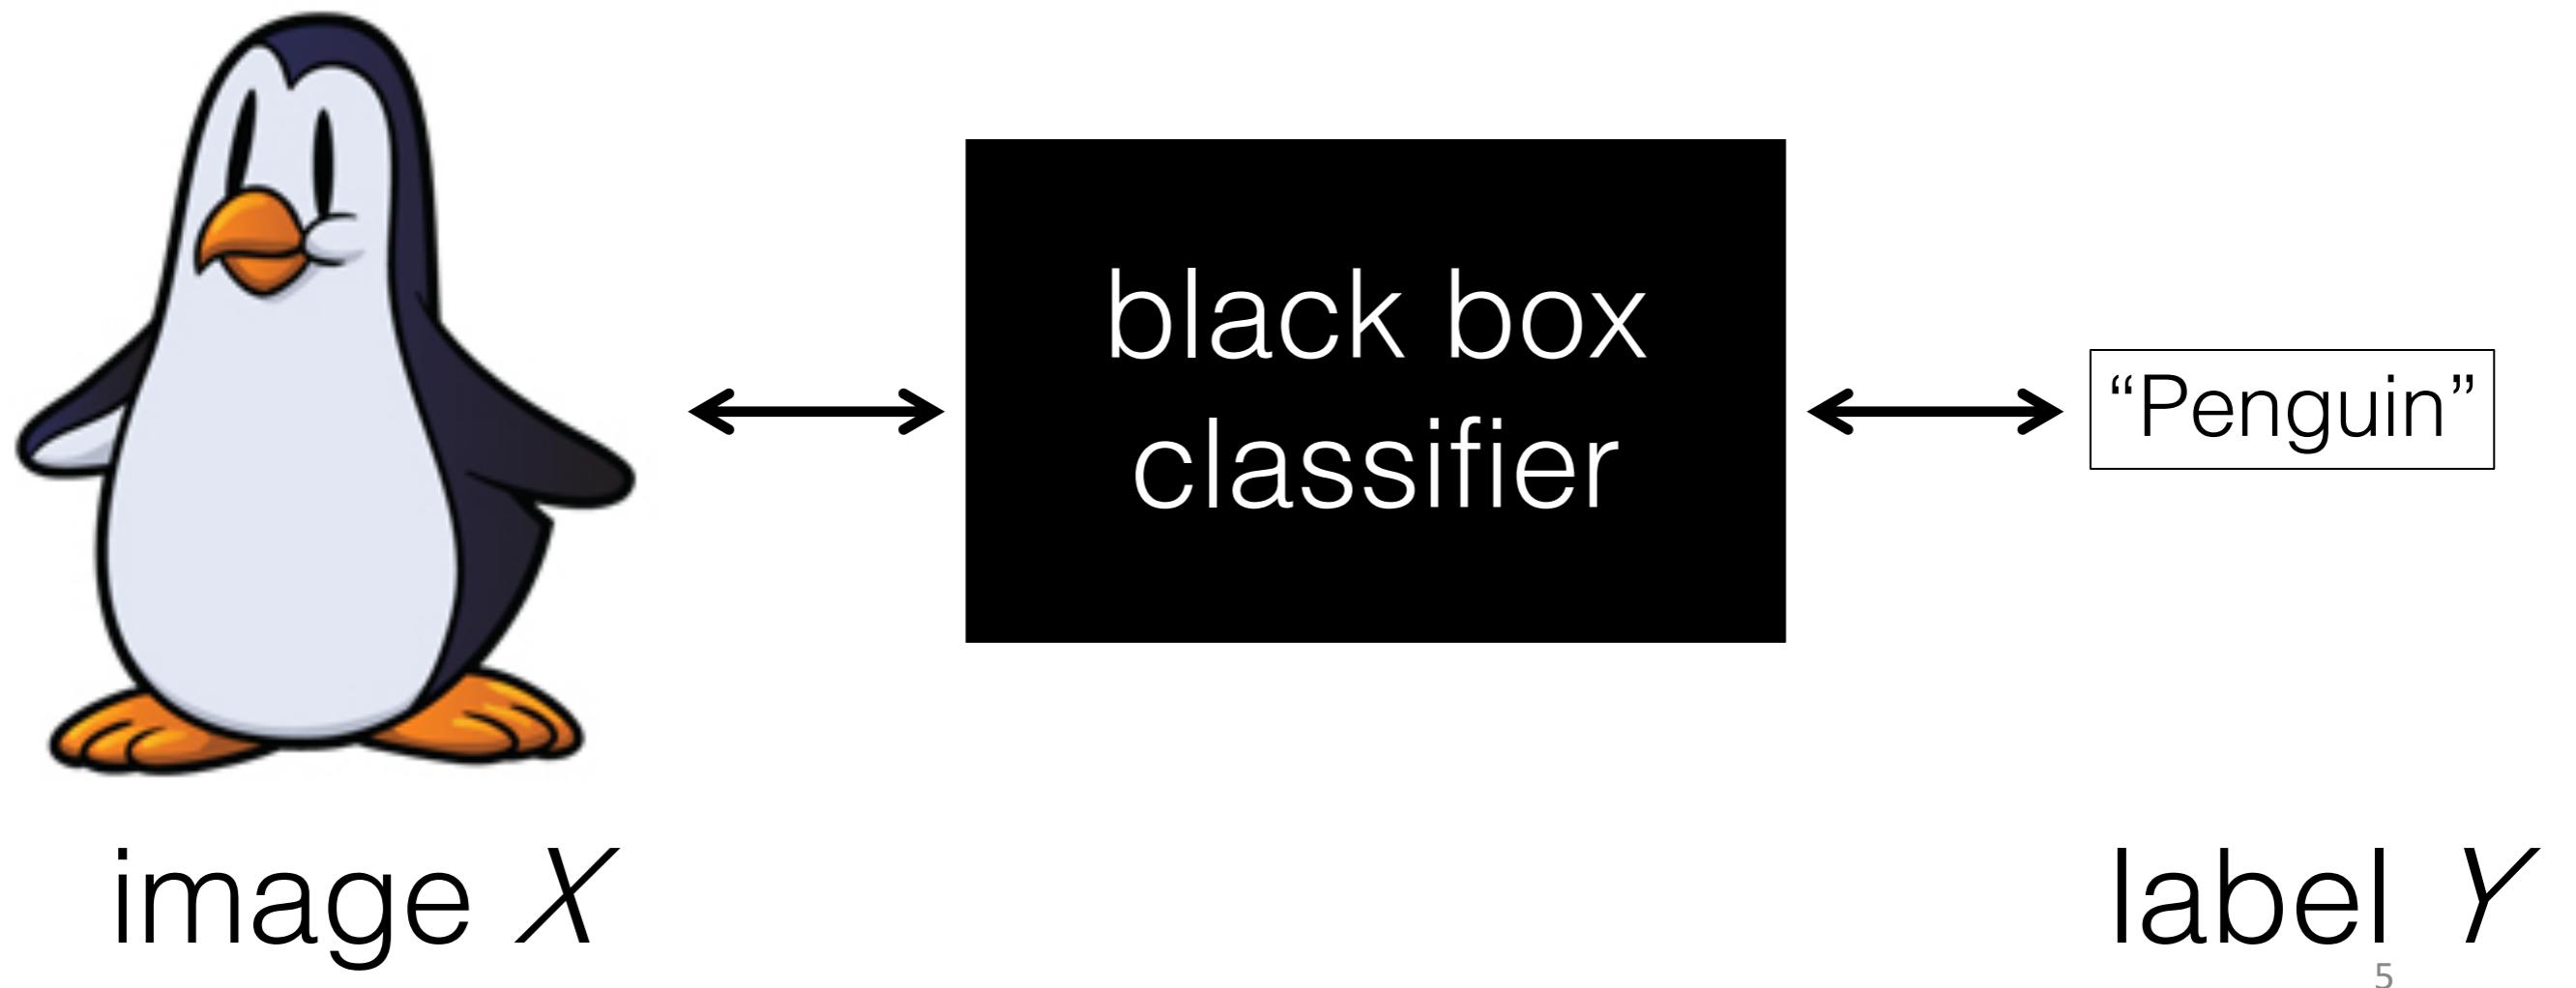
“Penguin”

?
==



“Penguin”

Machine Learning as data association



At test time...

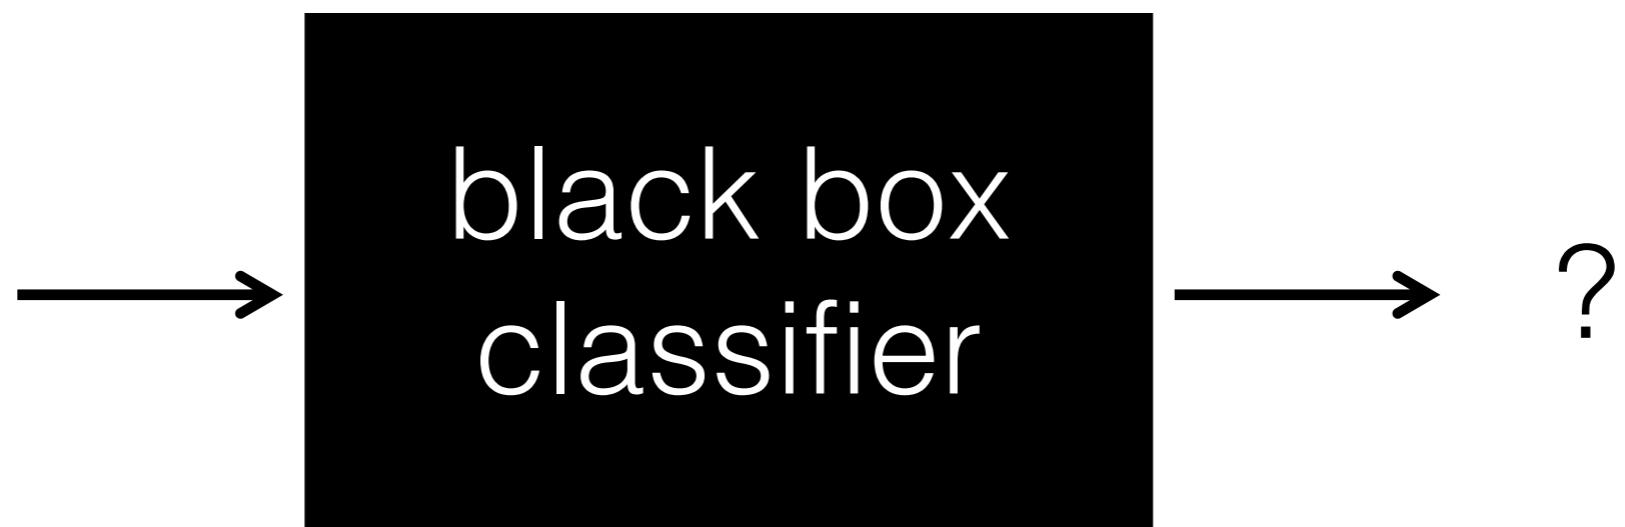
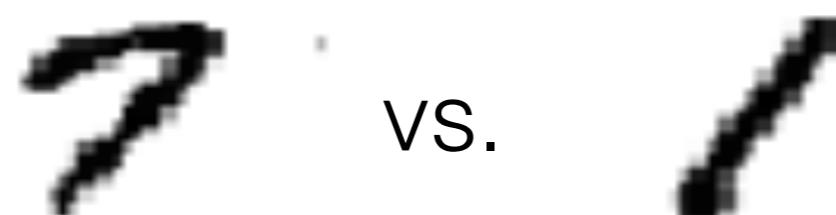


image X

3 6 8 1 7 9 6 6 9 1
6 7 5 7 8 6 3 4 8 5
2 1 7 9 7 1 2 8 4 6
4 8 1 9 0 1 8 8 9 4
7 6 1 8 6 4 1 5 6 0
7 5 9 2 6 5 8 1 9 7
1 2 2 2 2 3 4 4 8 0
0 2 3 8 0 7 3 8 5 7
0 1 4 6 4 6 0 2 4 3
7 1 2 8 7 6 9 8 6 1

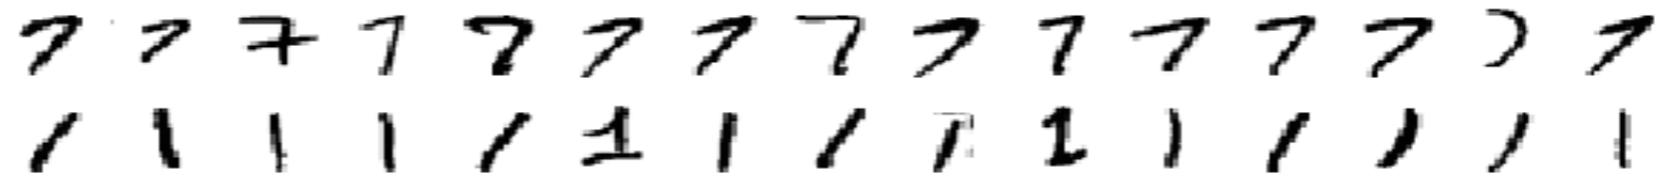
Warm-up Example: Binary Digit Classification



Learning Approach to Digit Recognition

- **Collect Training Images**

- Positive:



- Negative:



- **Training Time**

- Compute **feature vectors** for positive and negative example images
- Train a **classifier**

- **Test Time**

- Compute feature vector on new test image:
- Evaluate classifier



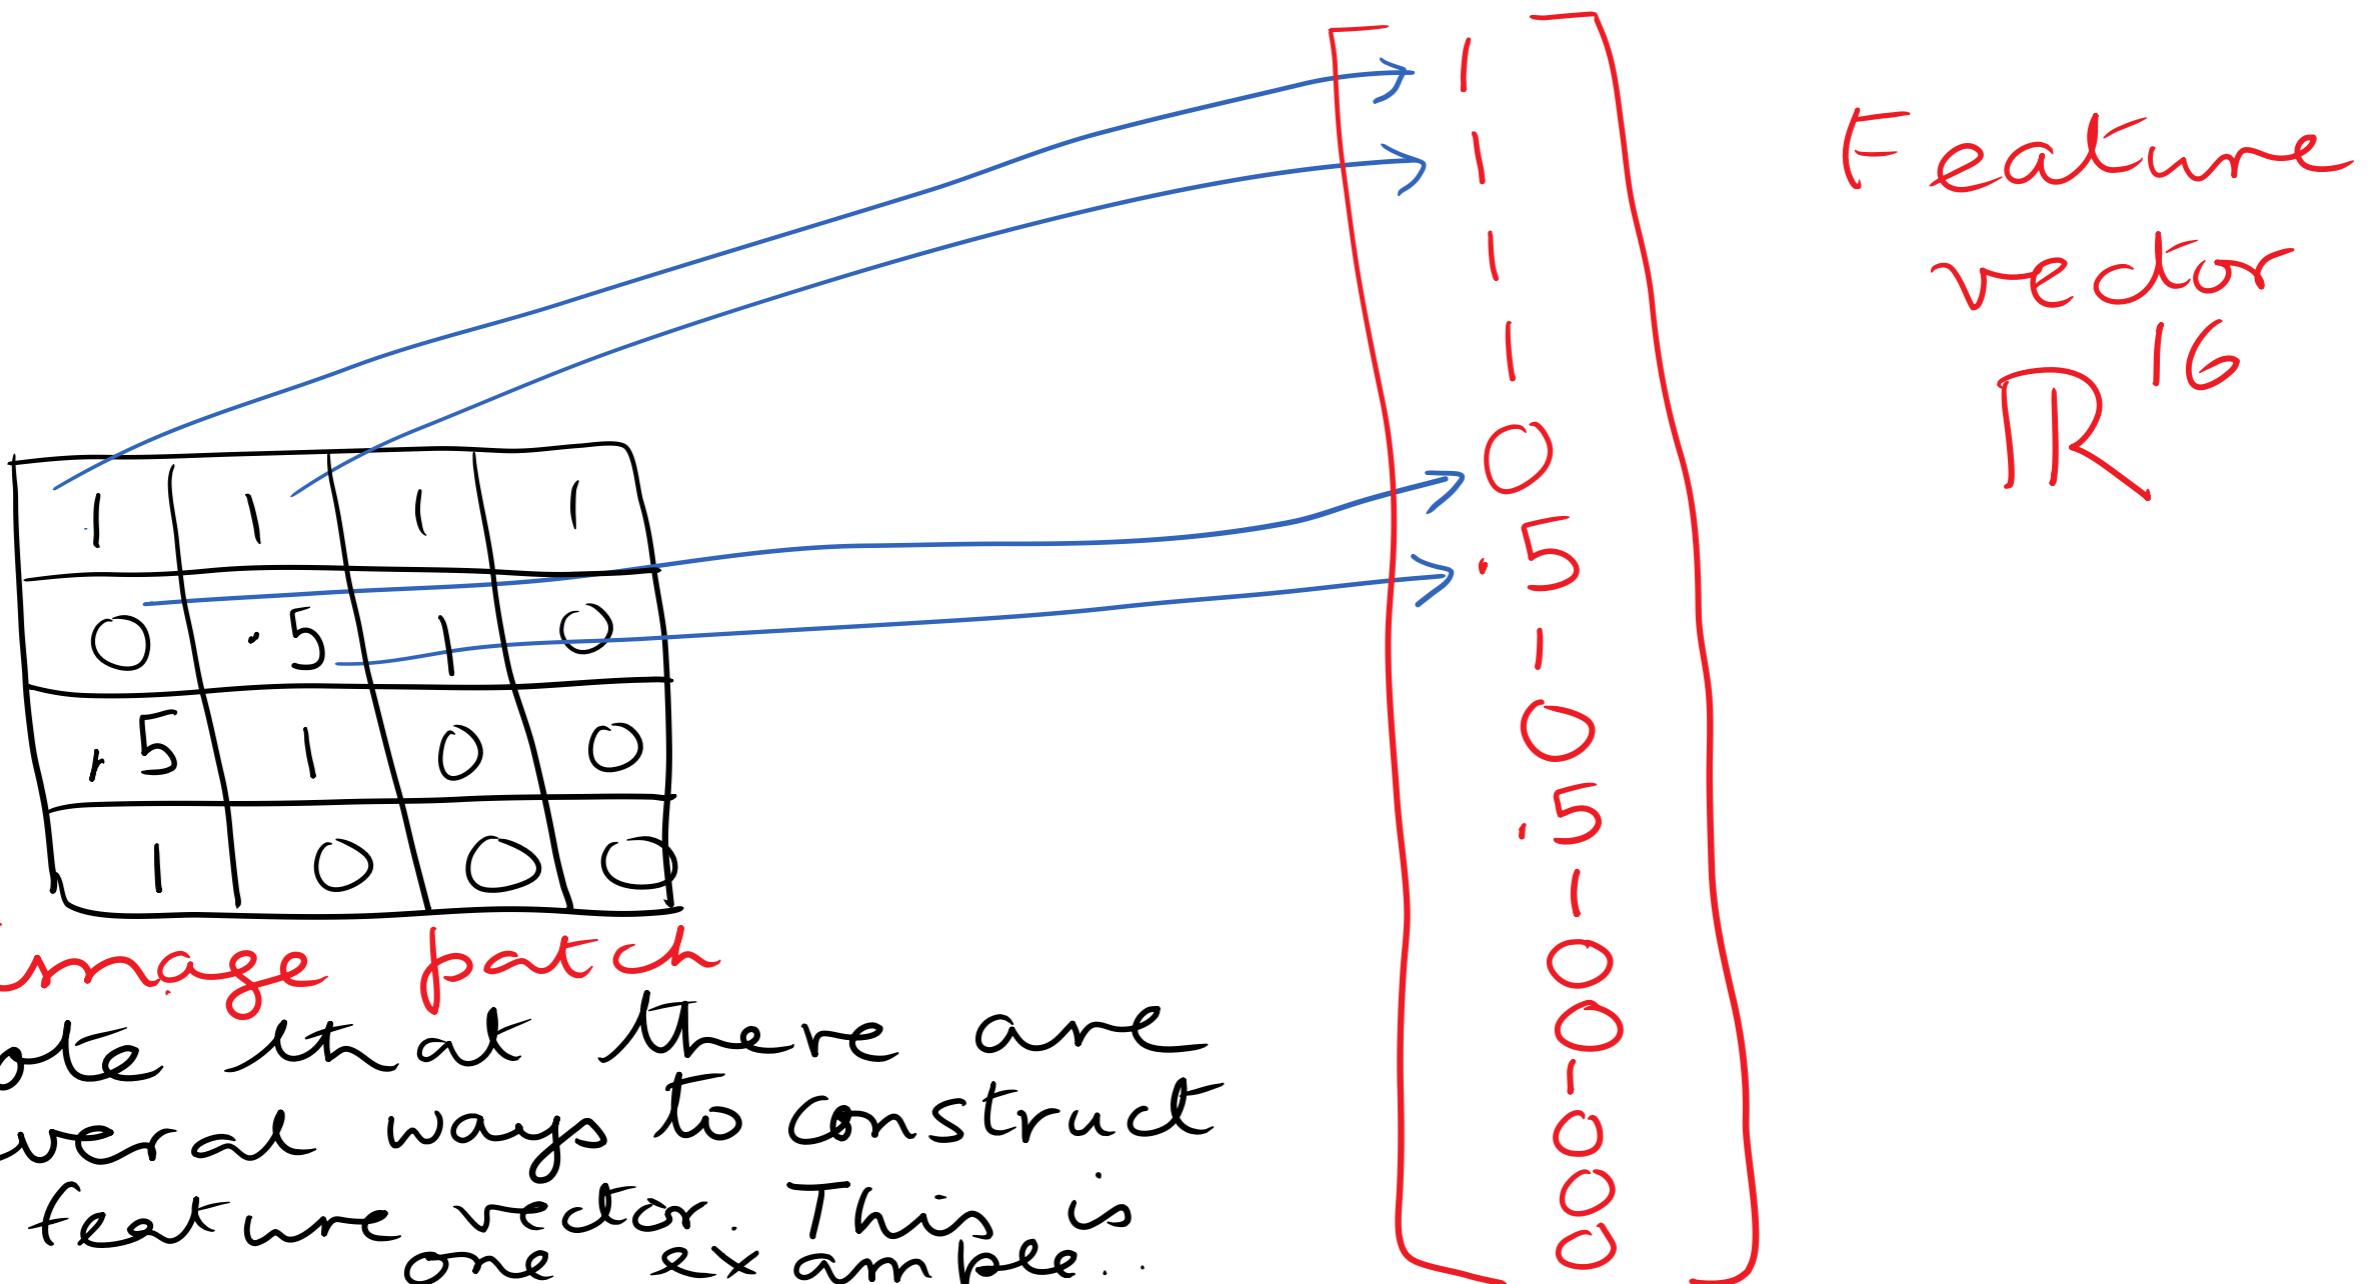
Let us take an example...



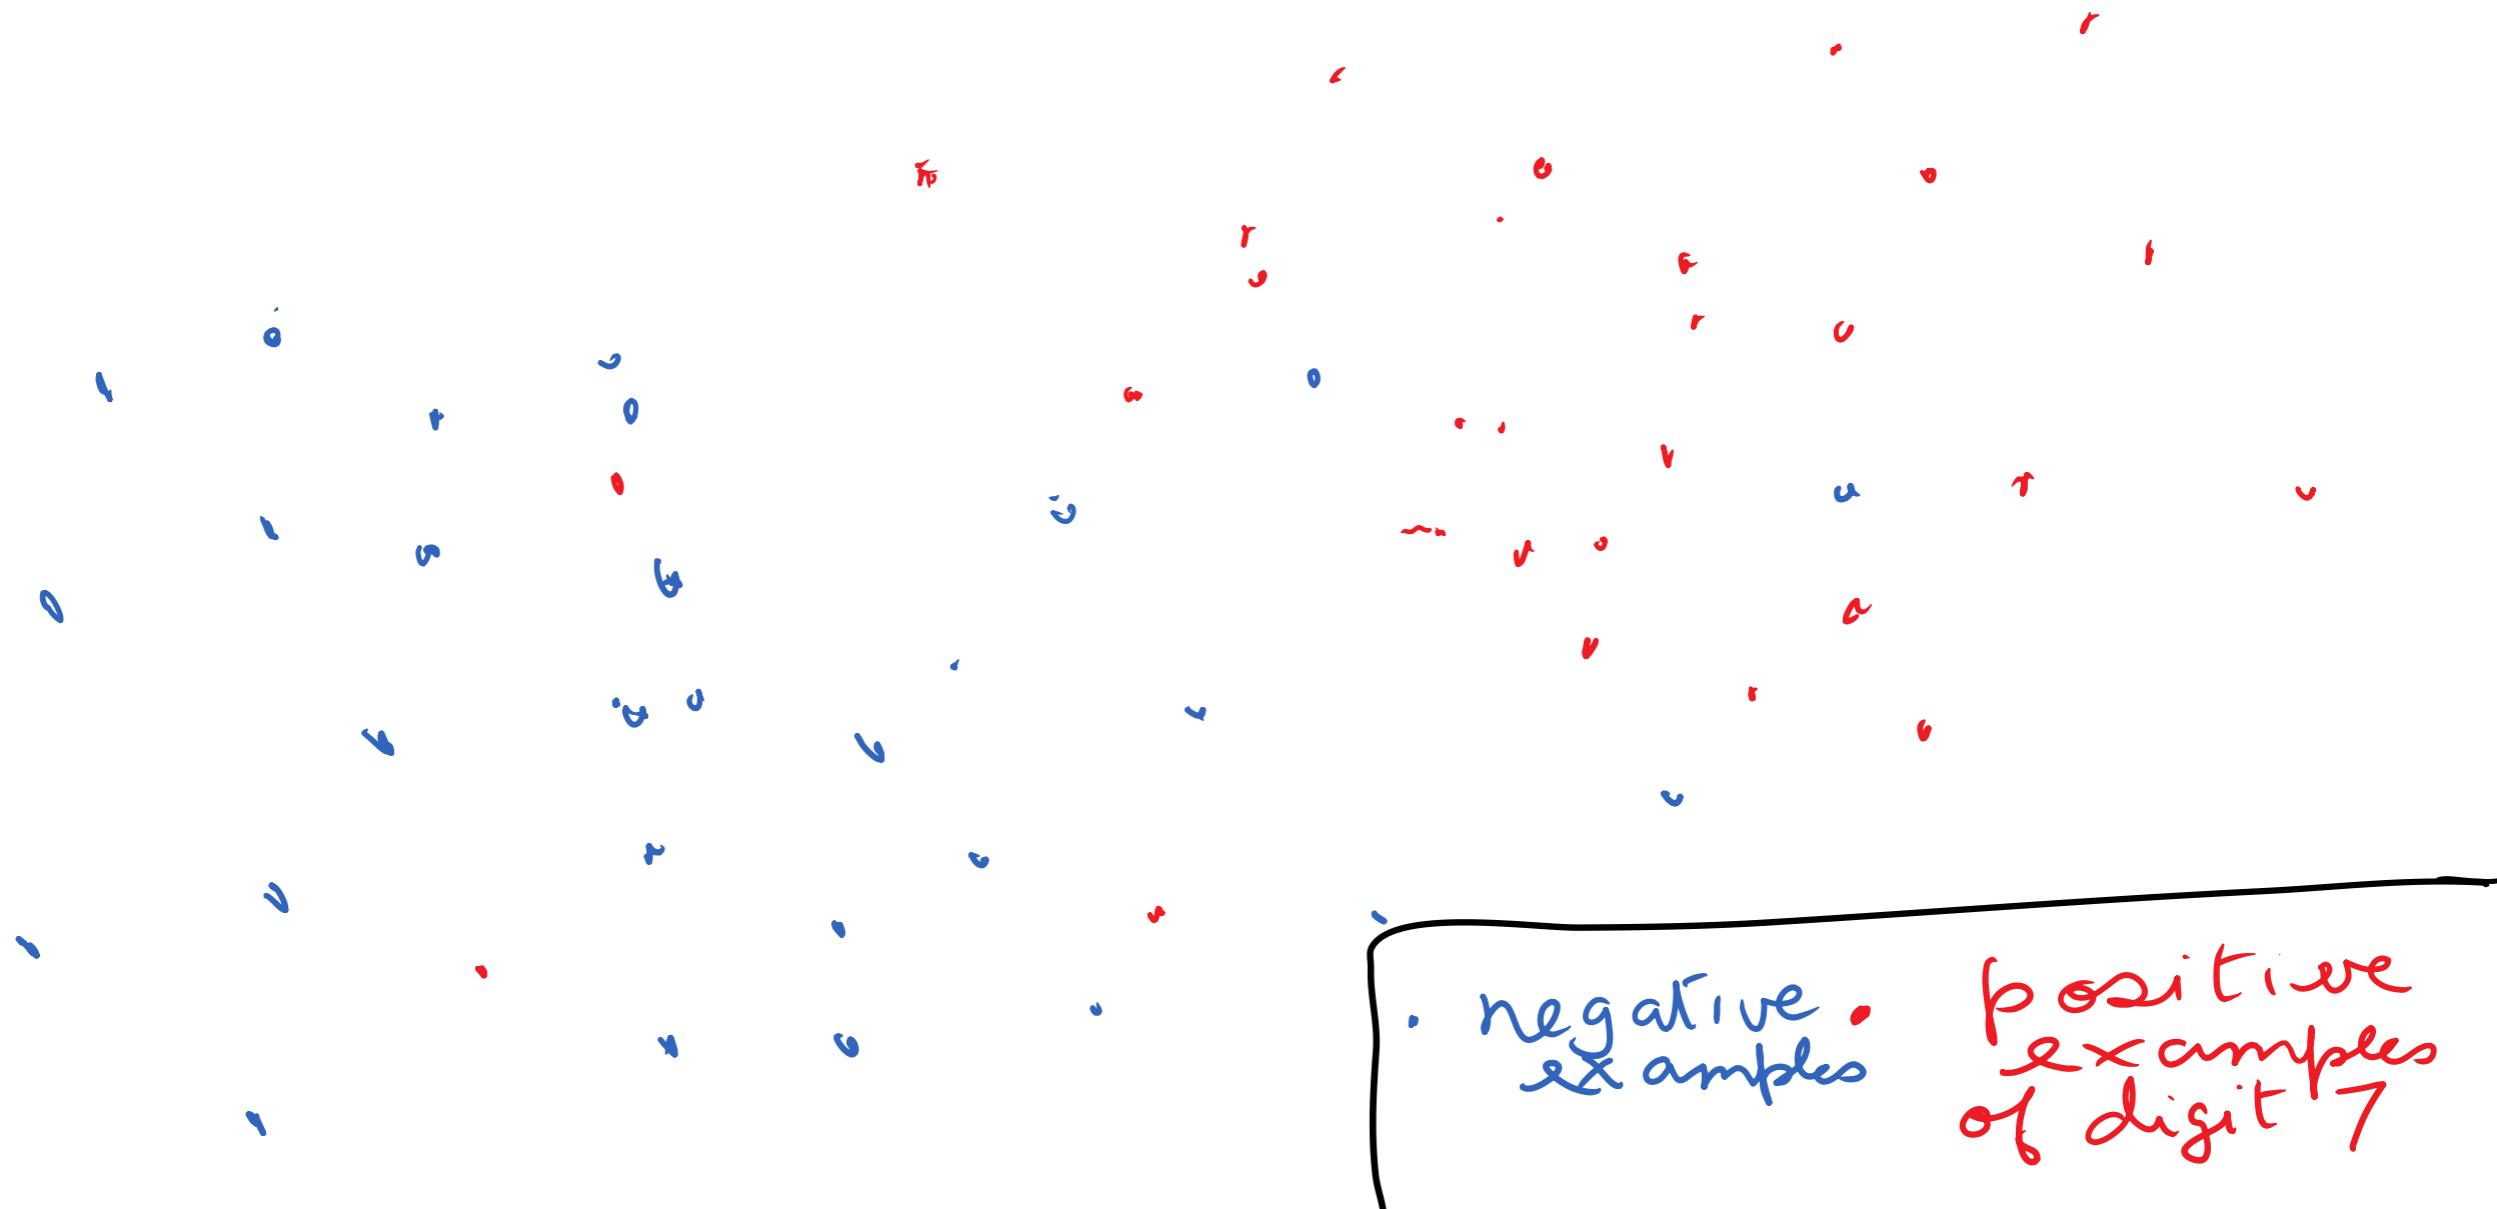
1	1	1	1
0	.5	1	0
.5	1	0	0
1	0	0	0

image
patch

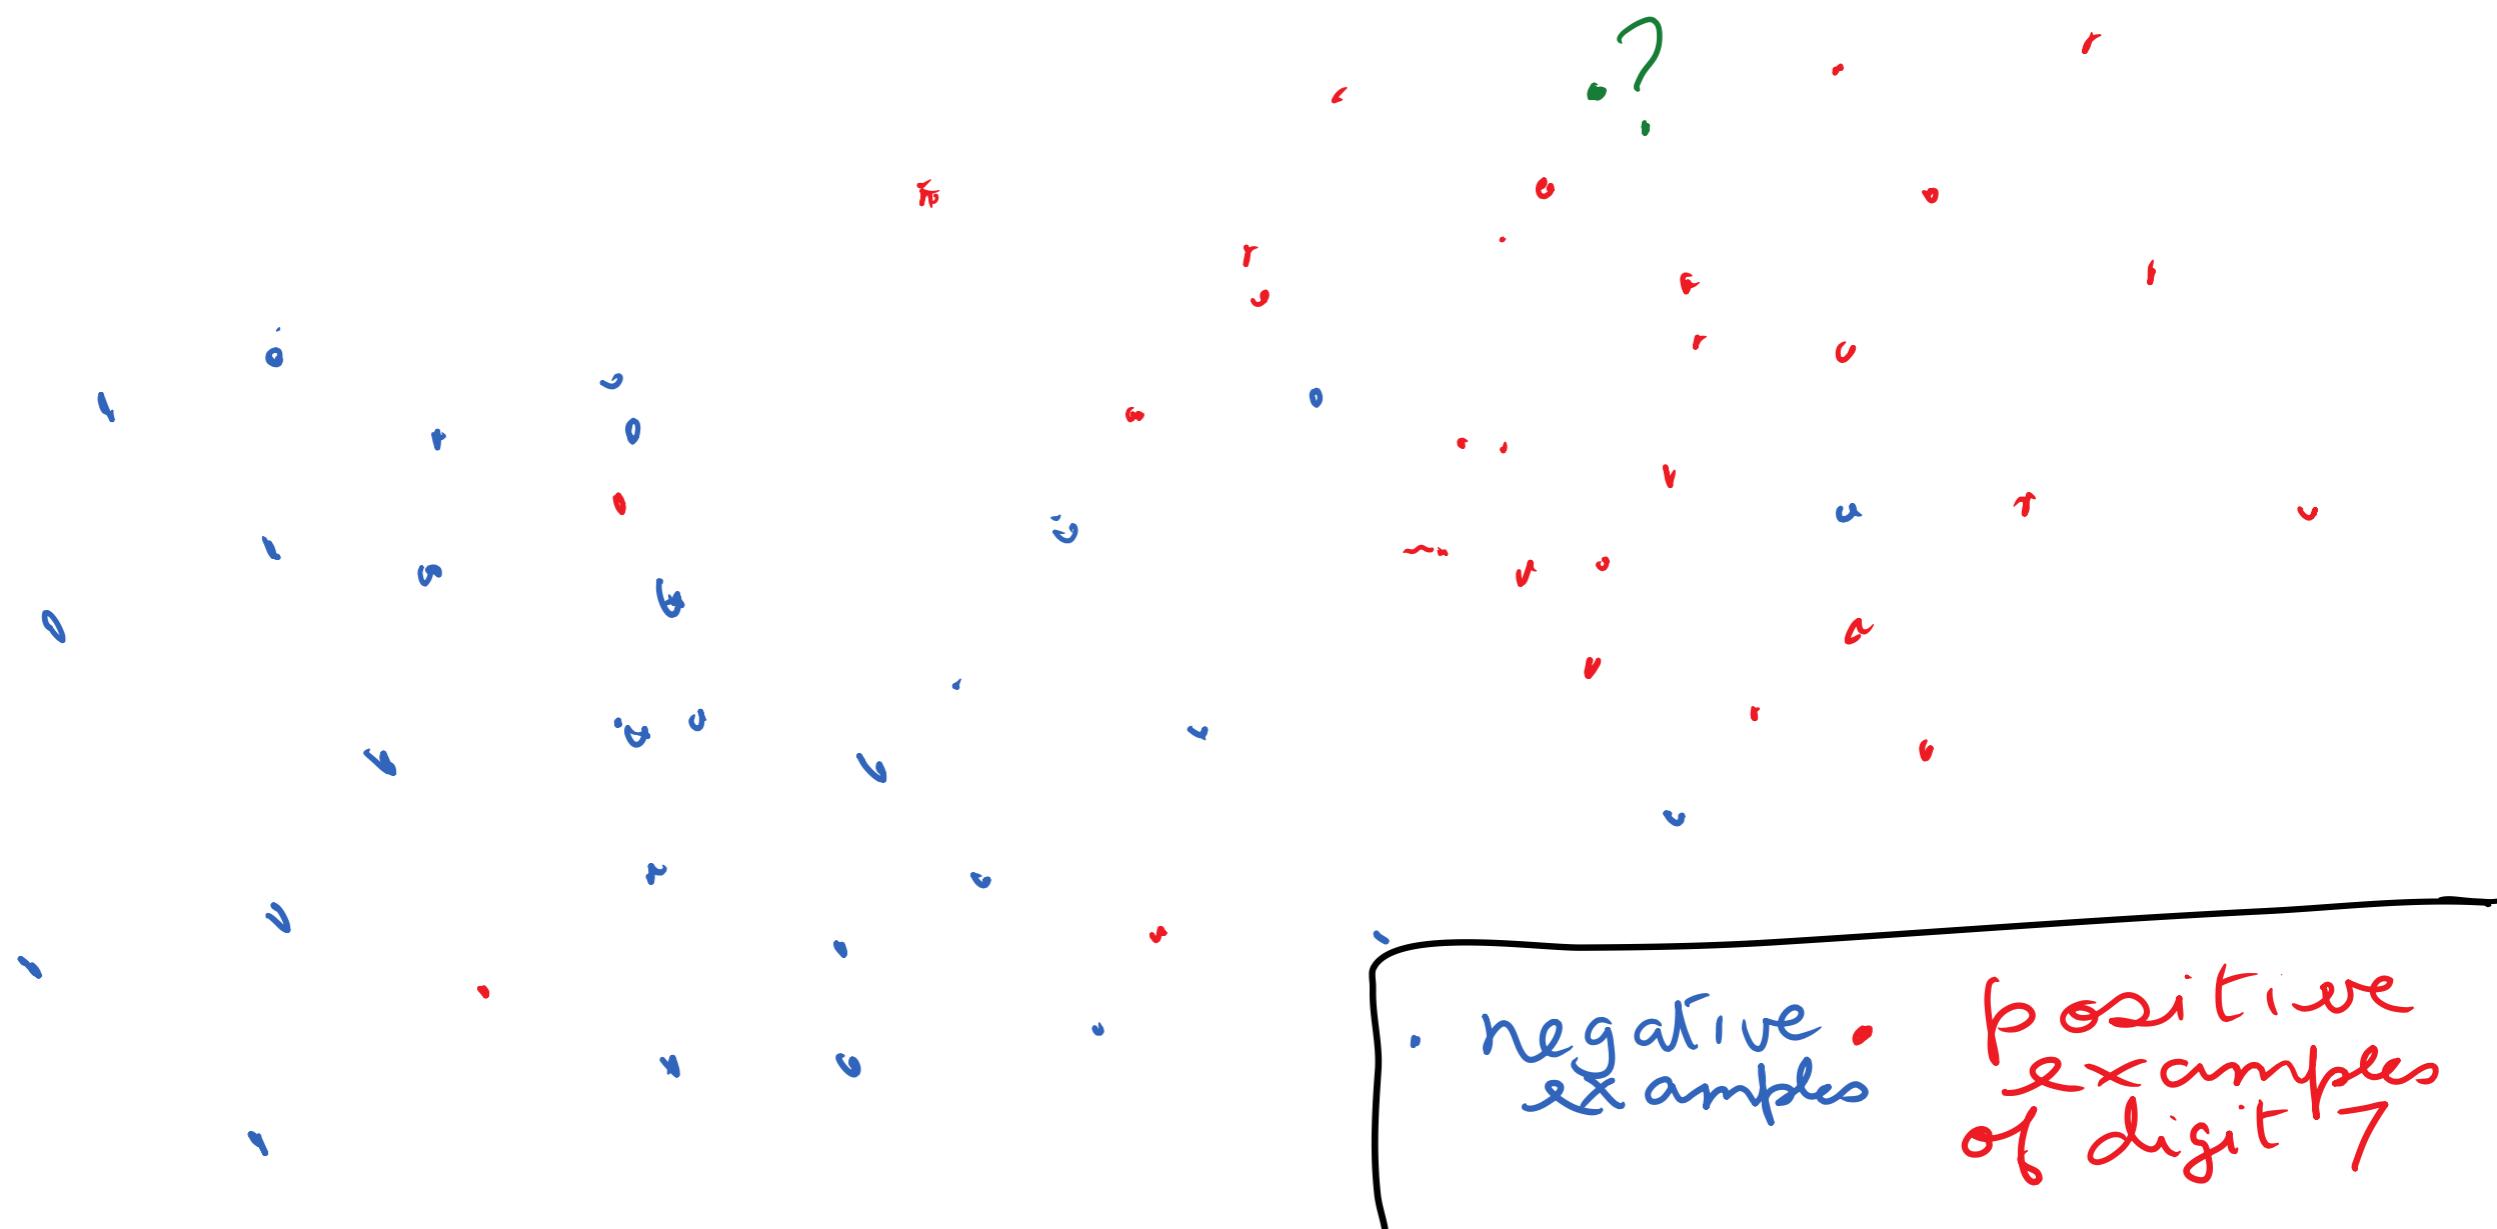
Let us take an example...



In feature space, positive and negative examples are just points...

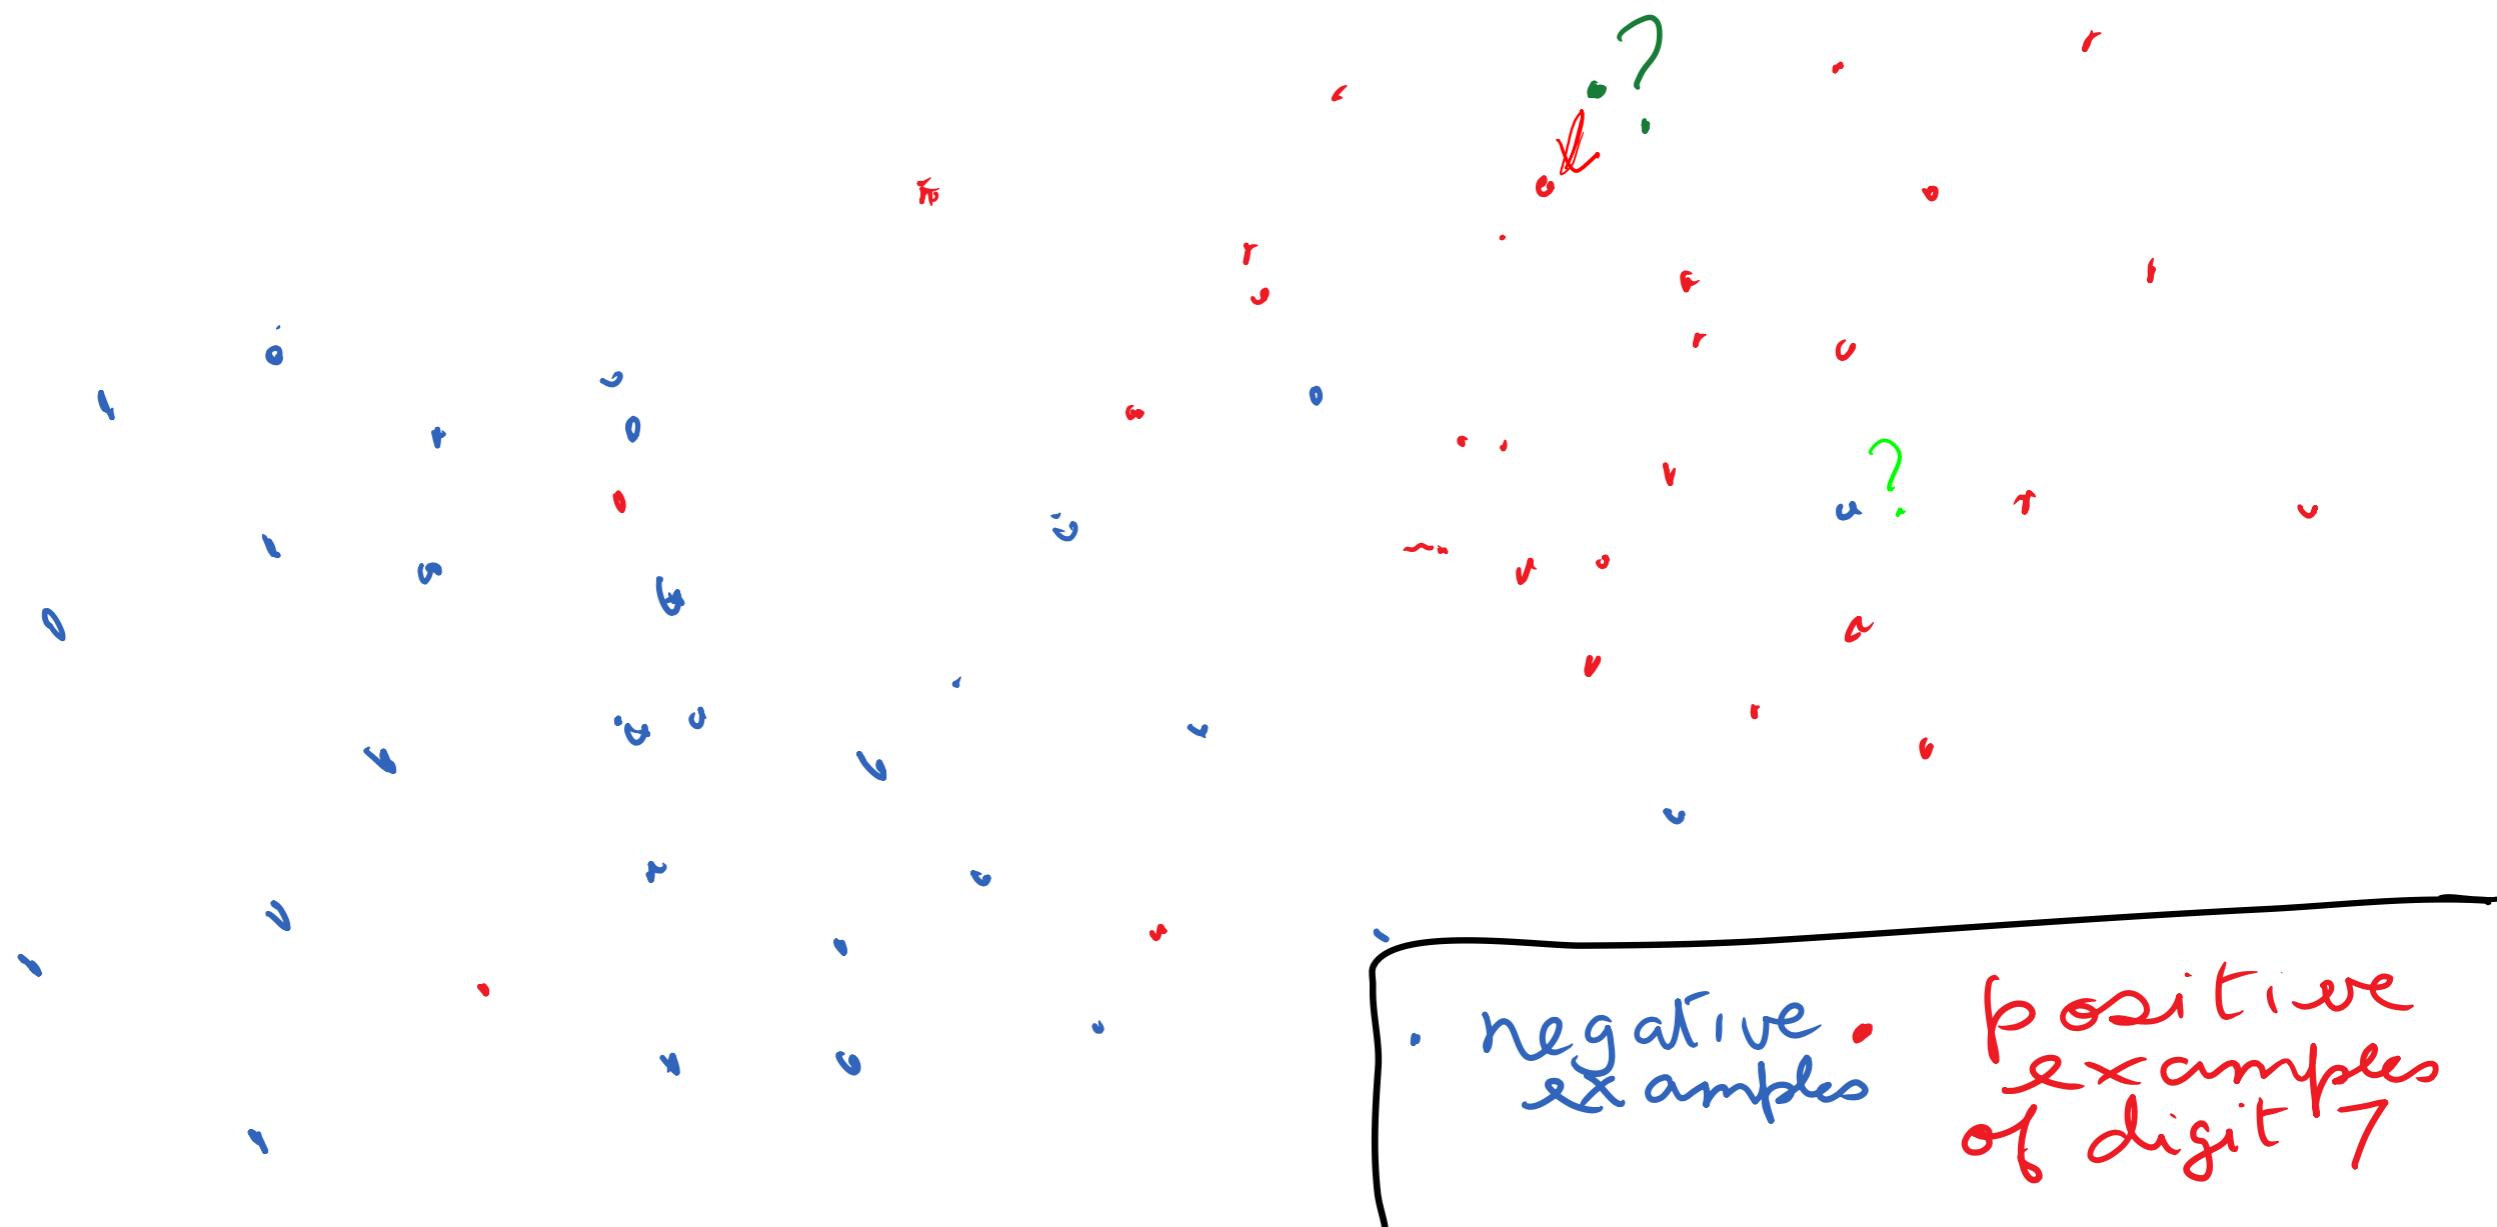


How do we classify a new point?

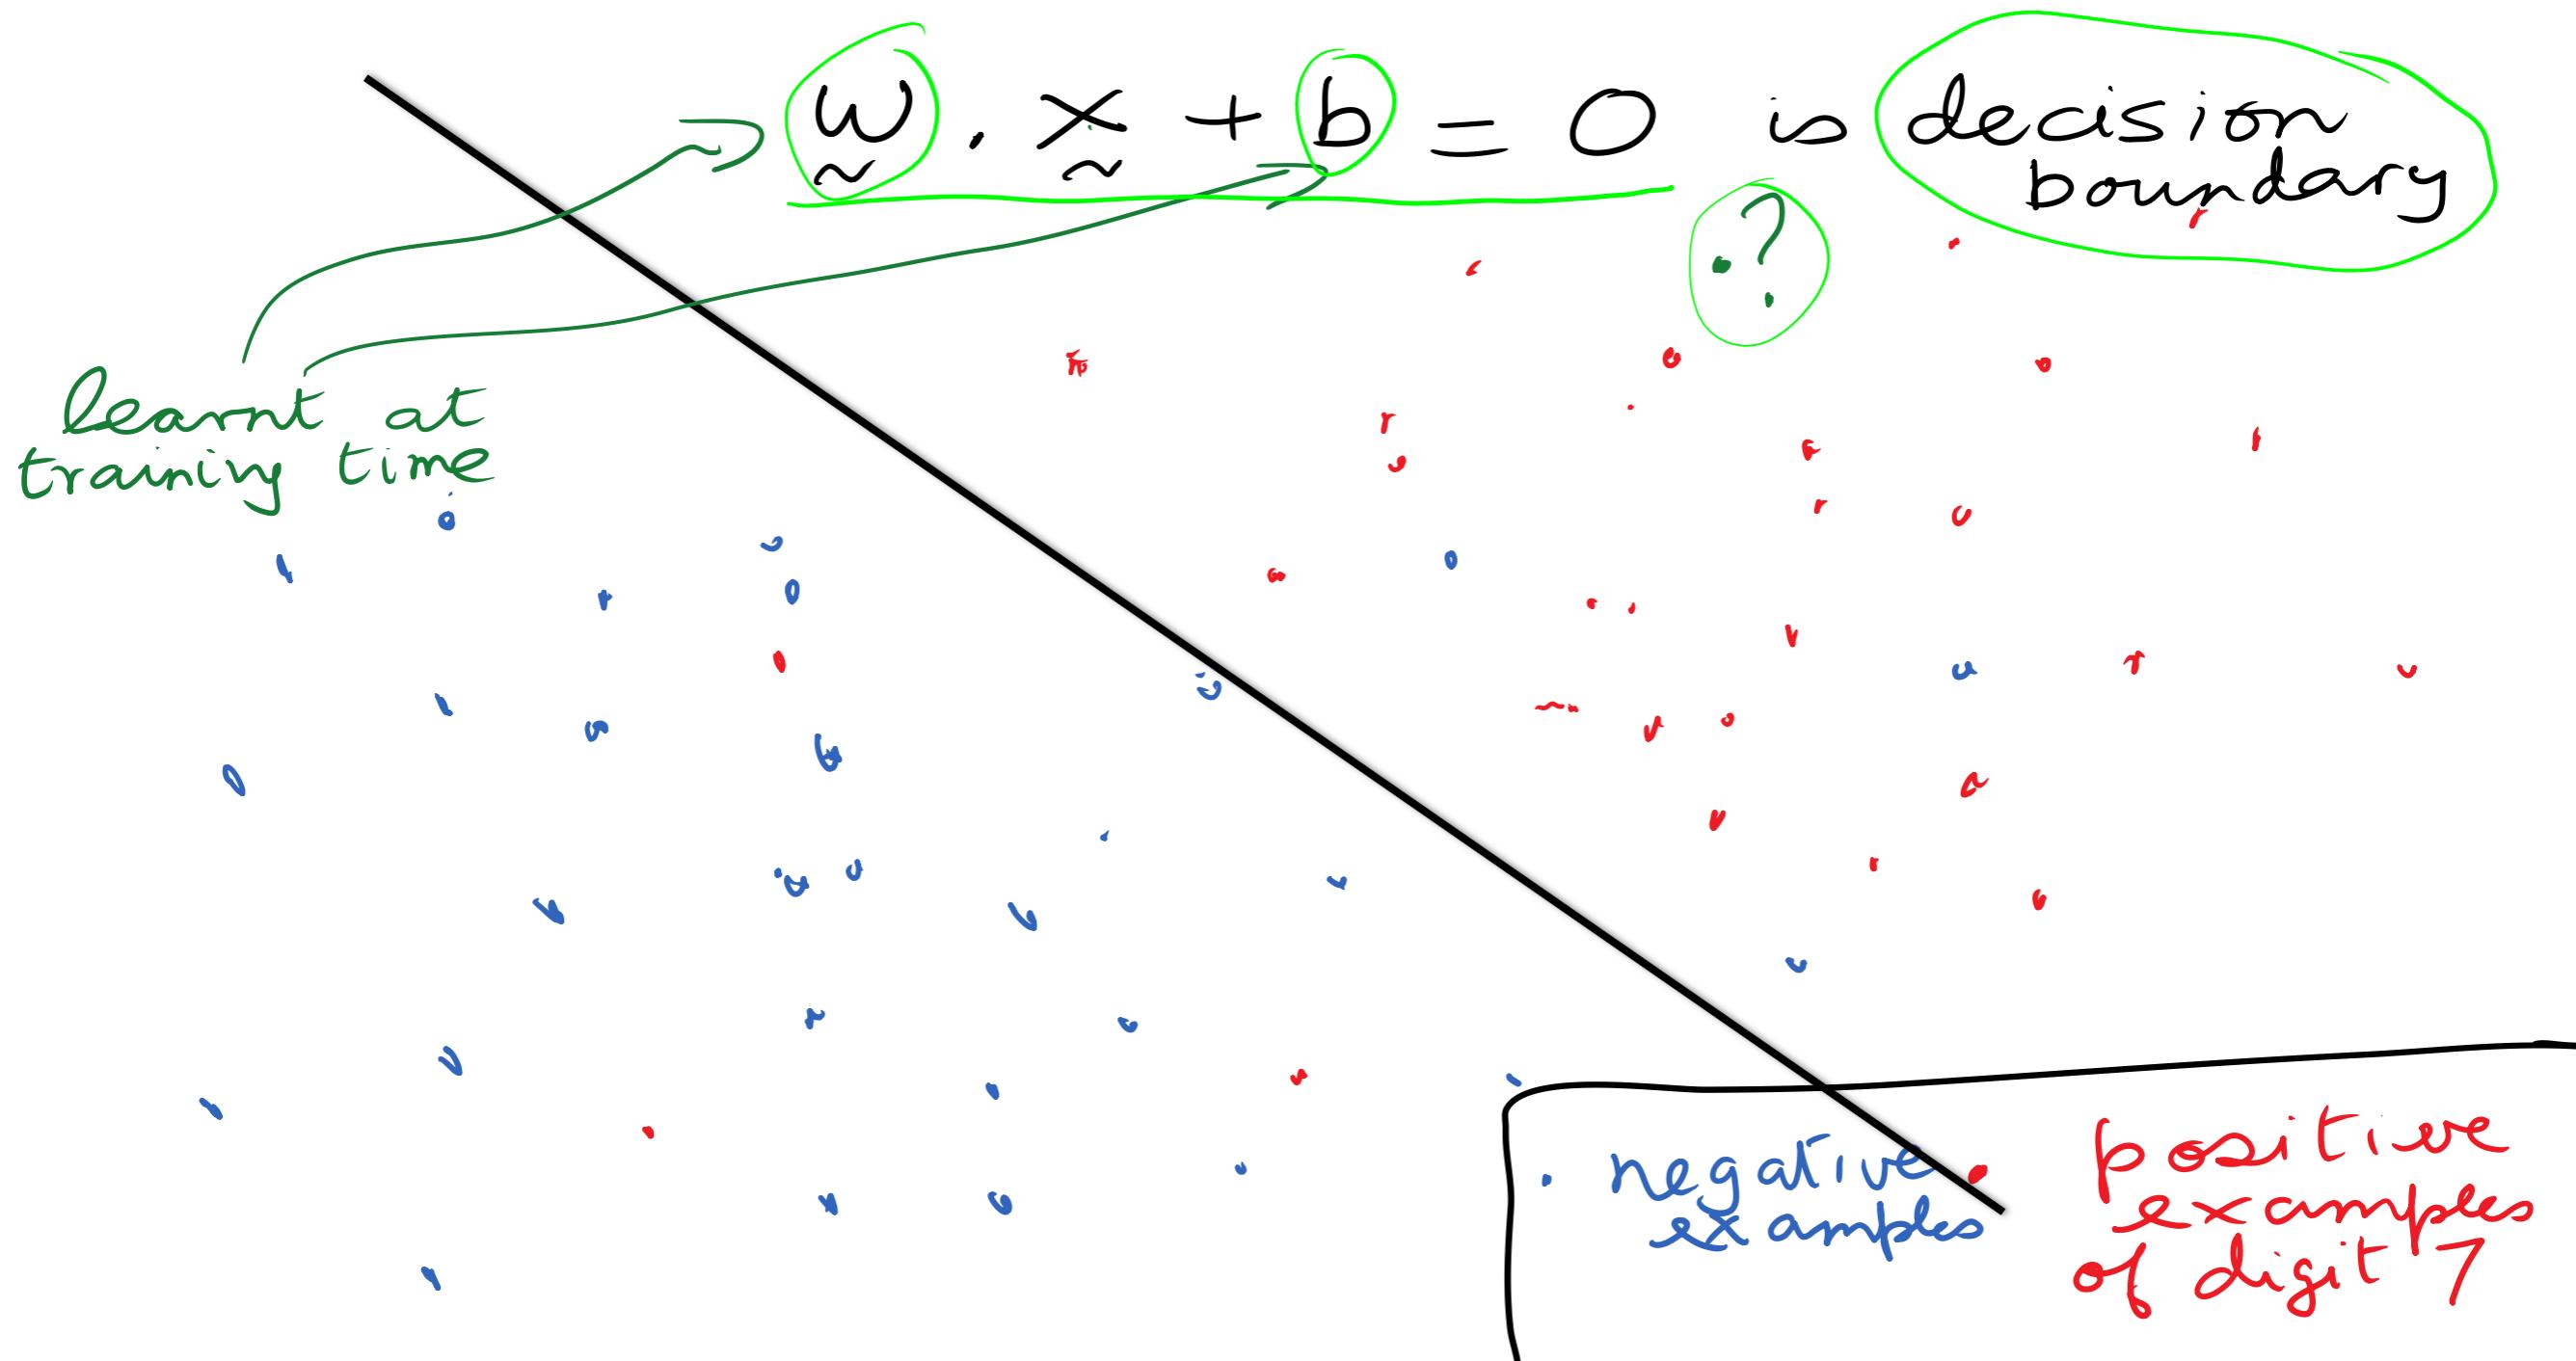


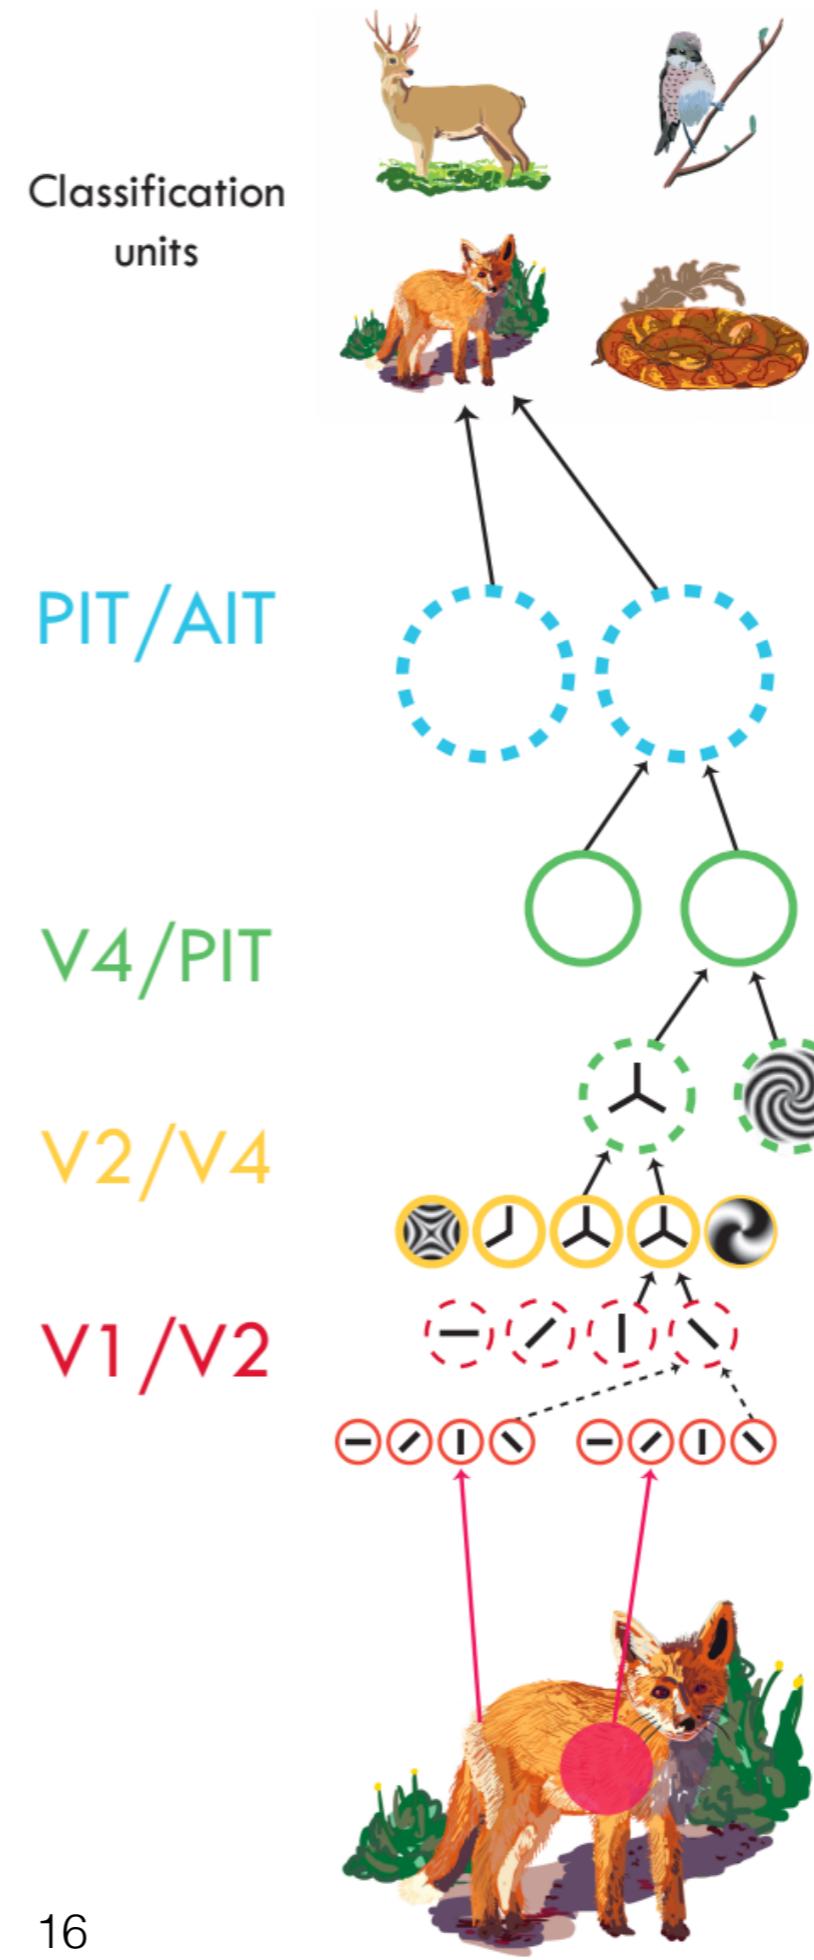
Nearest neighbor rule

“transfer label of nearest example”

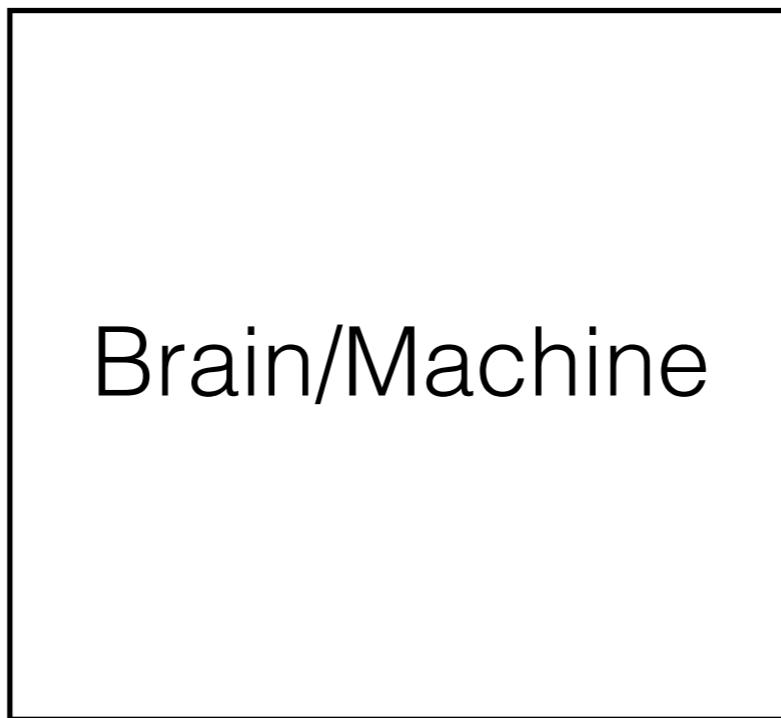
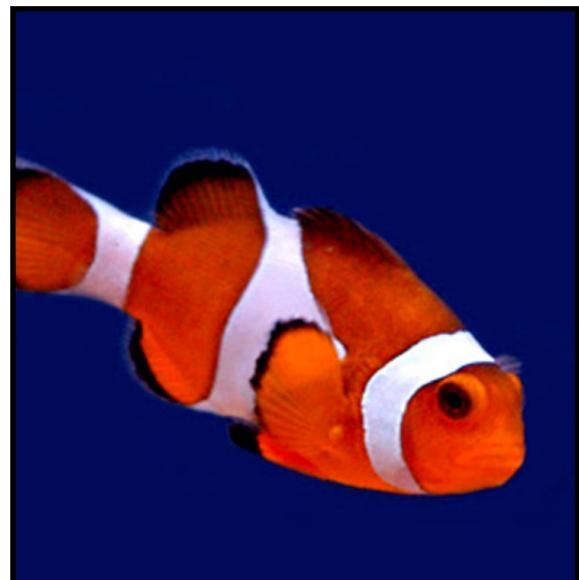


Linear classifier rule



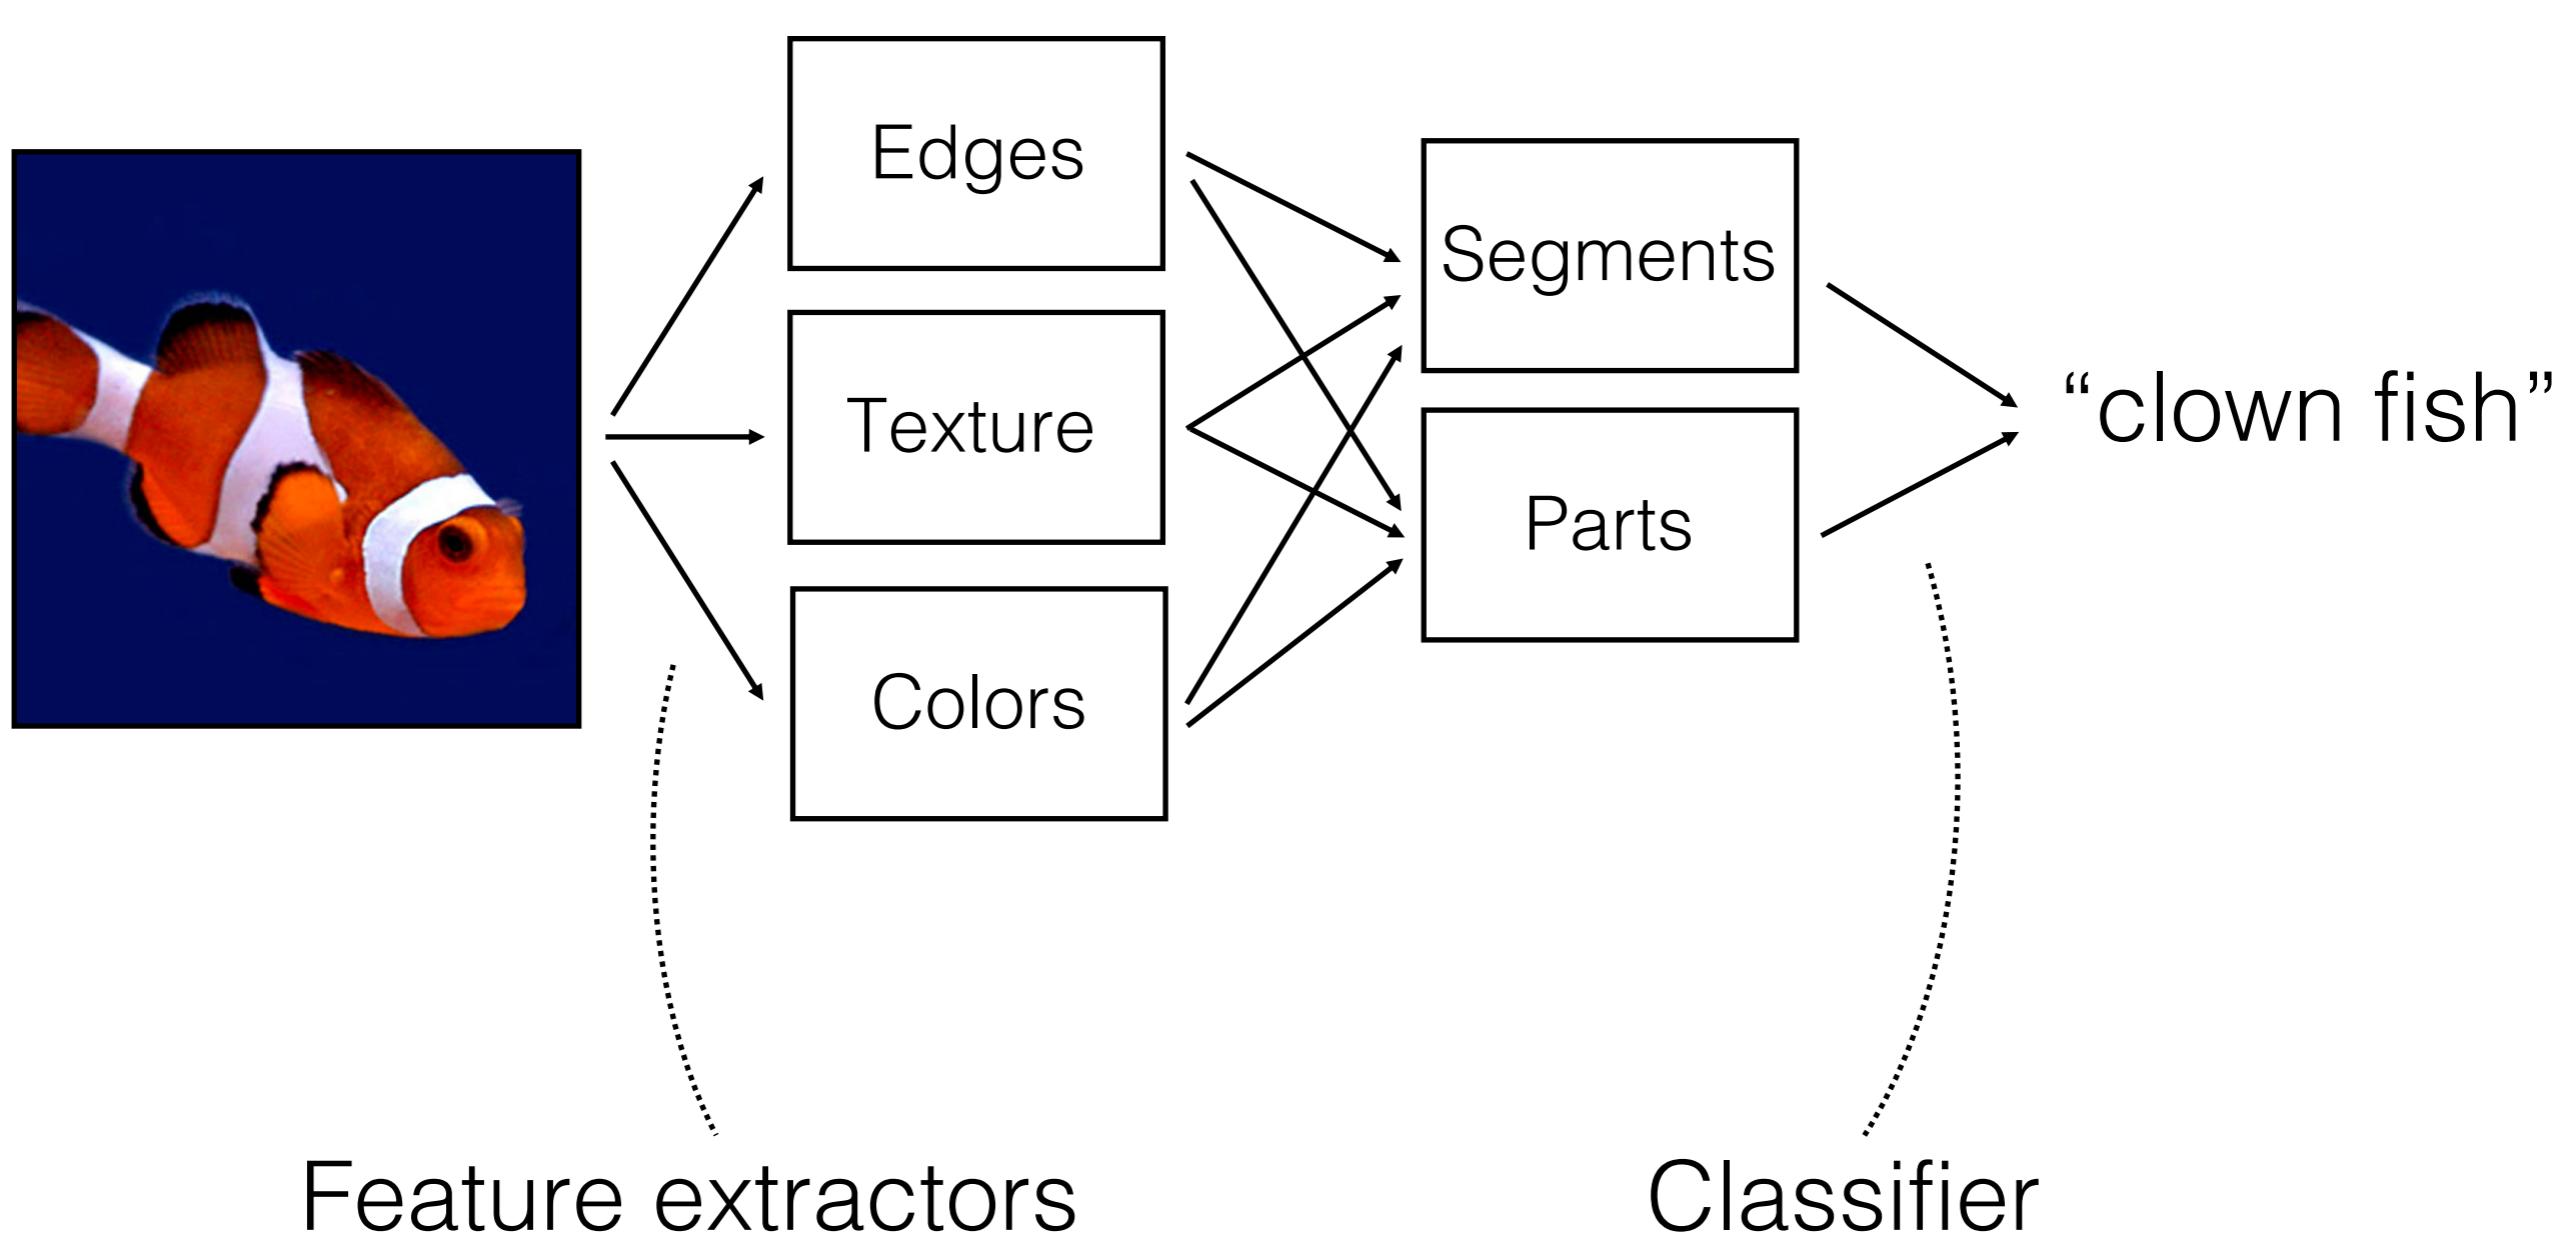


Basic idea

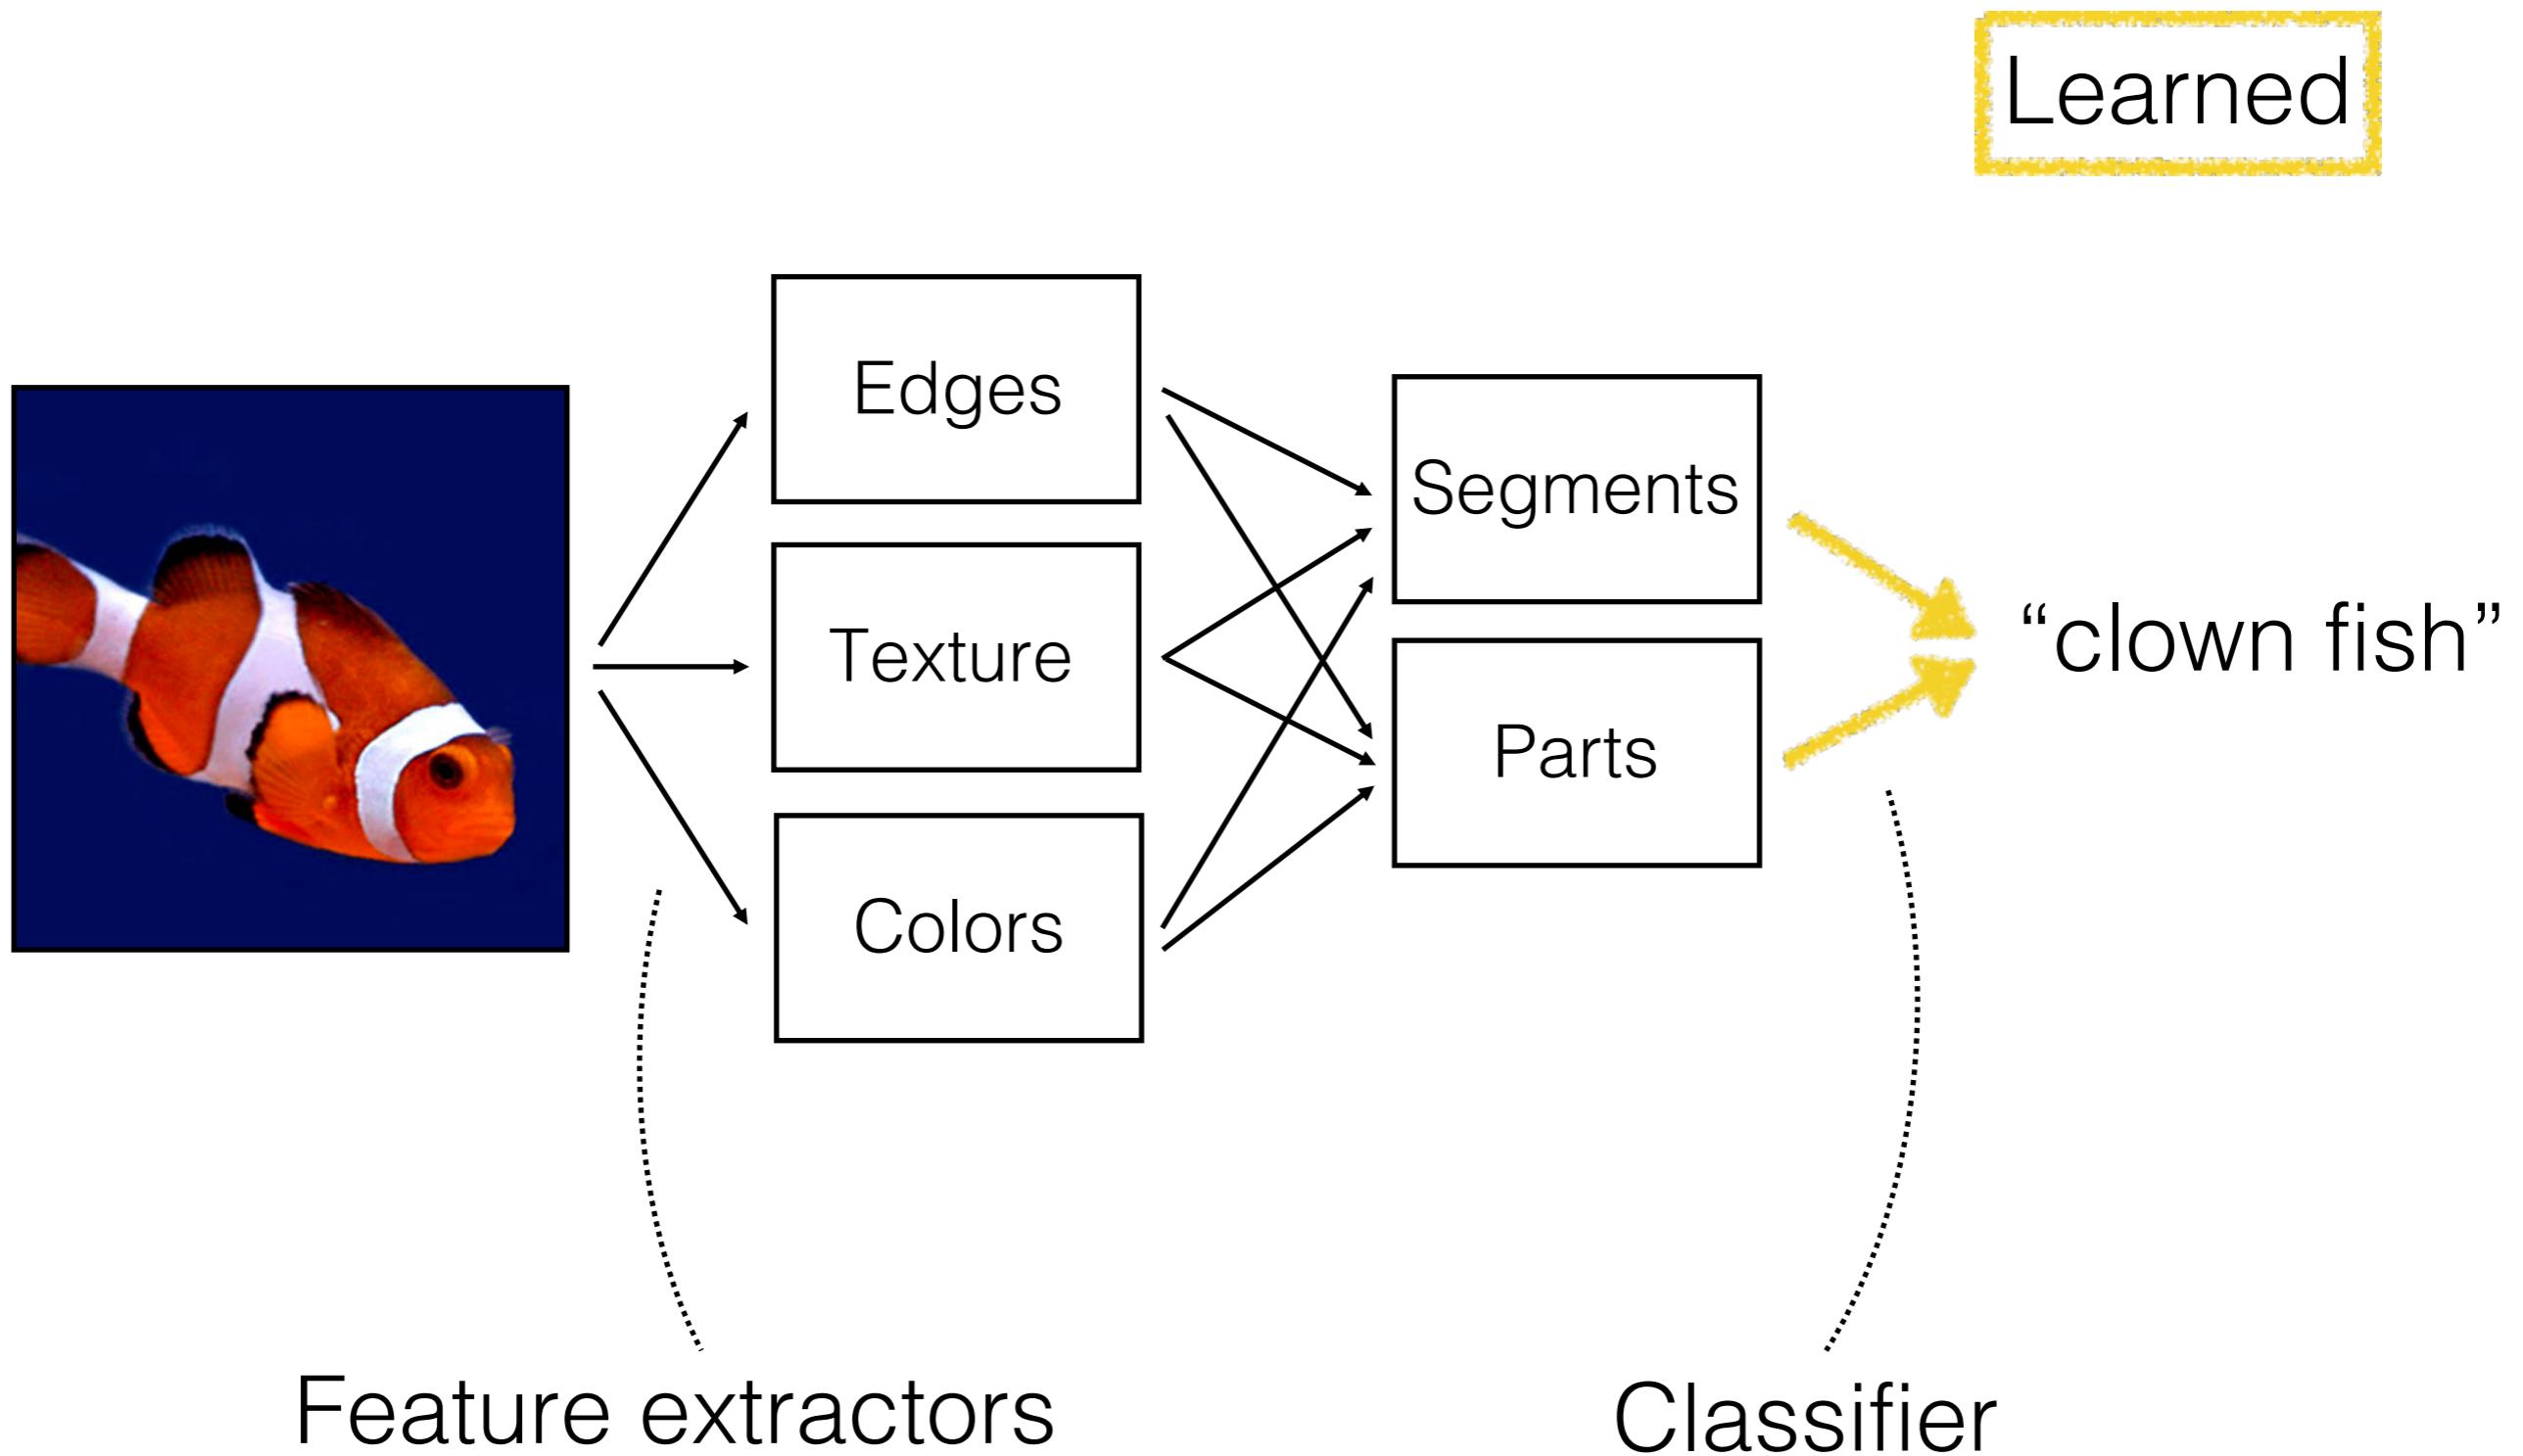


"clown fish"

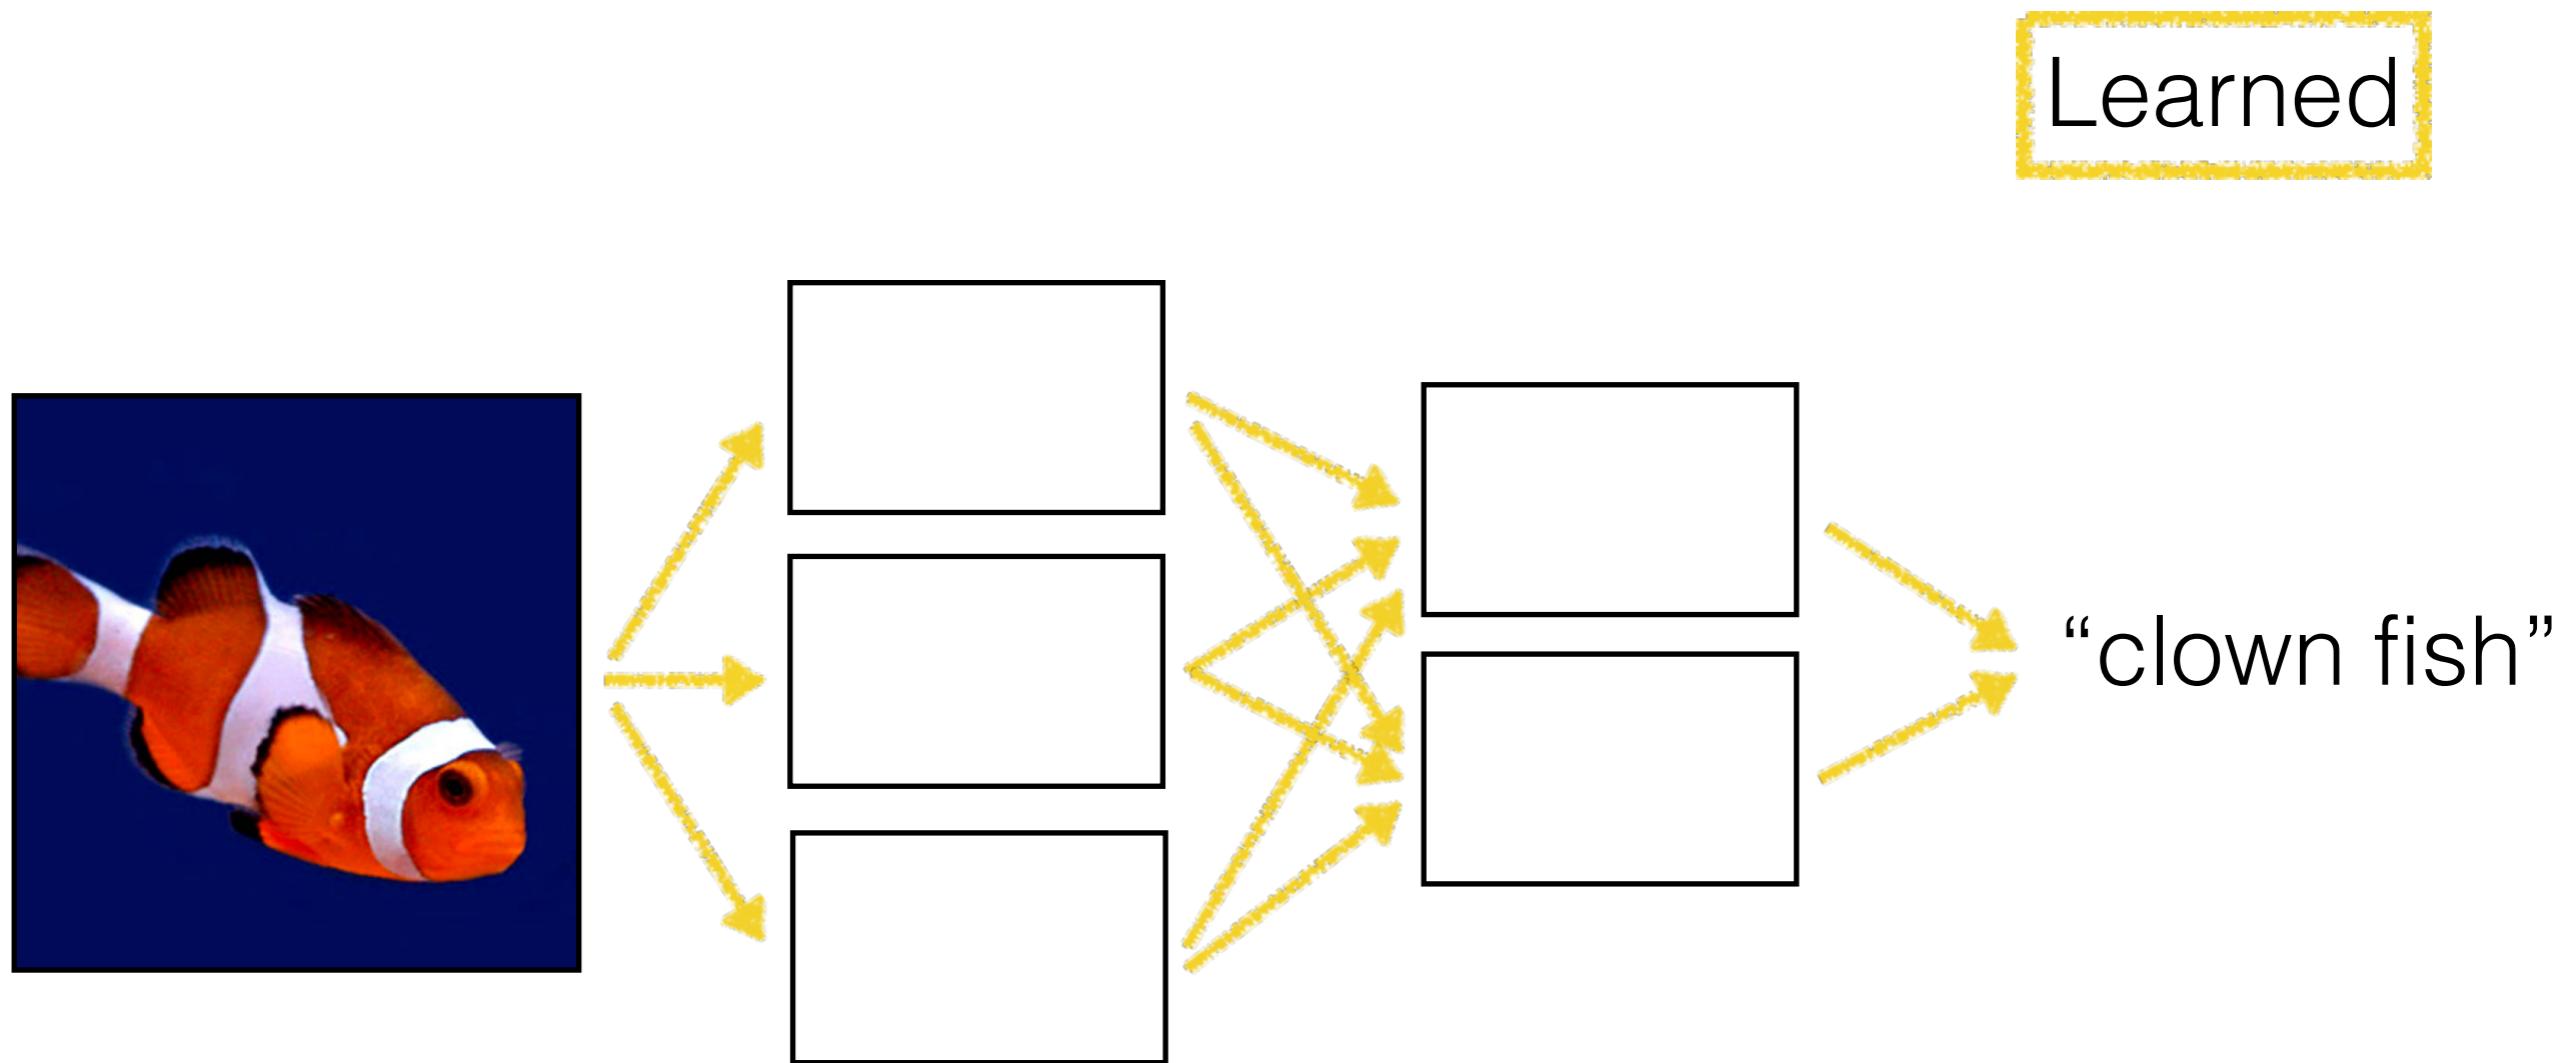
Object recognition



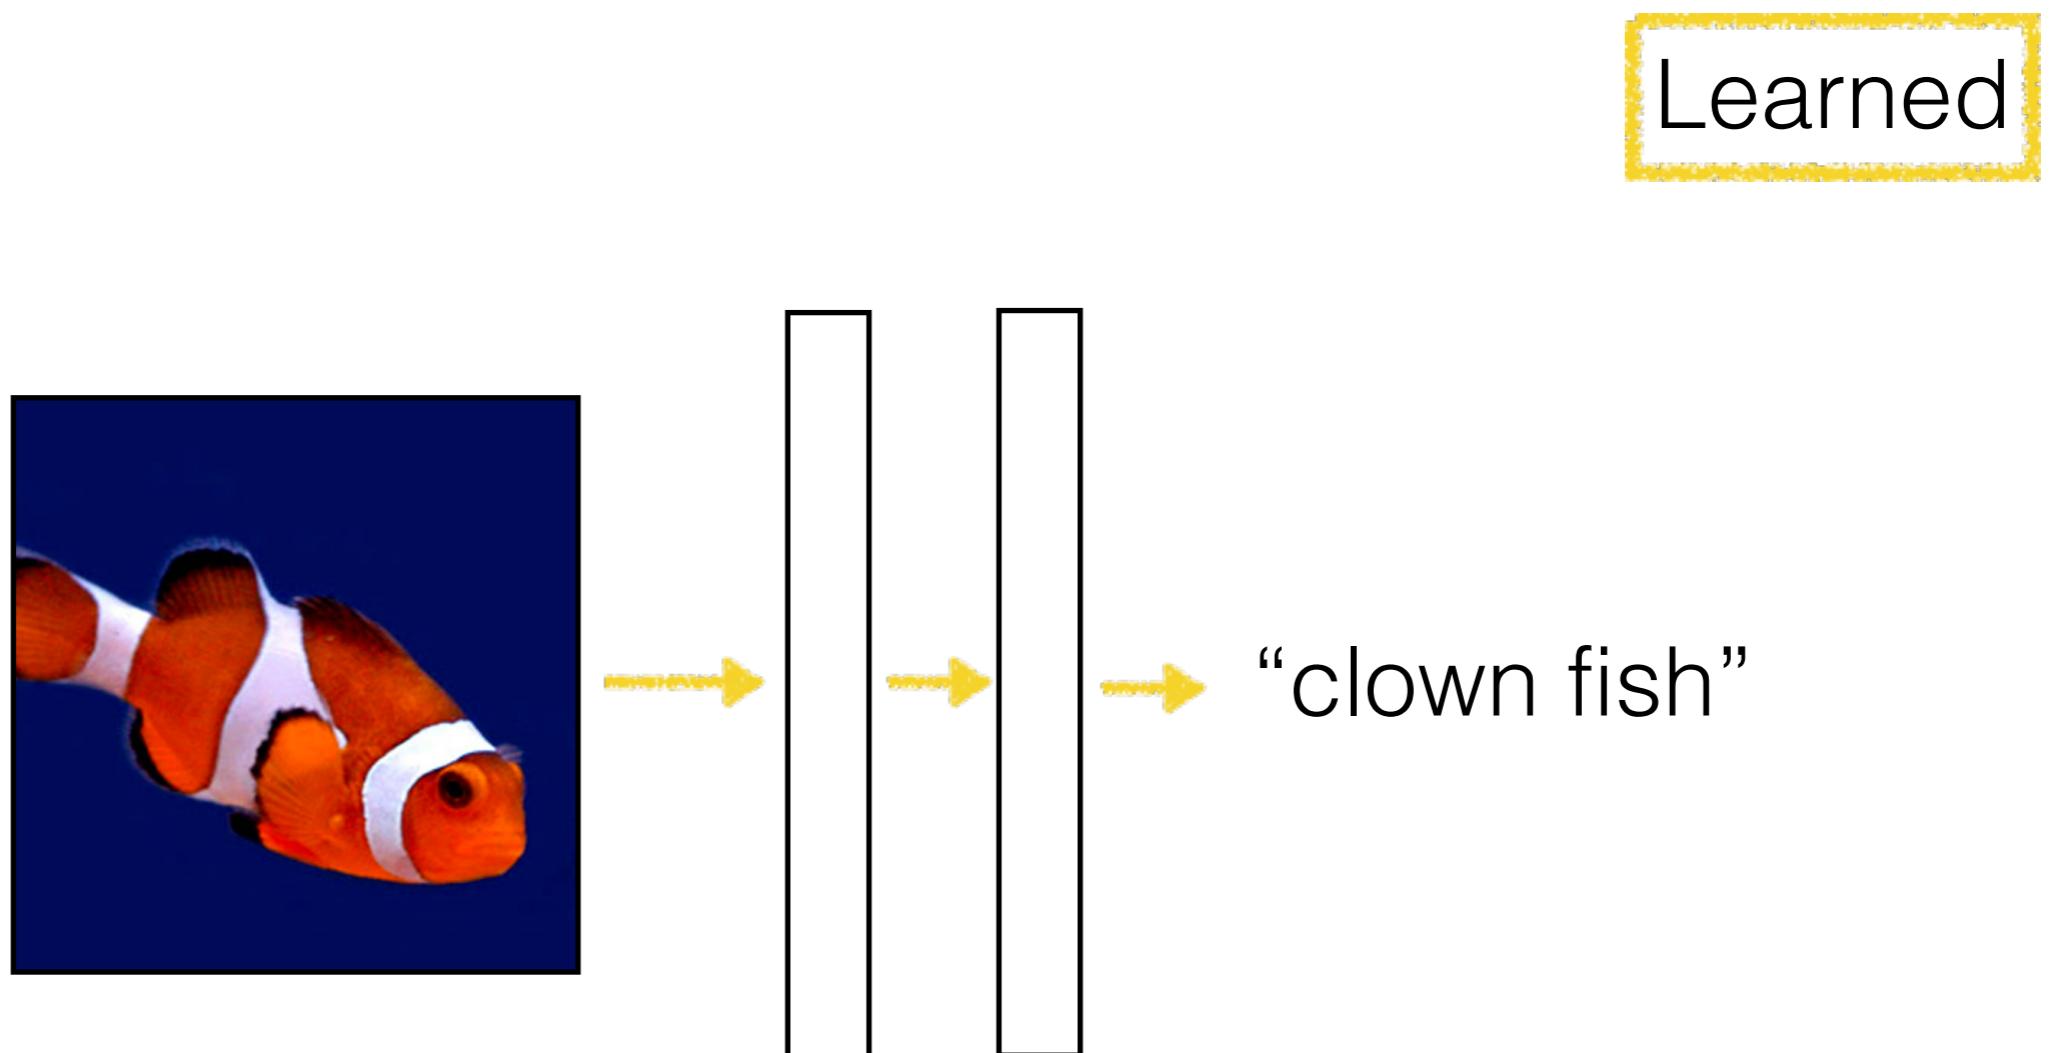
Object recognition



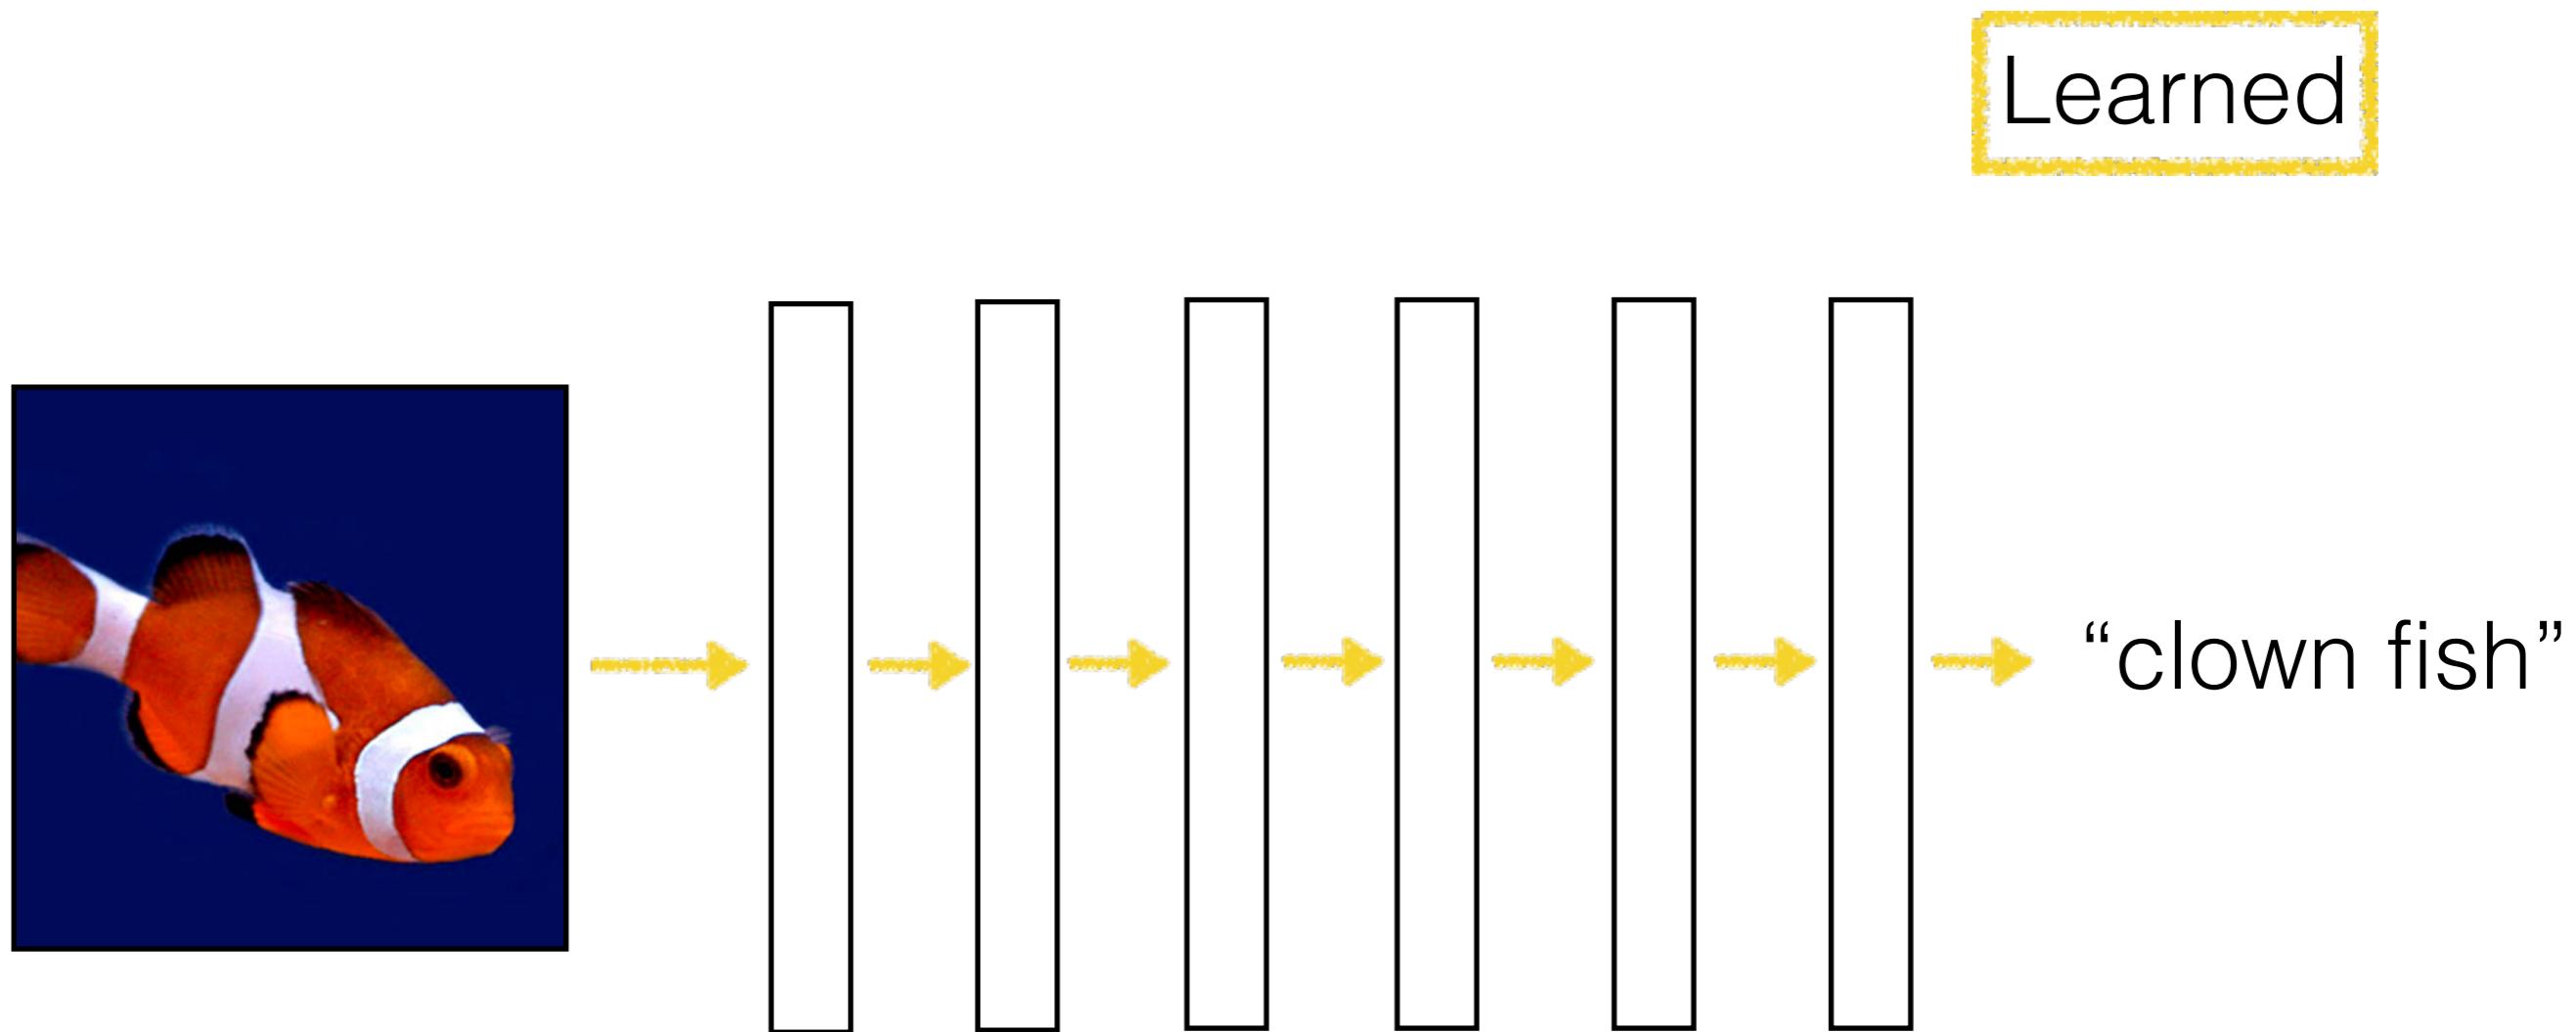
Neural network



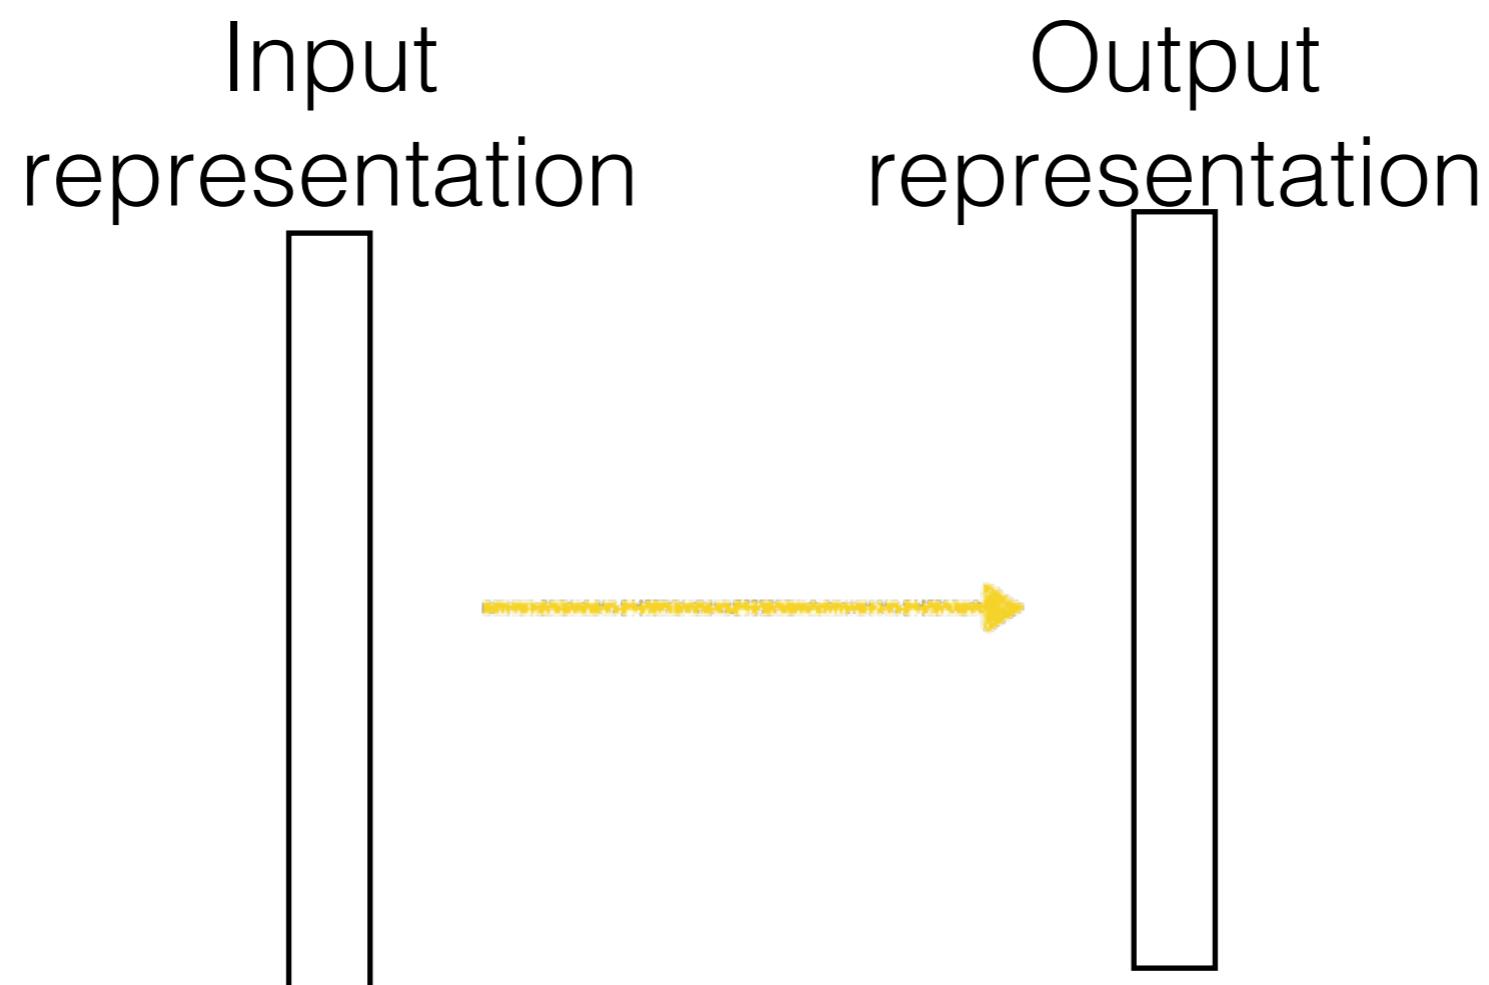
Neural network



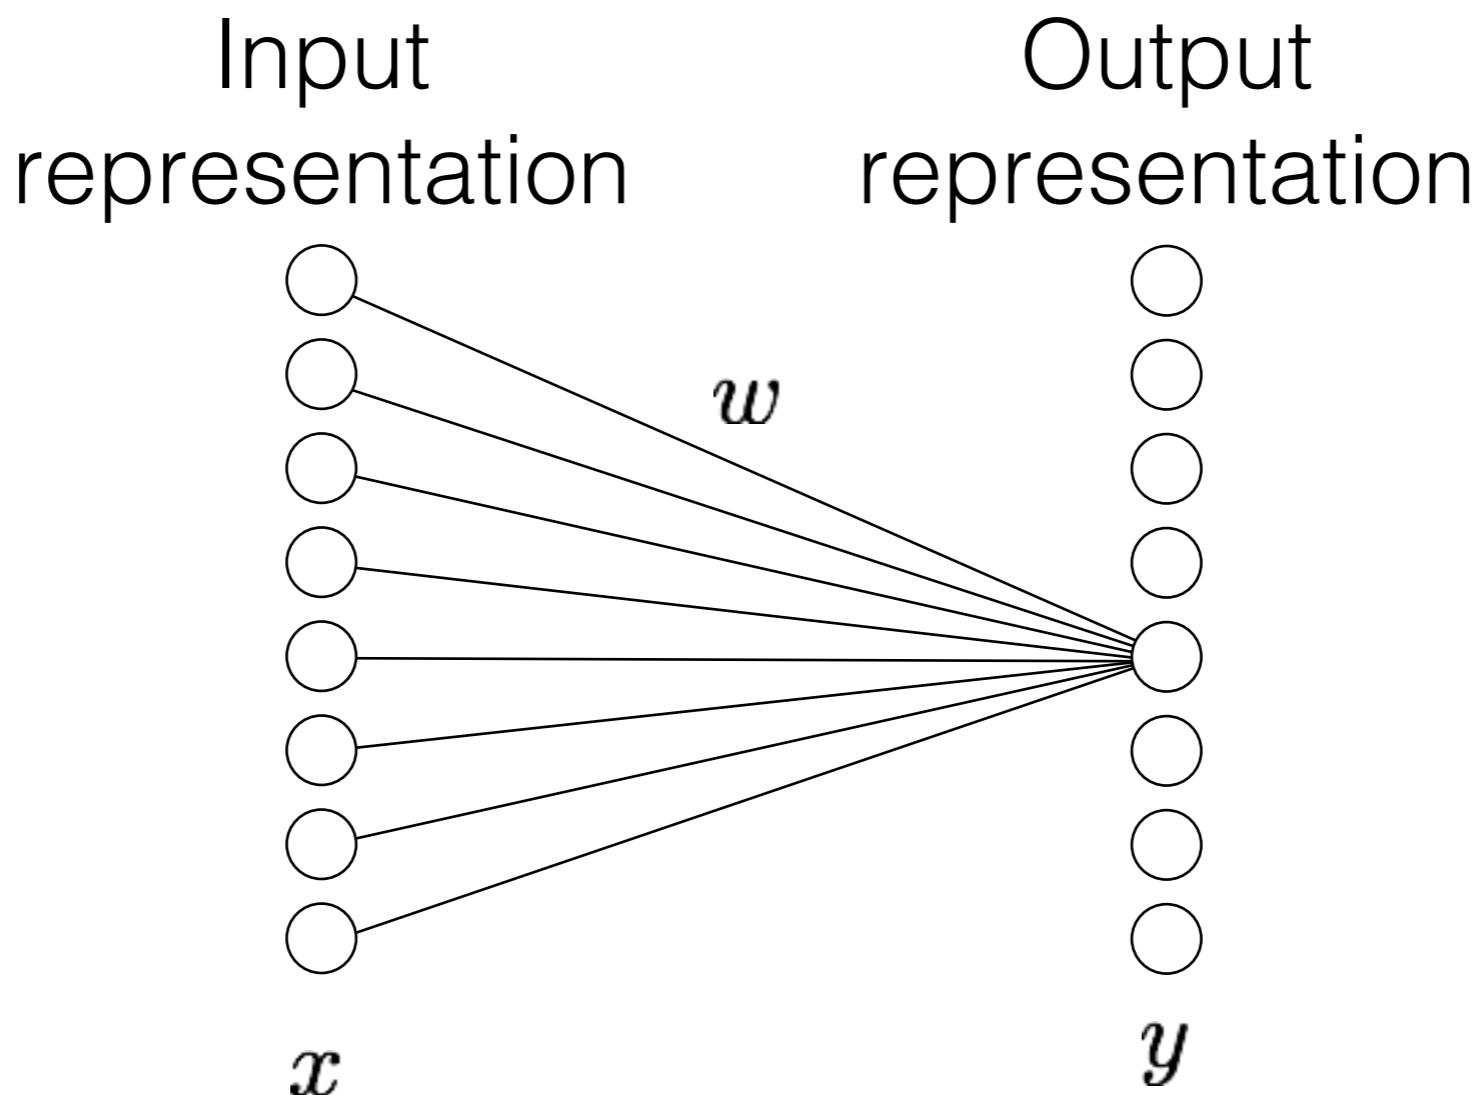
Deep neural network



Computation in a neural net



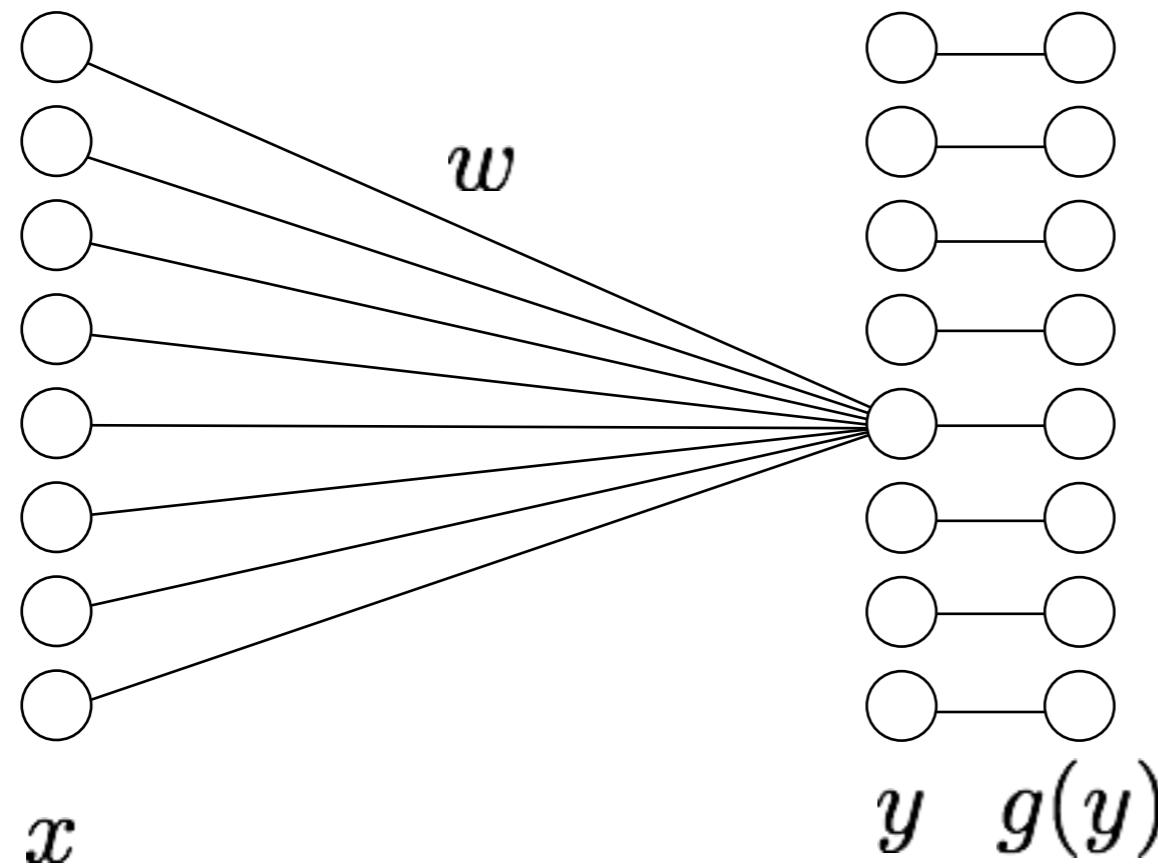
Computation in a neural net



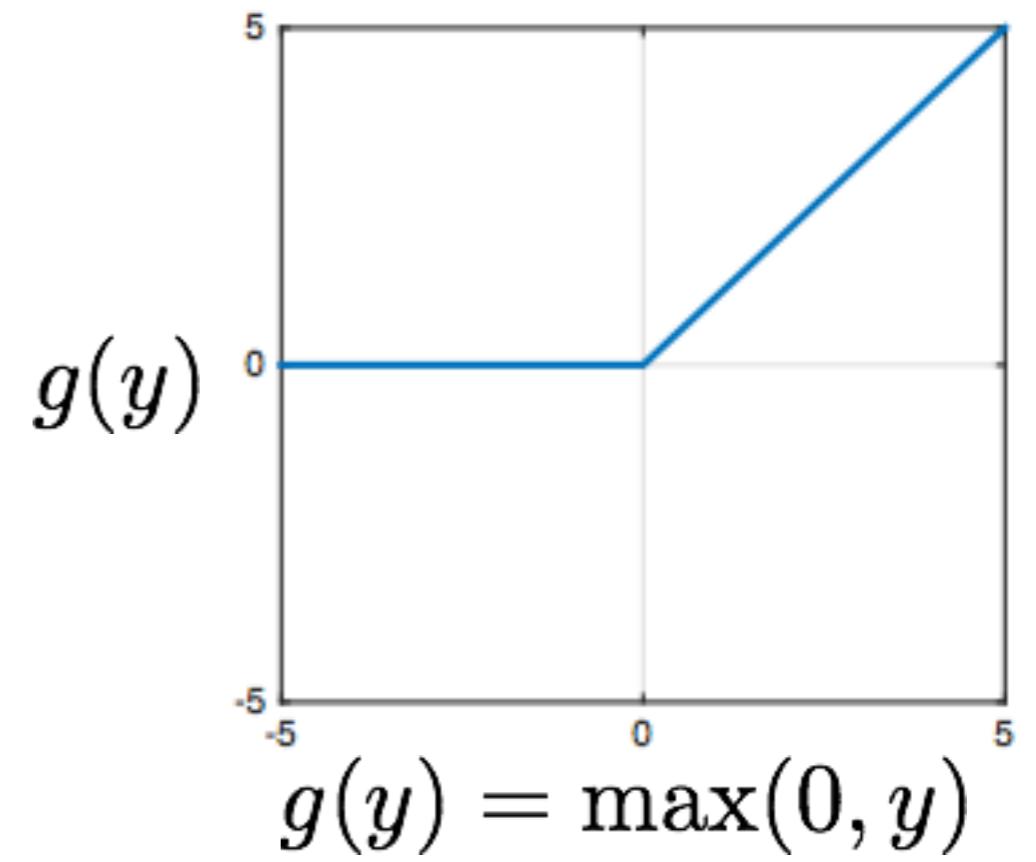
$$y_j = \sum_i w_{ij} x_i$$

i: the i^{th} dimension of x , j; the j^{th} dimension of y

Computation in a neural net

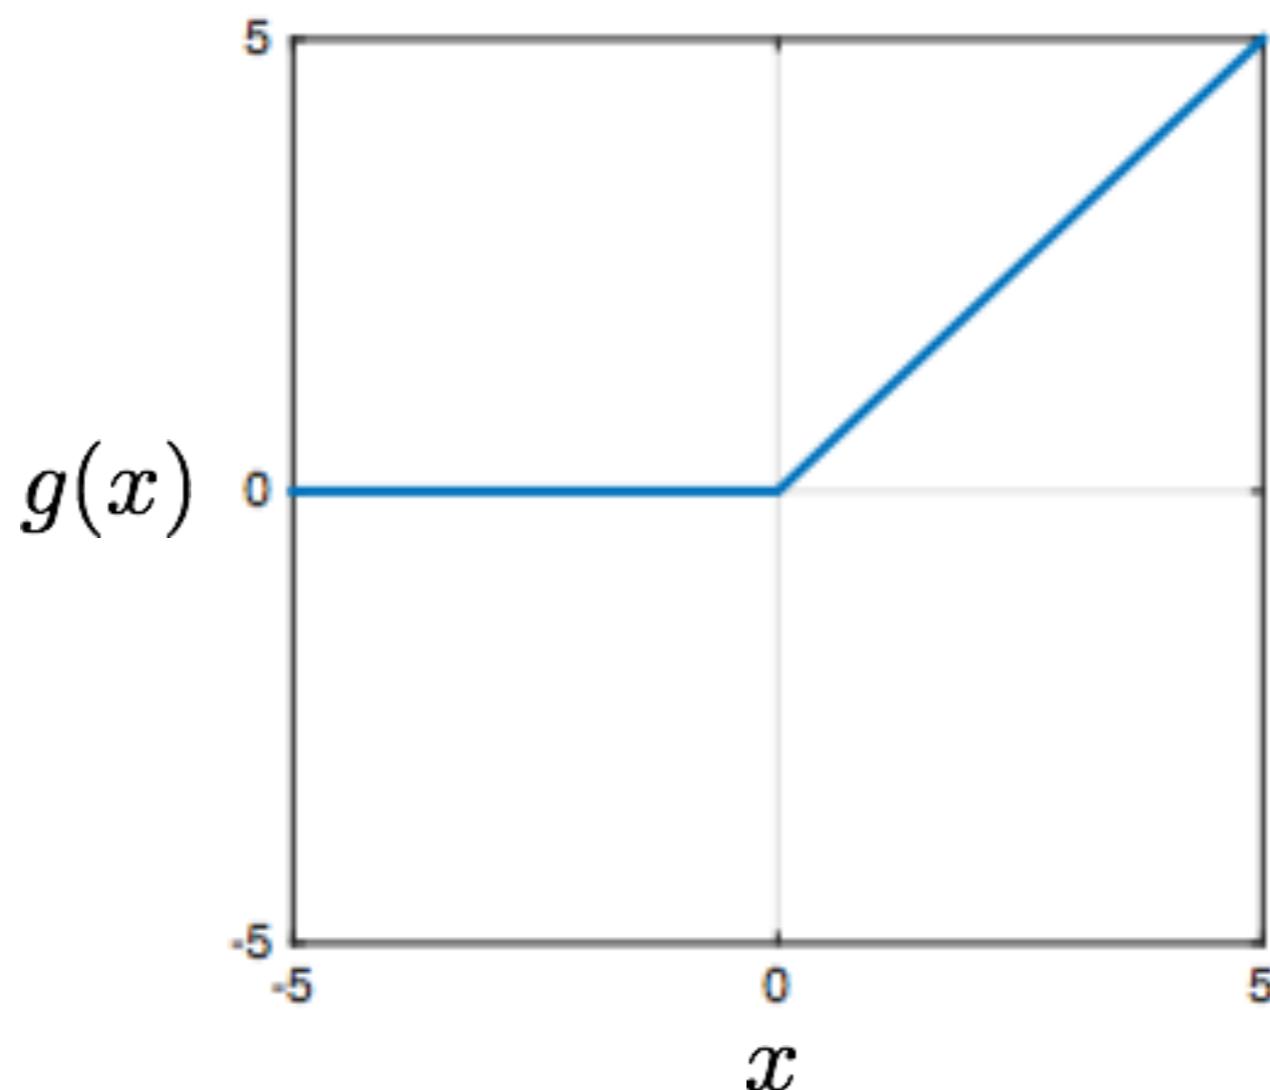


Rectified linear unit (ReLU)



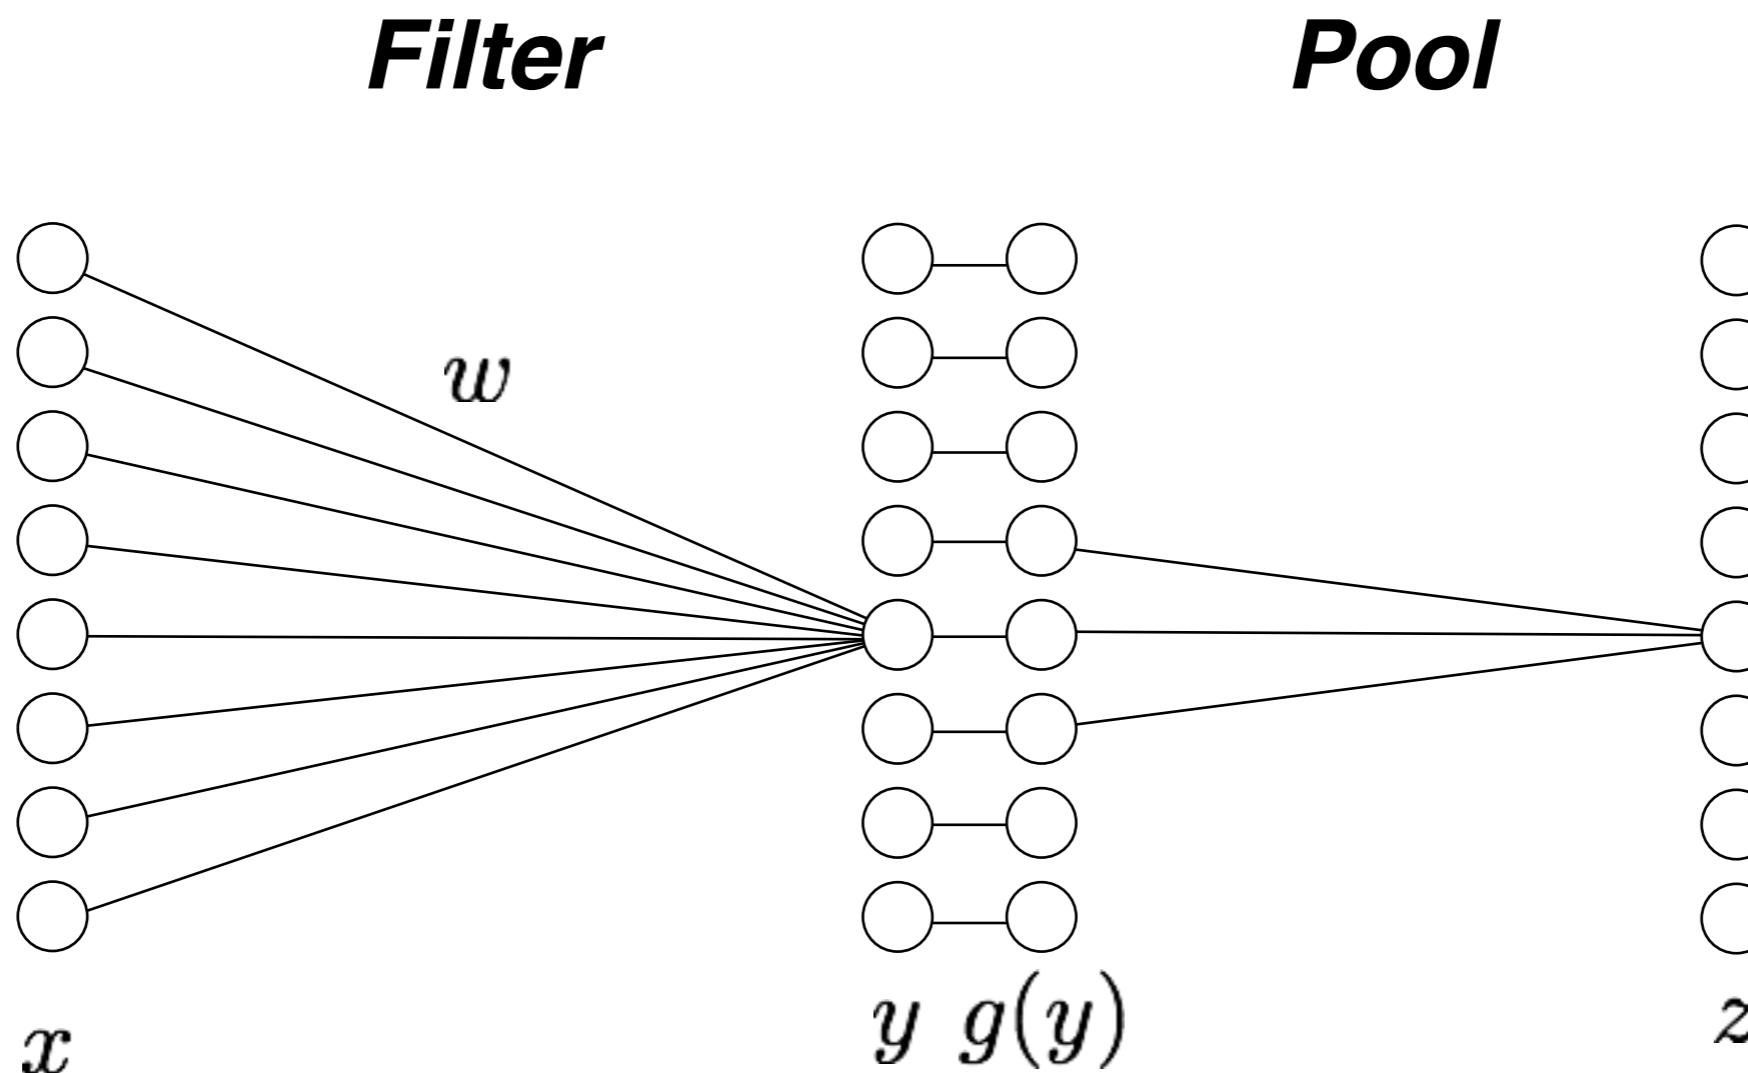
Computation in a neural net

Rectified linear unit (ReLU)



$$g(x) = \max(0, x)$$

Computation in a neural net



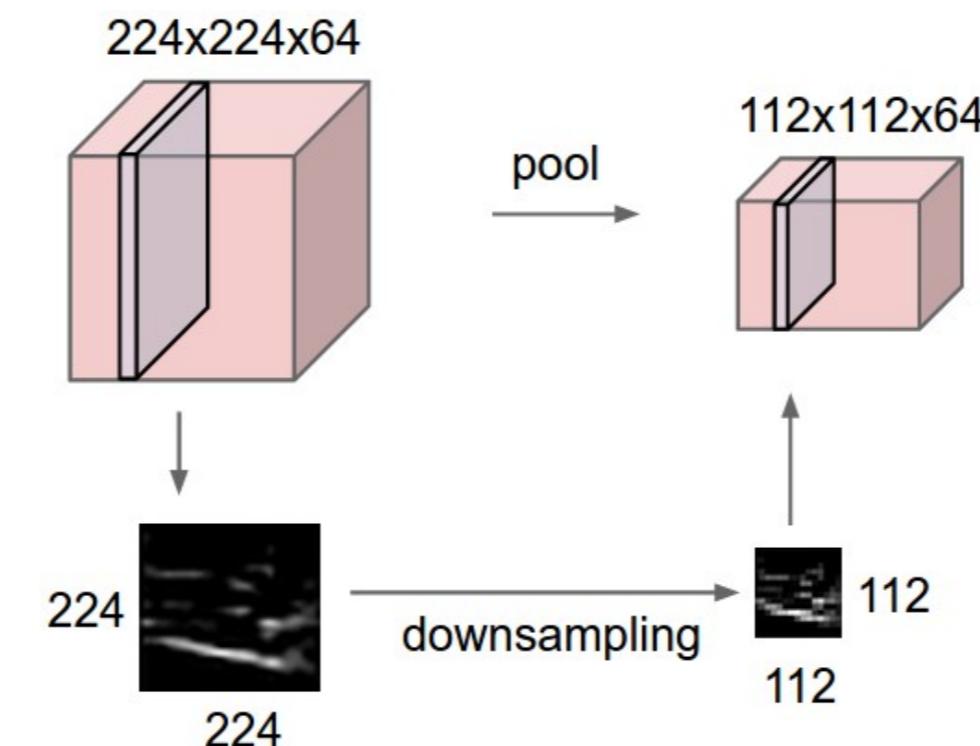
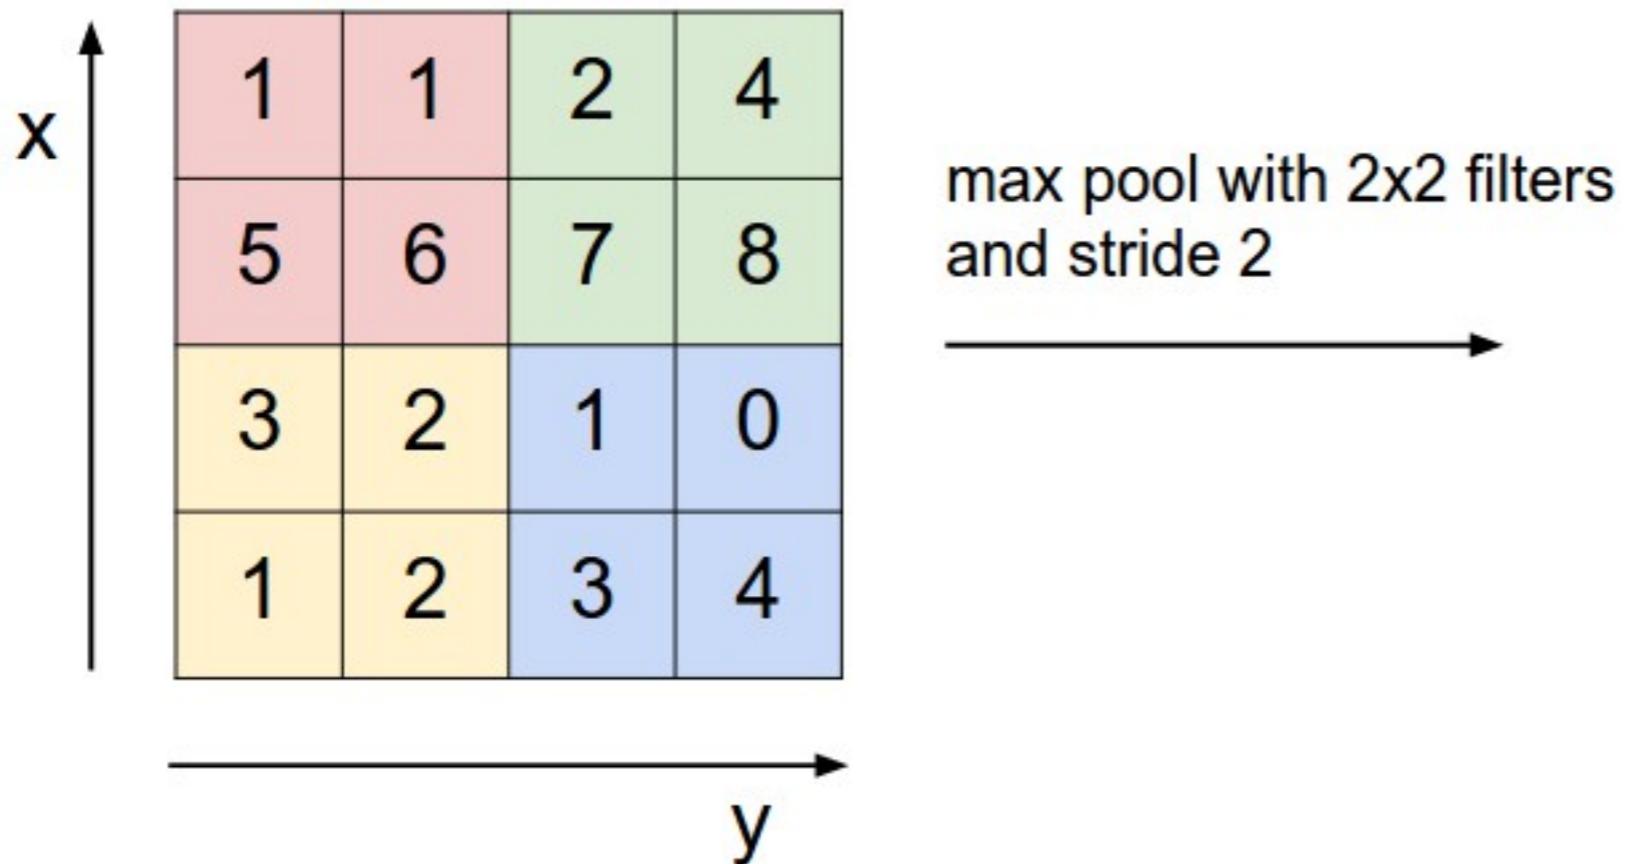
$$y_j = \sum_i w_{ij} x_i$$

$$z_k = \max_{j \in \mathcal{N}(k)} g(y_j)$$

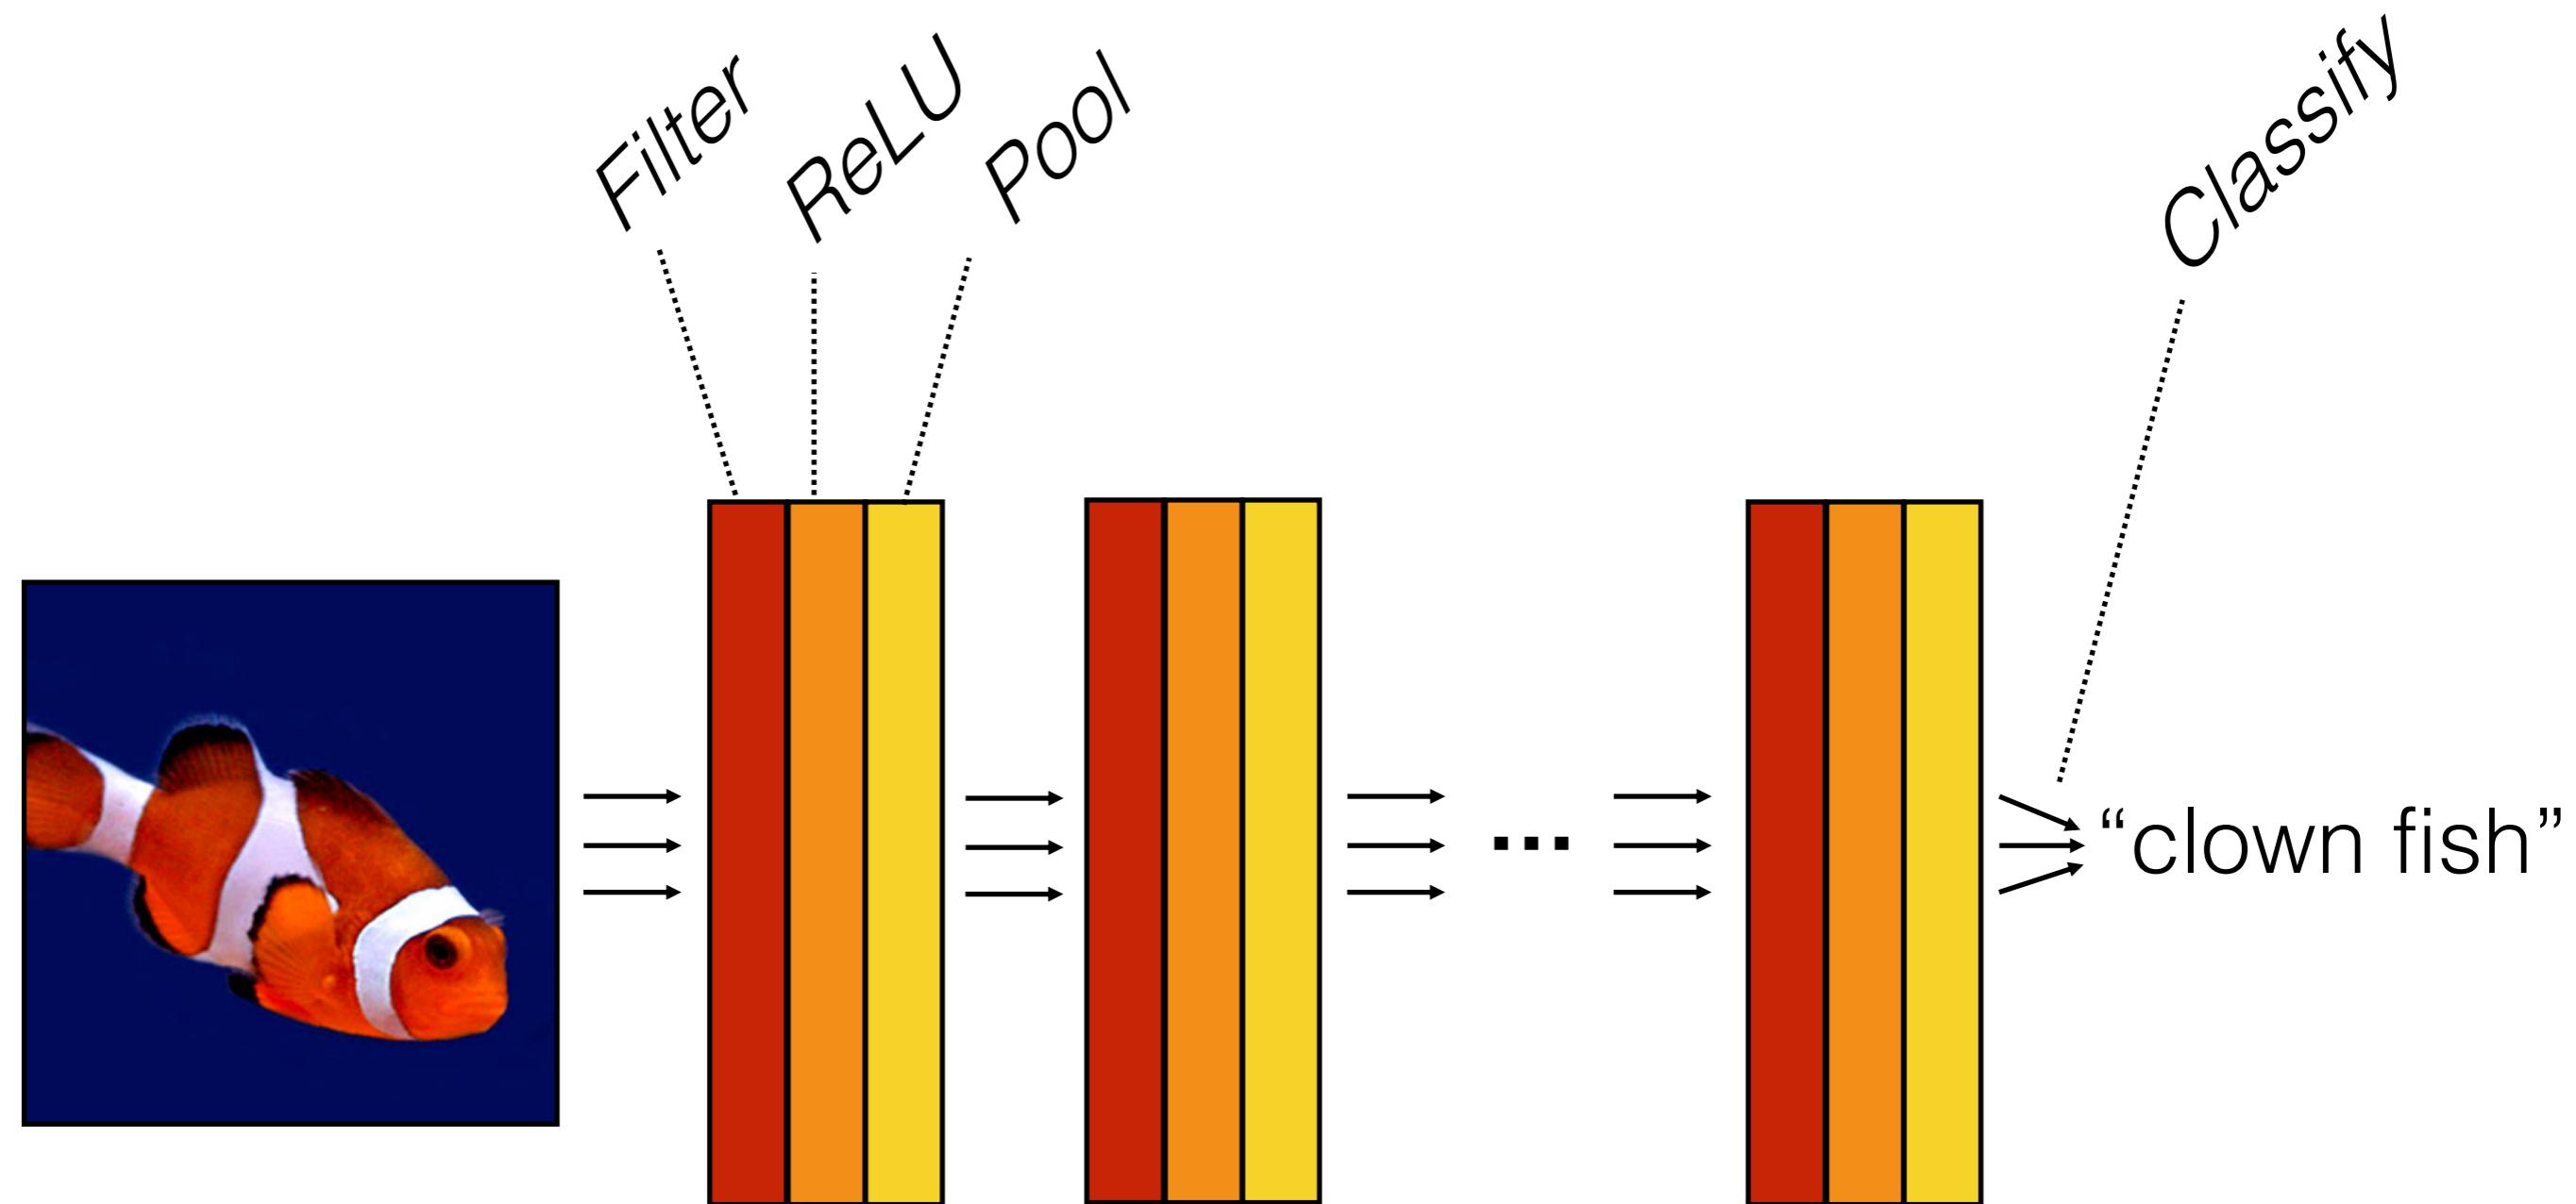
i: the i^{th} dimension of x , j; the j^{th} dimension of y

Computation in a neural net

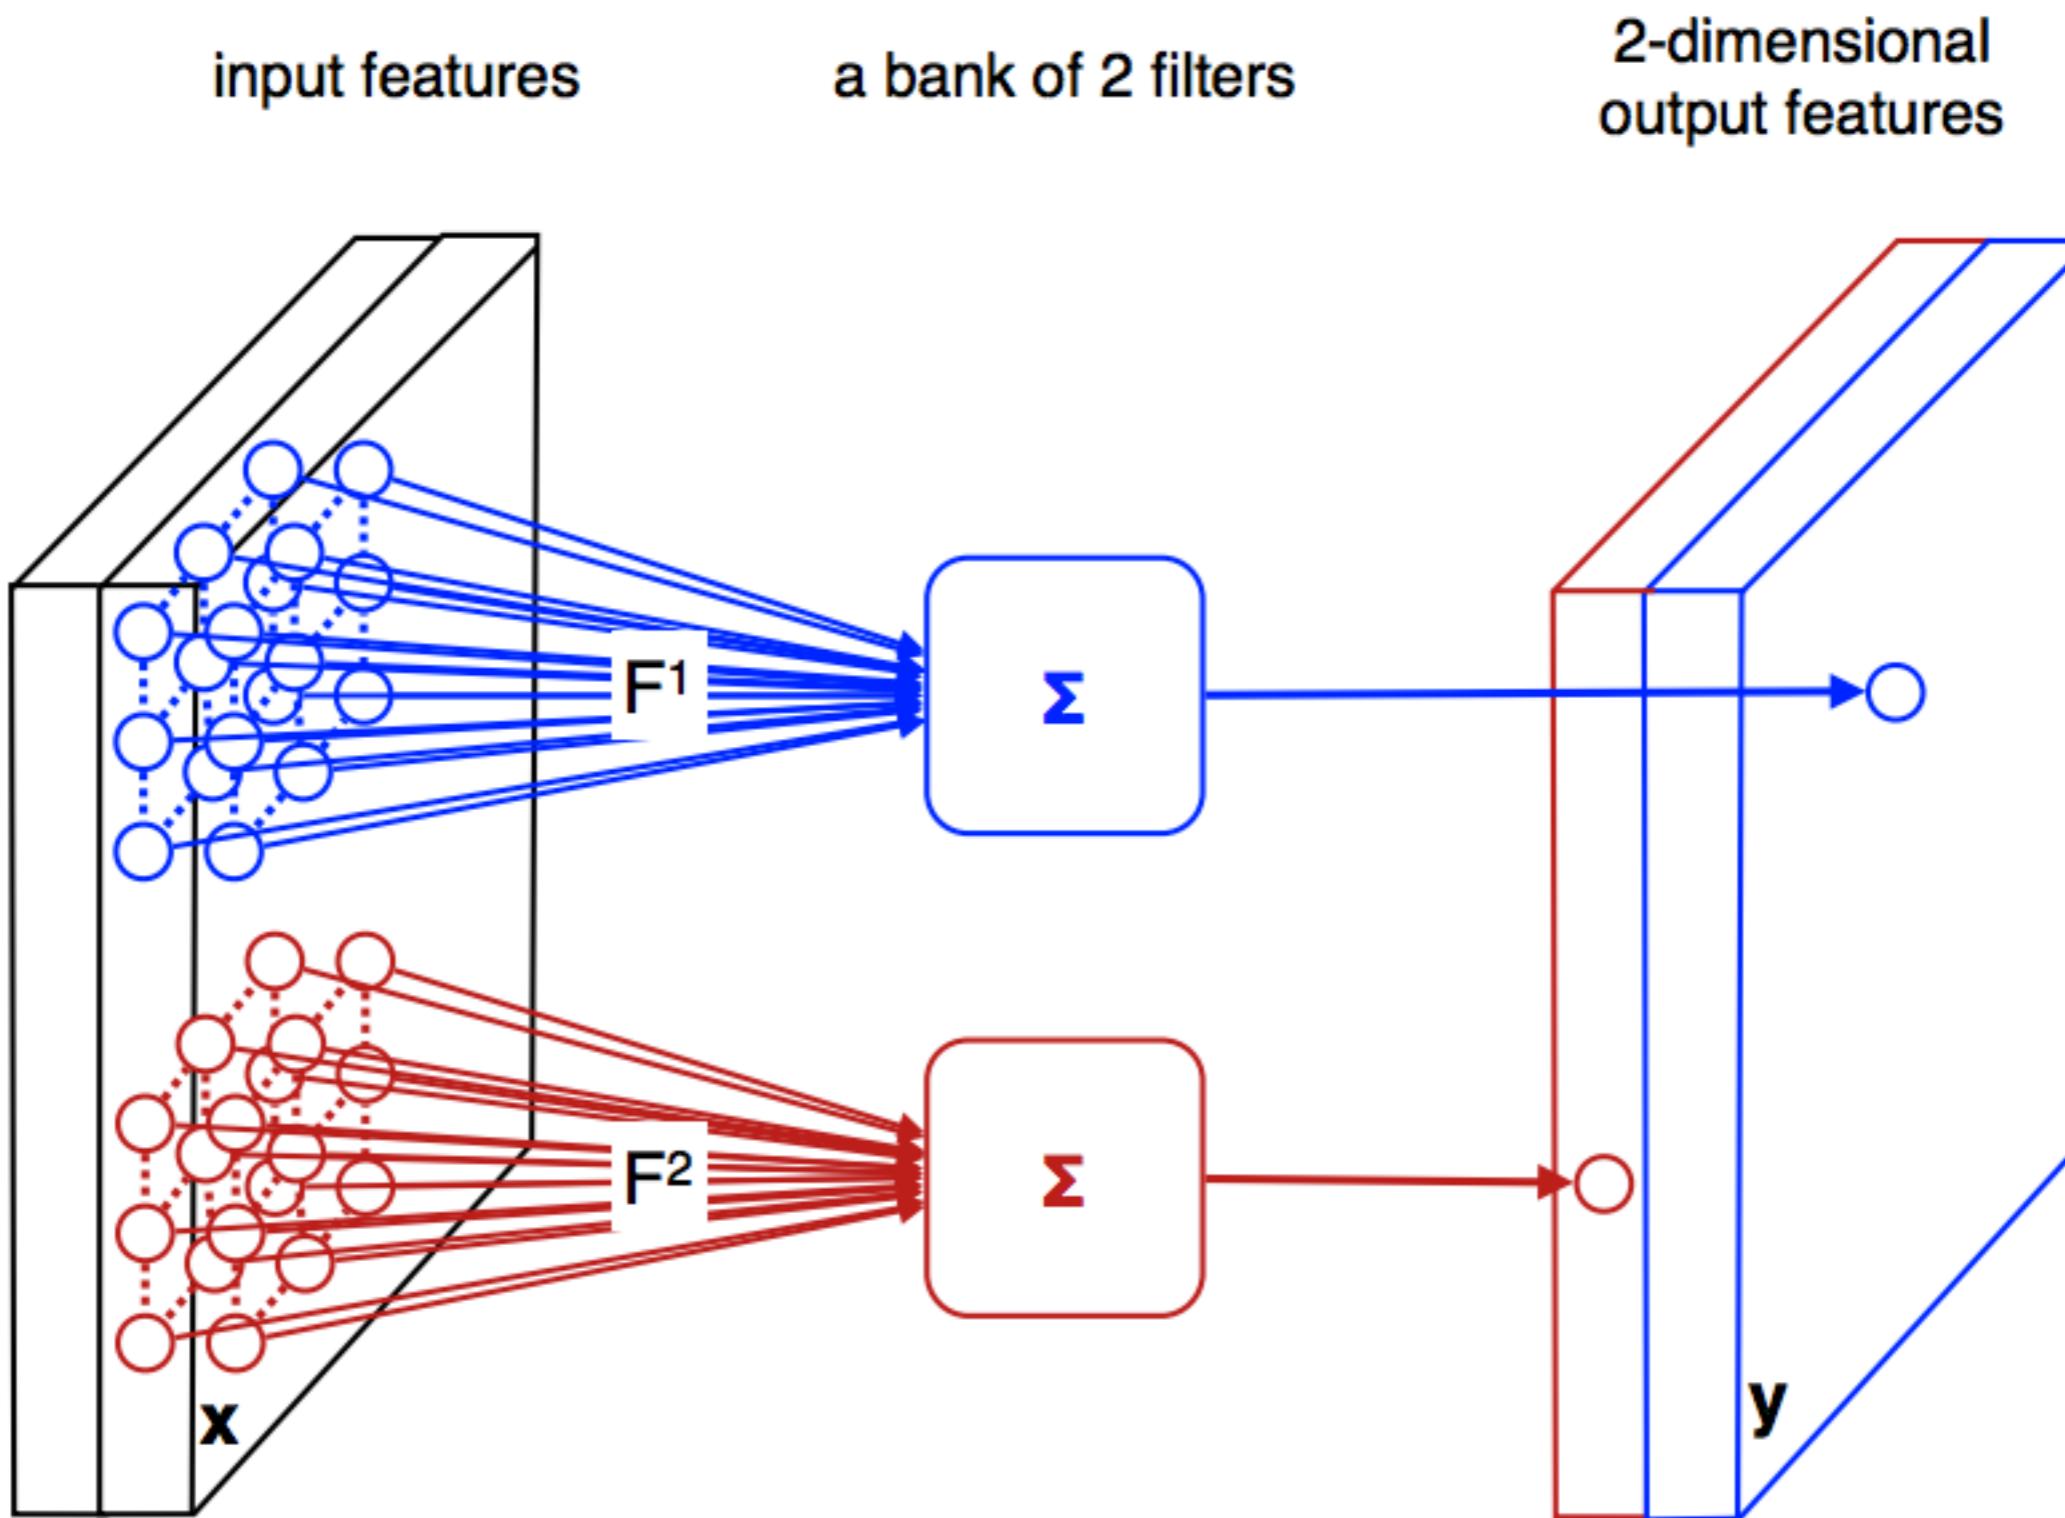
Single depth slice



Computation in a neural net

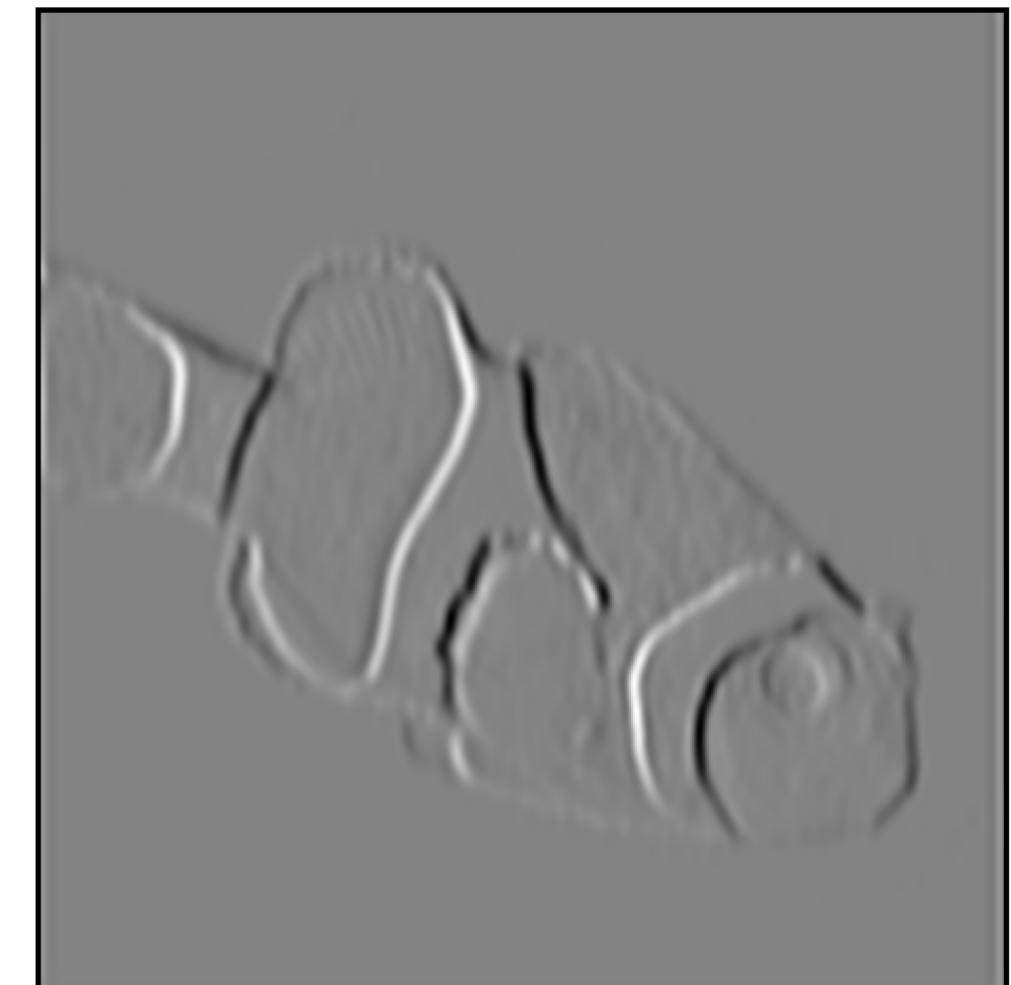
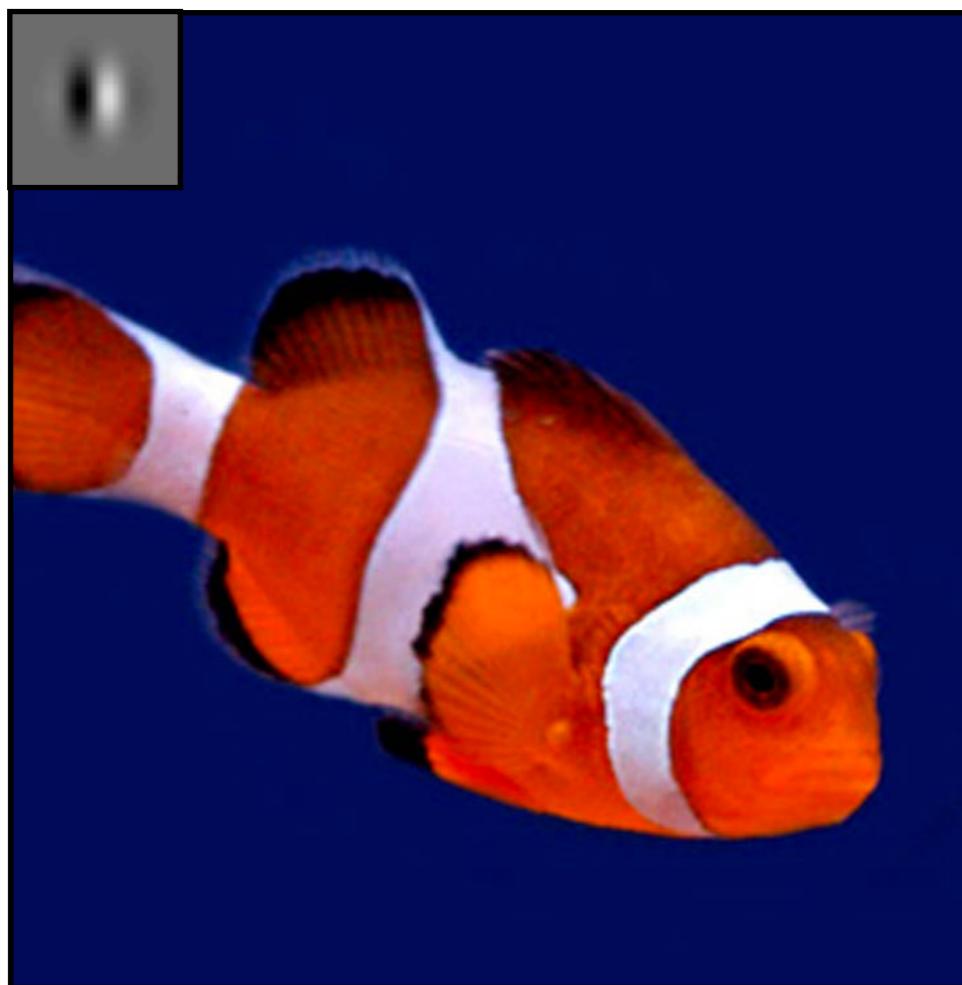


Convolutional Neural Nets

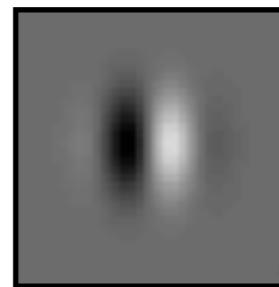


Convolutional Neural Nets

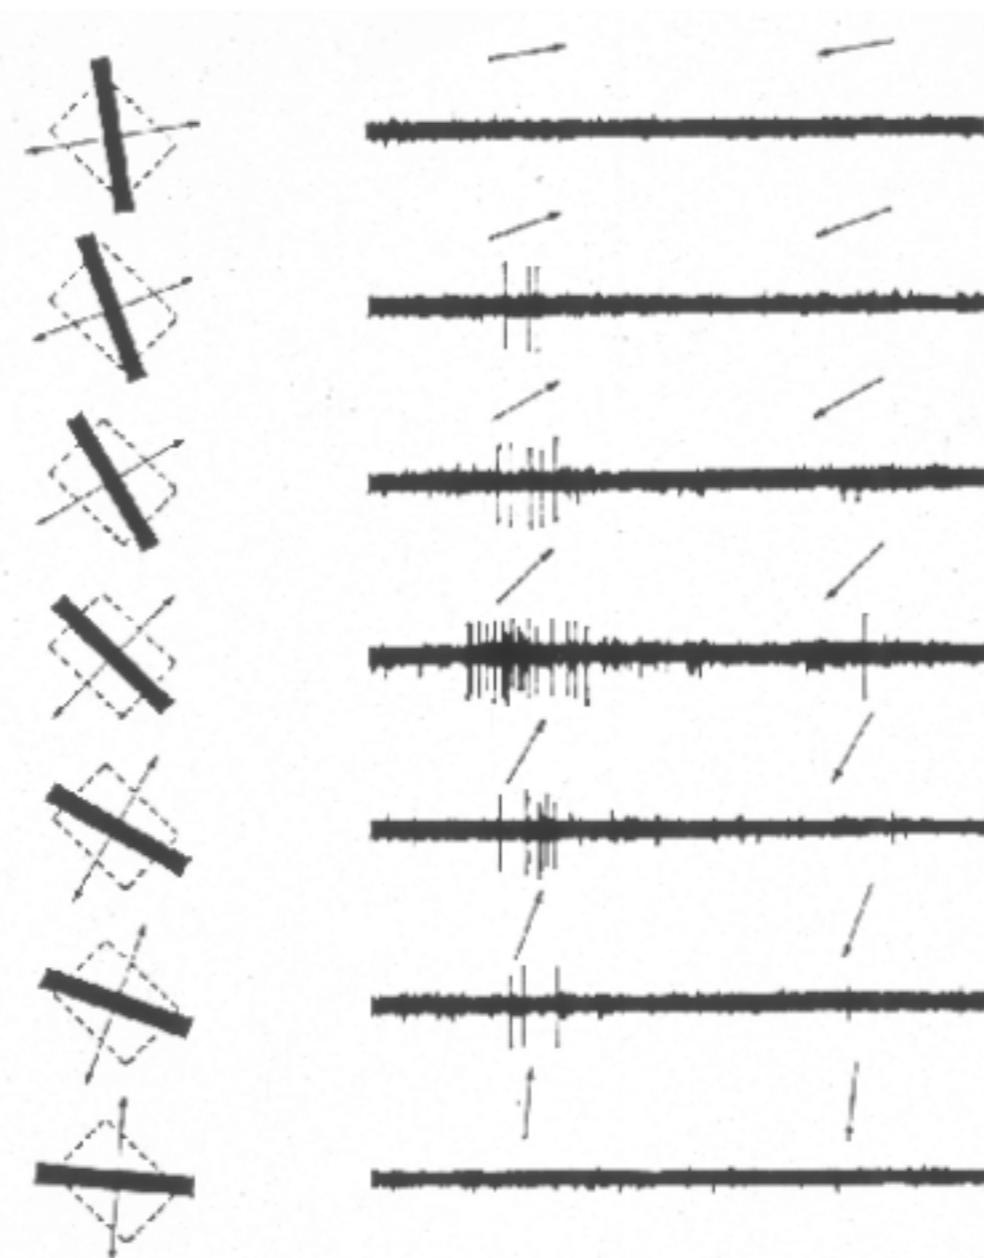
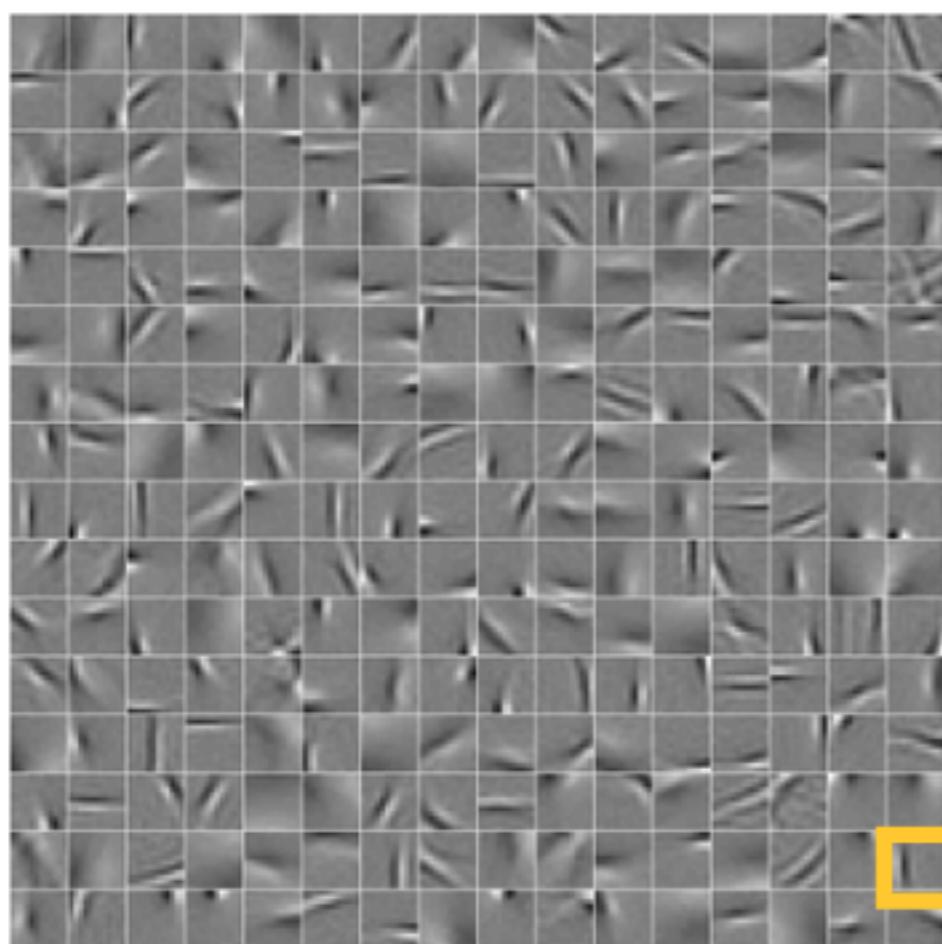
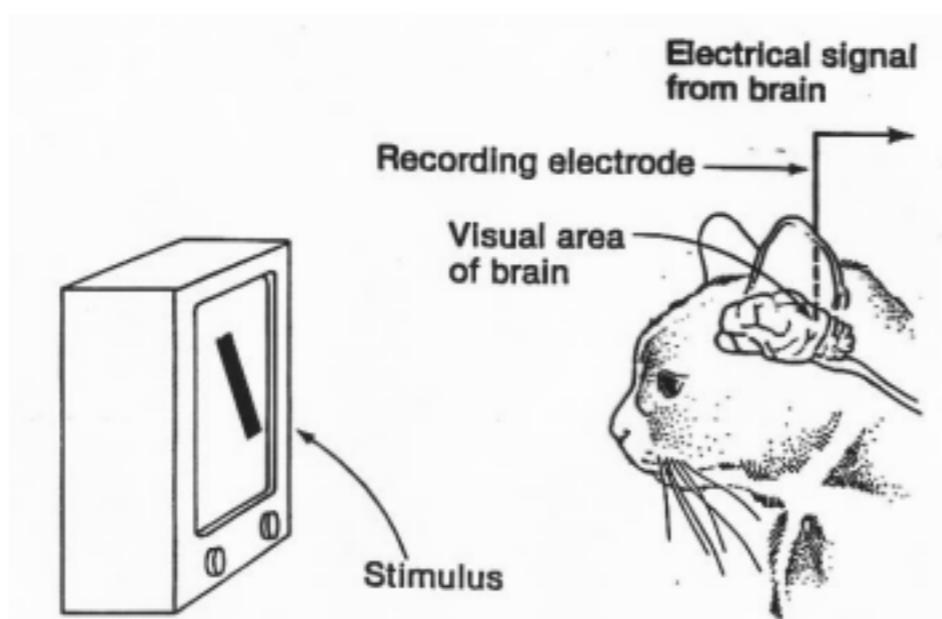
Convolution



filter



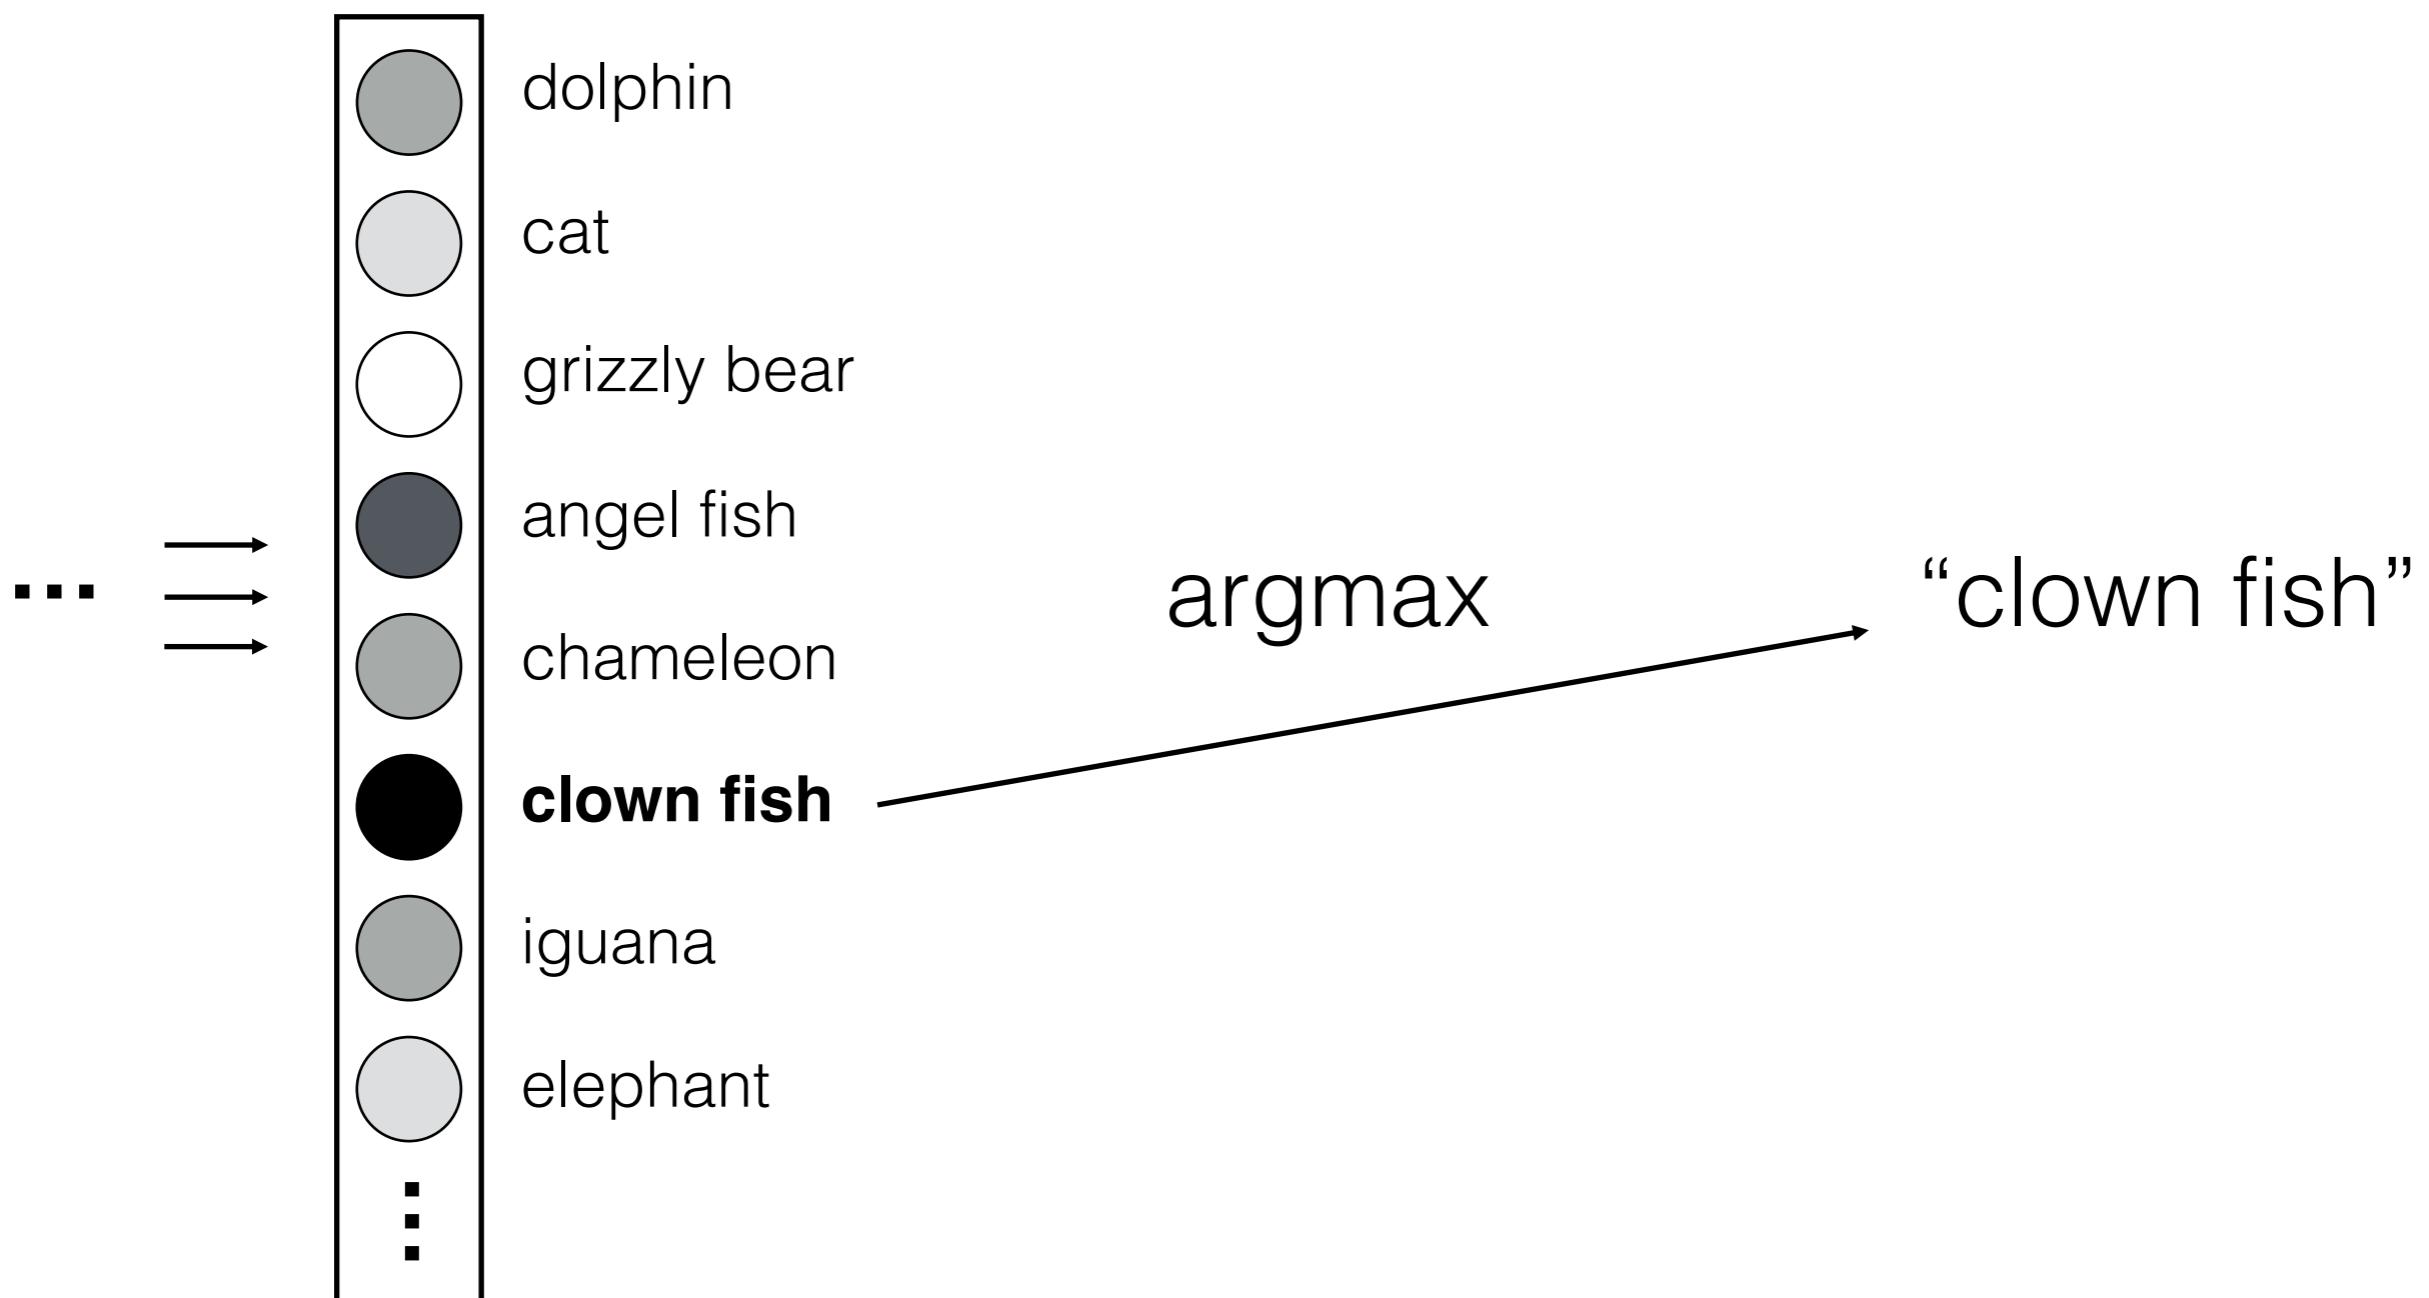
[Hubel and Wiesel 59]



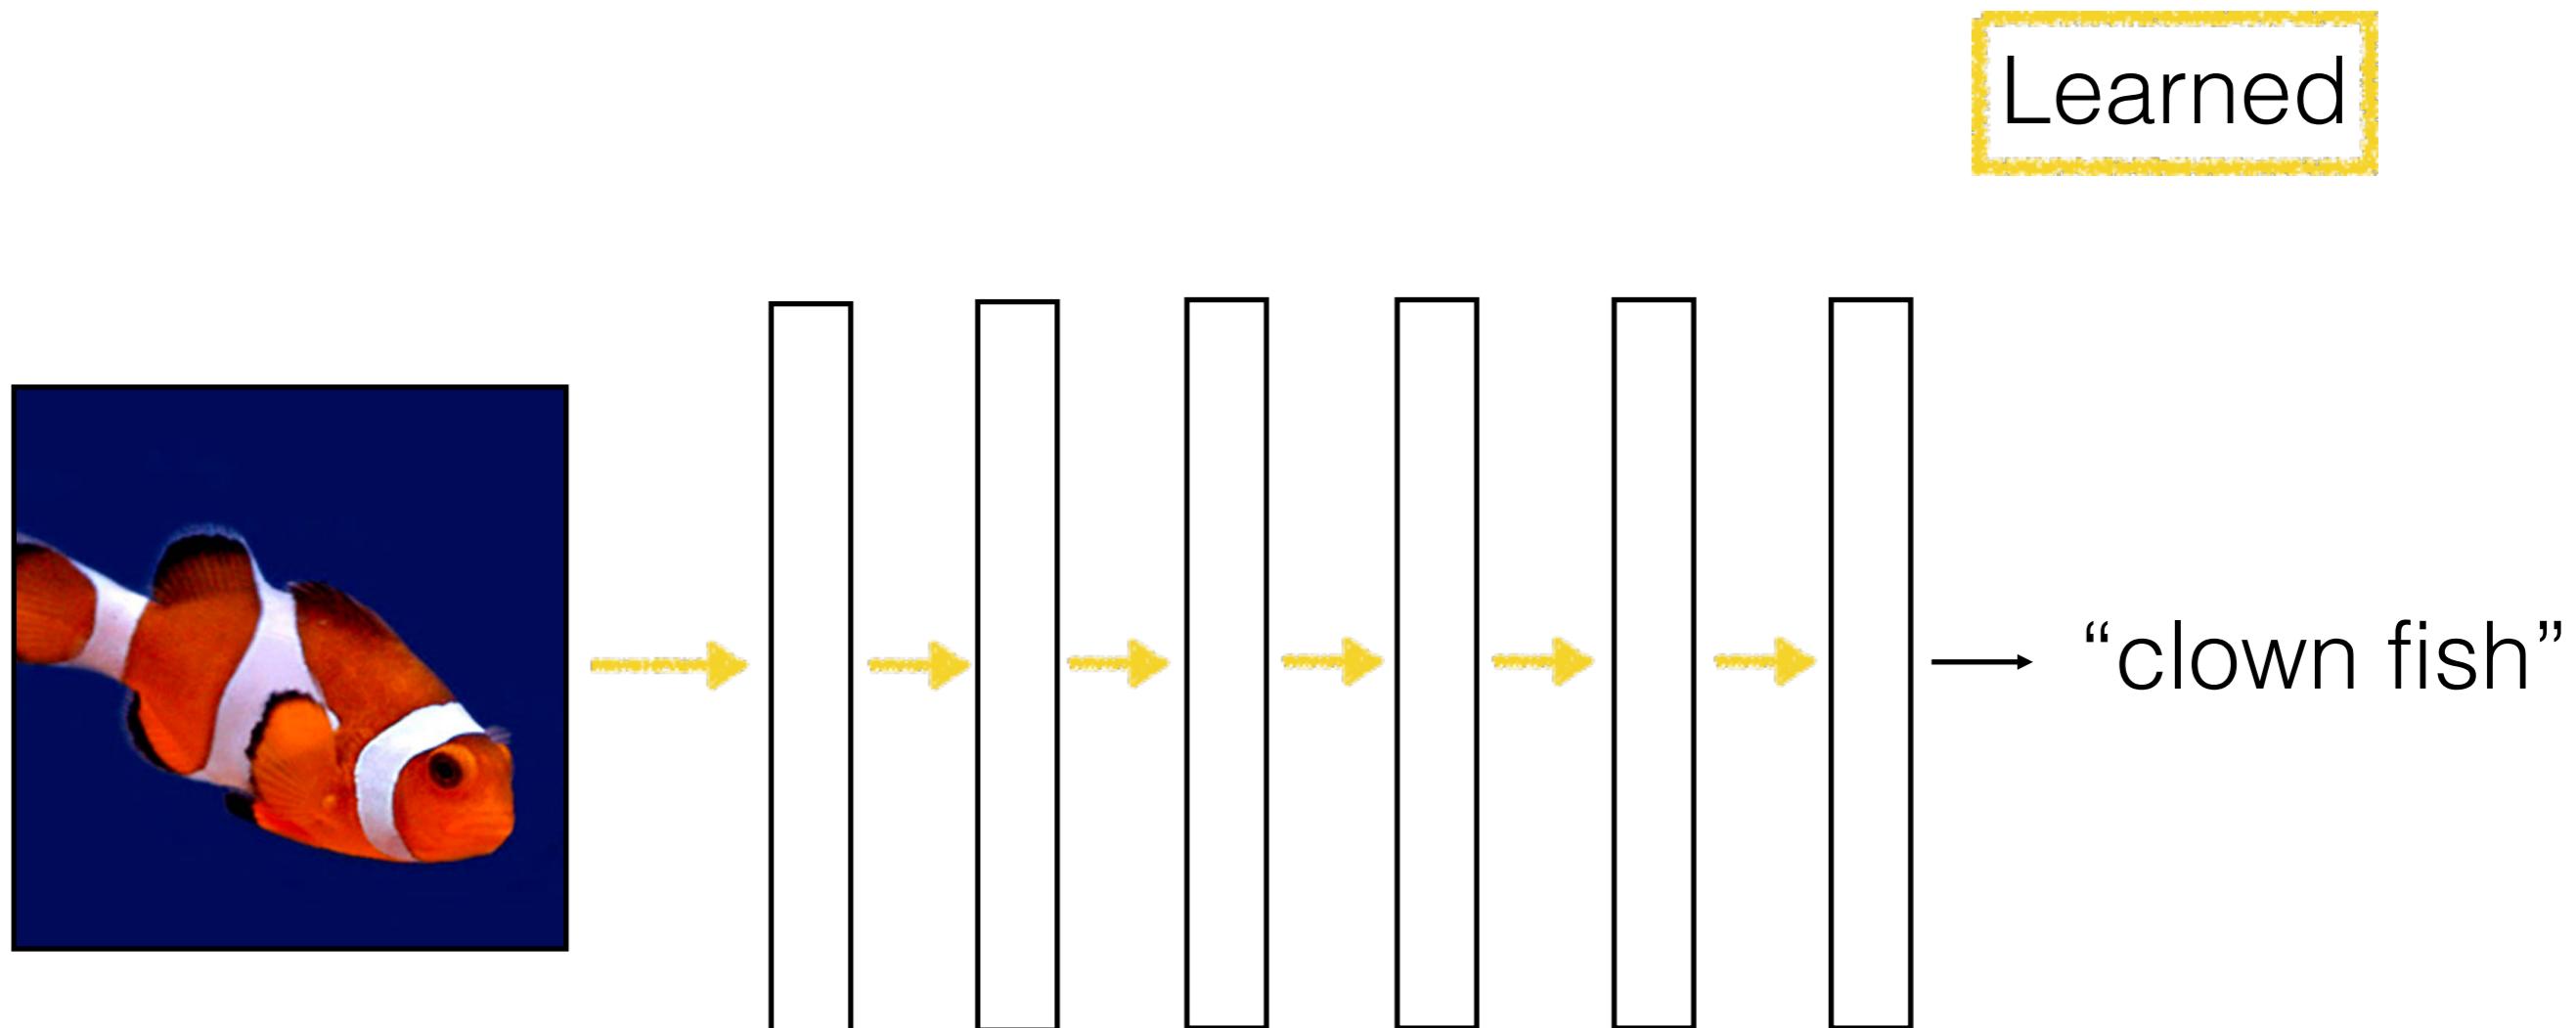
oriented filter

Computation in a neural net

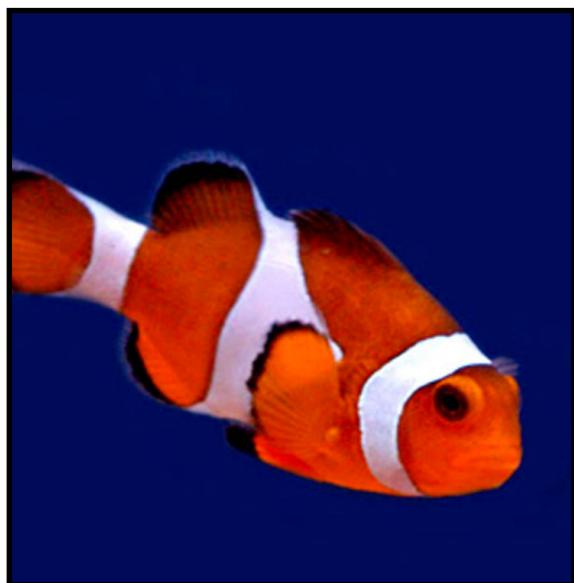
Last layer



Learning with deep nets



Learning with deep nets



→ “clown fish”



→ “grizzly bear”



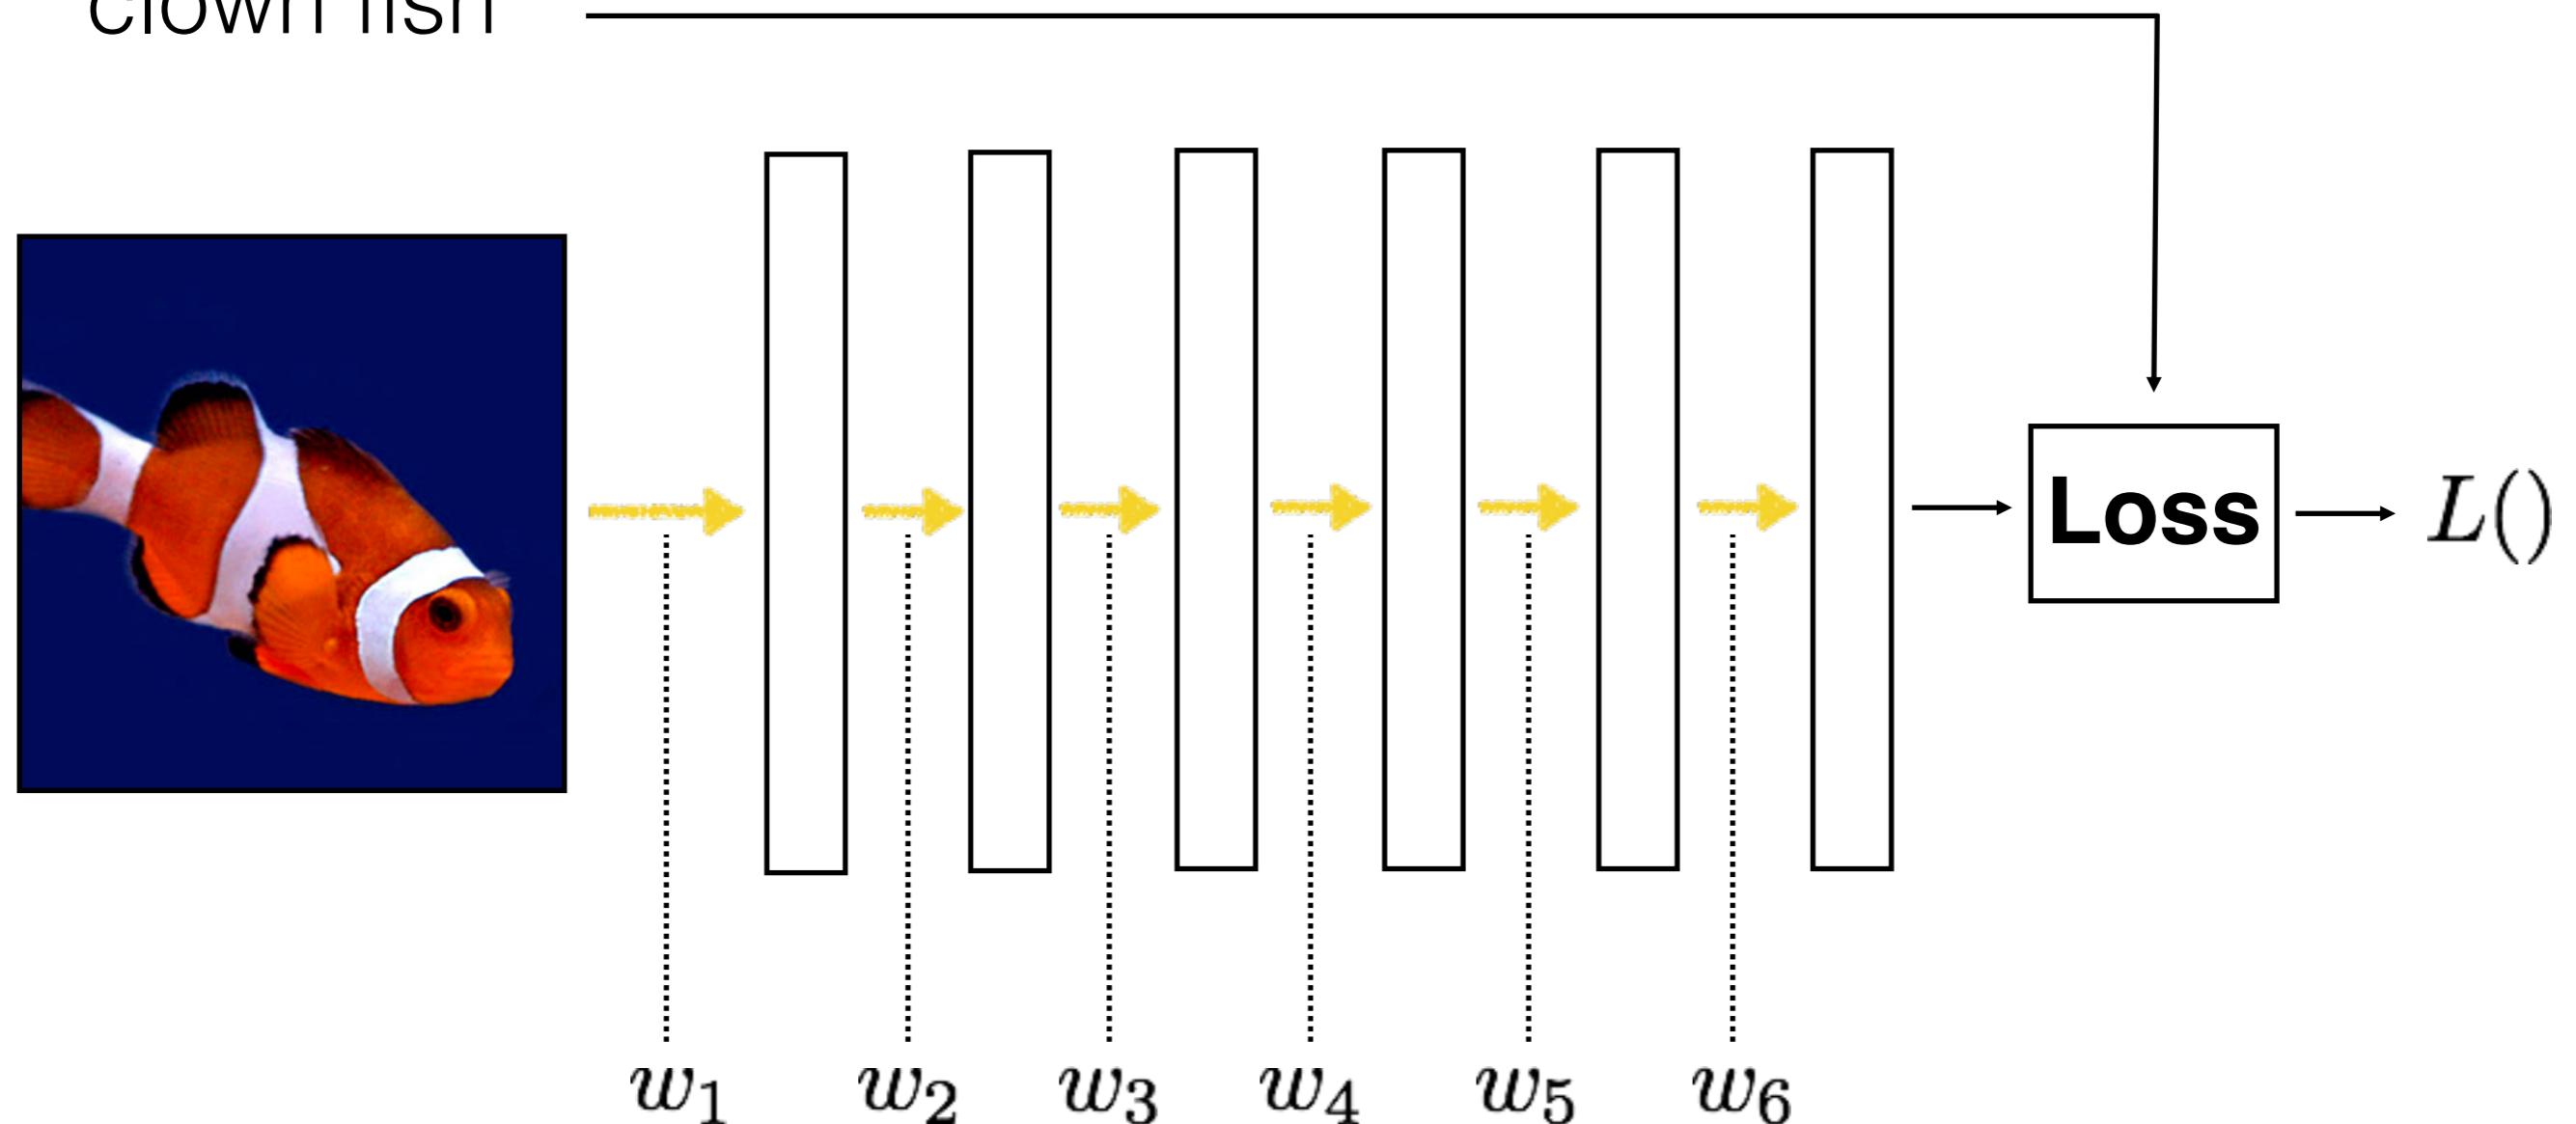
→ “chameleon”

Train network to
associate the right
label with each image

Learning with deep nets

Learned

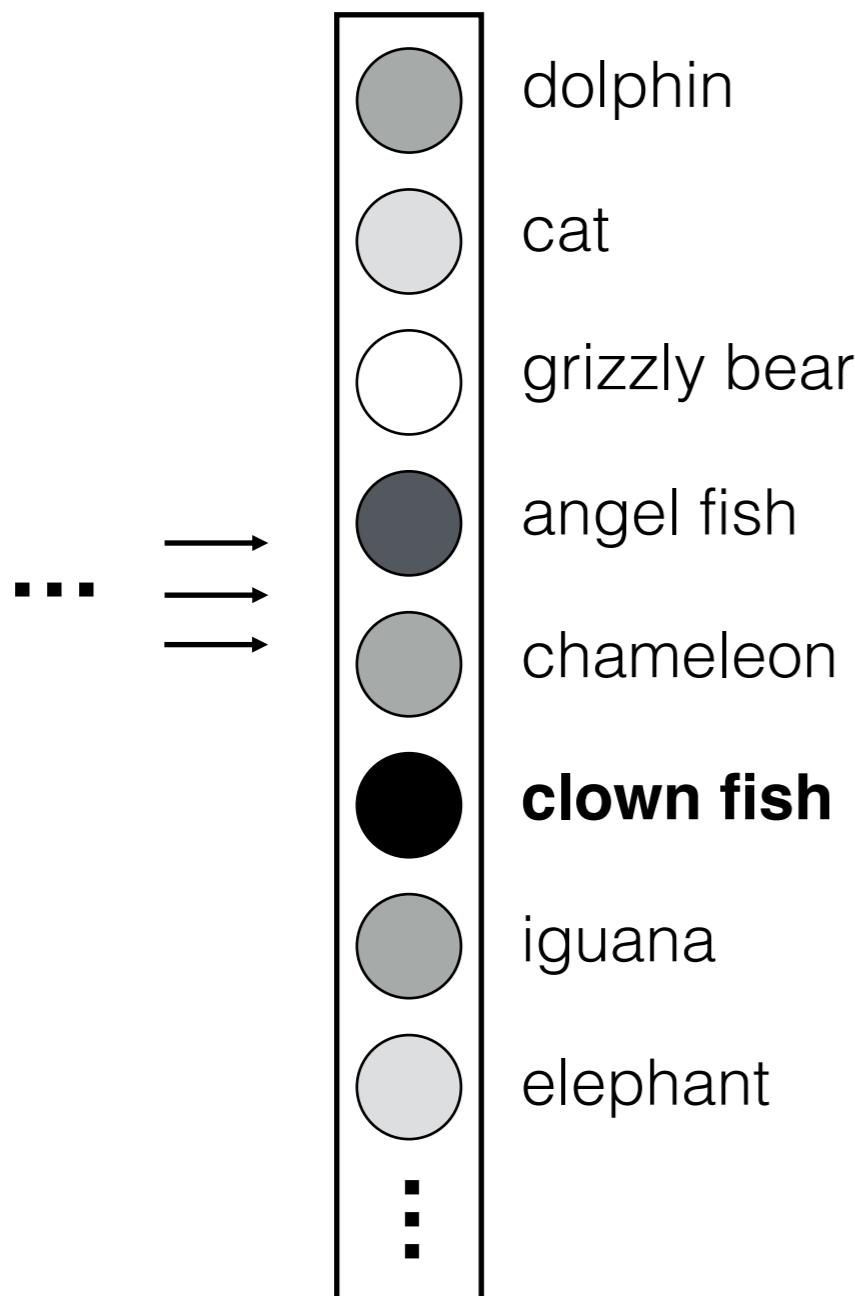
“clown fish”



$$\underset{\mathbf{w}}{\operatorname{argmin}} \quad L(w_1, \dots, w_6)$$

Loss function

Network output



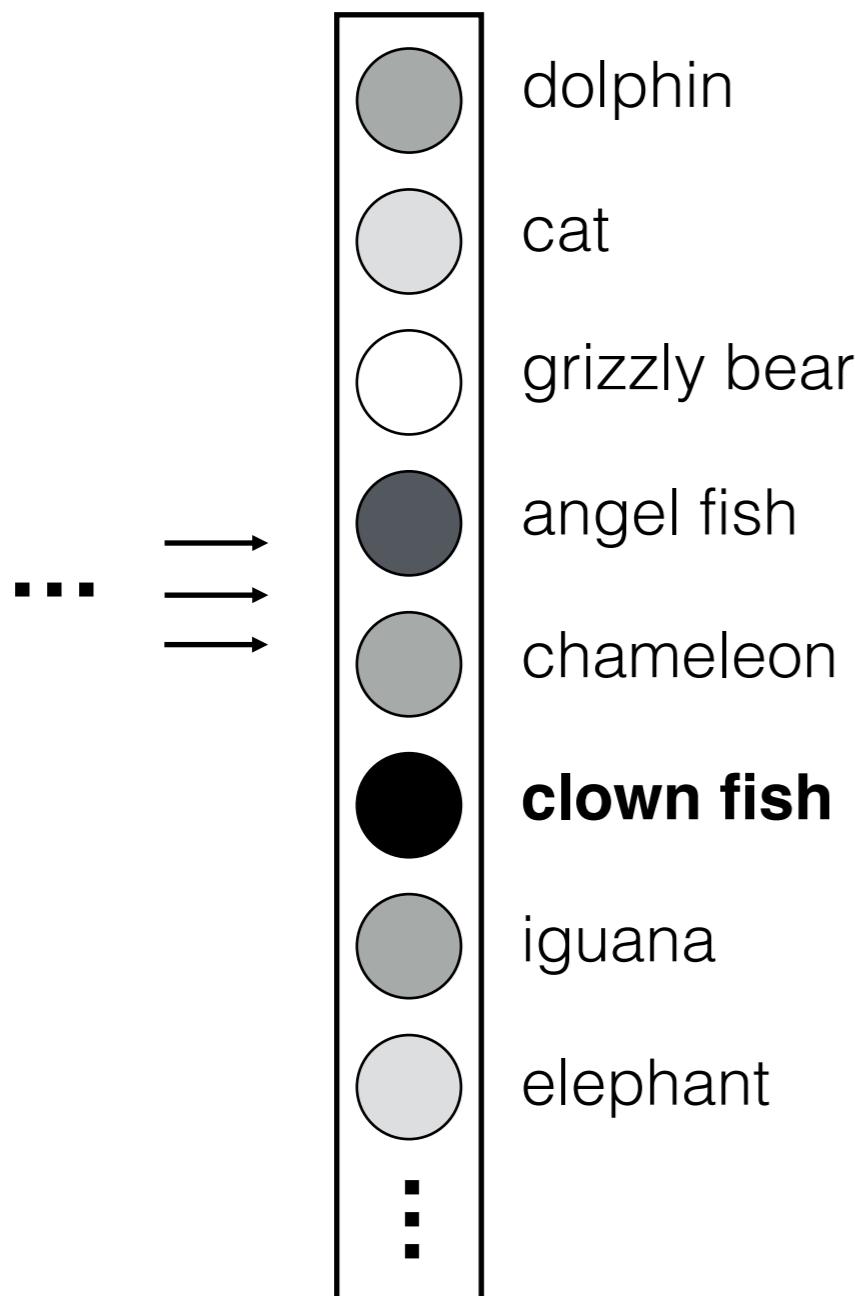
Ground truth label

“clown fish”

Loss → error

Loss function

Network output



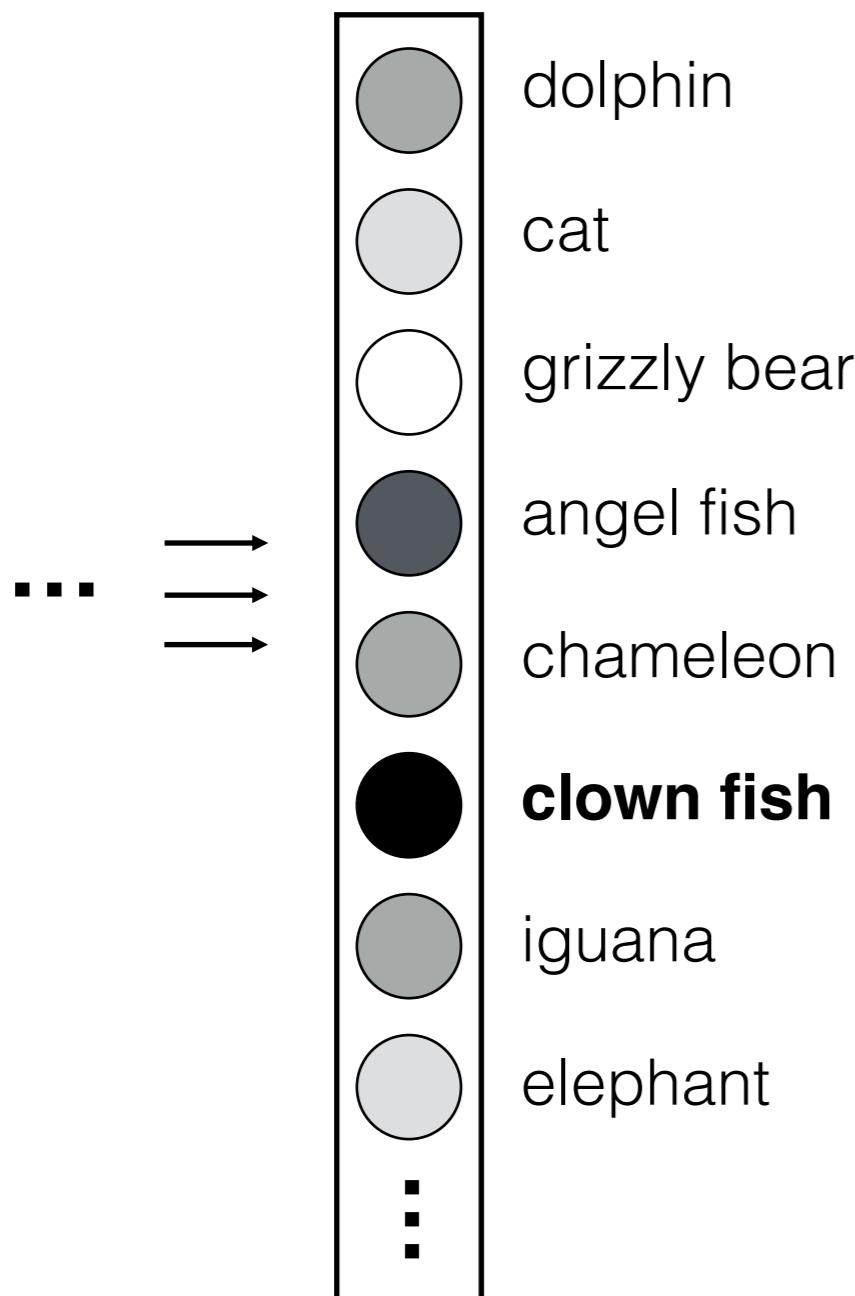
Ground truth label

“clown fish”

Loss → **small**

Loss function

Network output



Ground truth label

“grizzly bear”

Loss → **large**

Loss function for classification

Network output

Ground truth label



Probability of the observed data under the model

$$H(\hat{z}, z) = - \sum_c z_c \log \hat{z}_c$$

Cross-entropy loss

c is the c^{th} class in the output

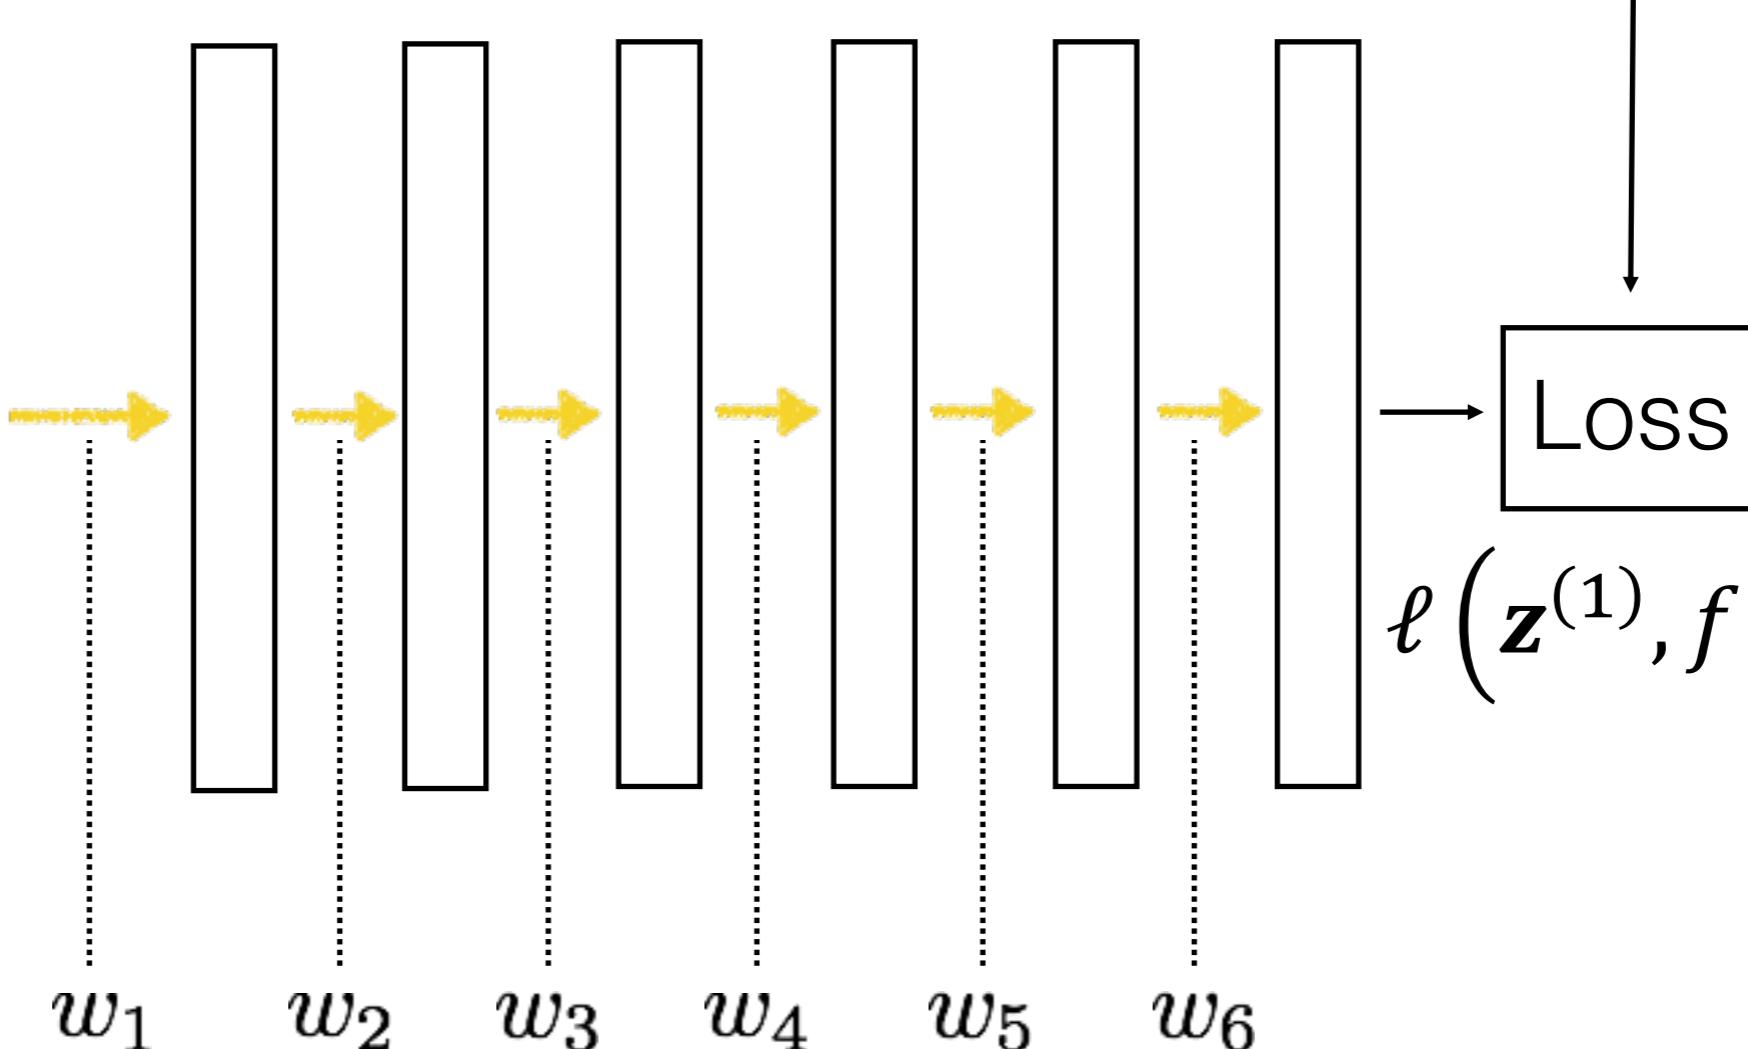
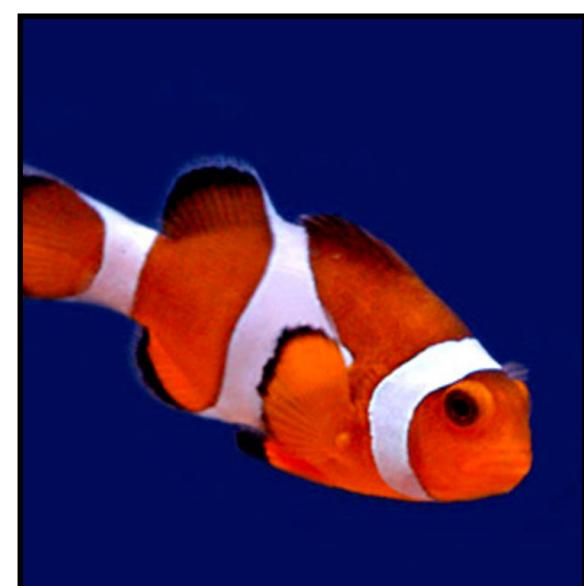
Learning with deep nets

$\mathbf{z}^{(1)}$

“clown fish”

Learned

$\mathbf{x}^{(1)}$



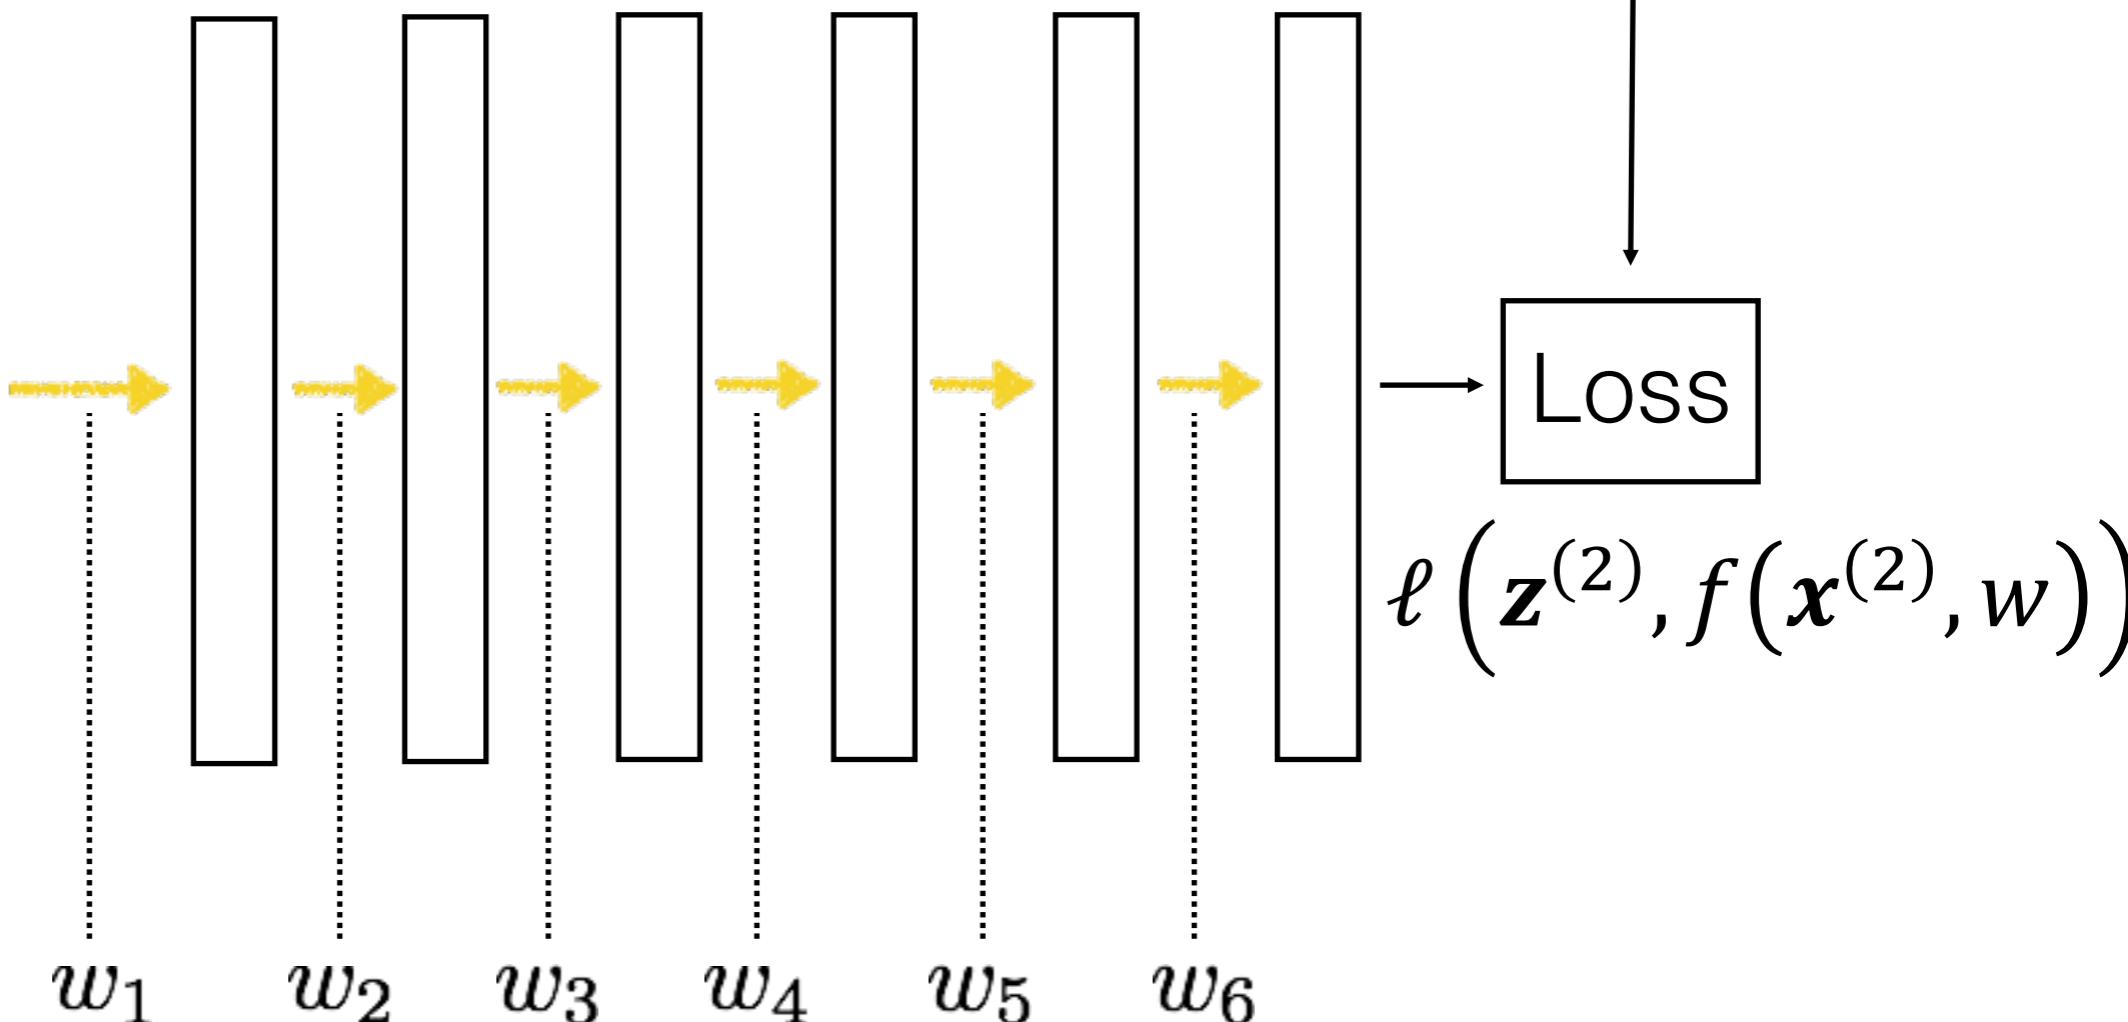
$\mathbf{x}^{(1)}, \mathbf{z}^{(1)}$ is the input and label
of the 1st training image

Learning with deep nets

$\mathbf{z}^{(2)}$
“grizzly bear” —

Learned

$\mathbf{x}^{(2)}$



$\mathbf{x}^{(2)}, \mathbf{z}^{(2)}$ is the input and label
of the 2nd training image

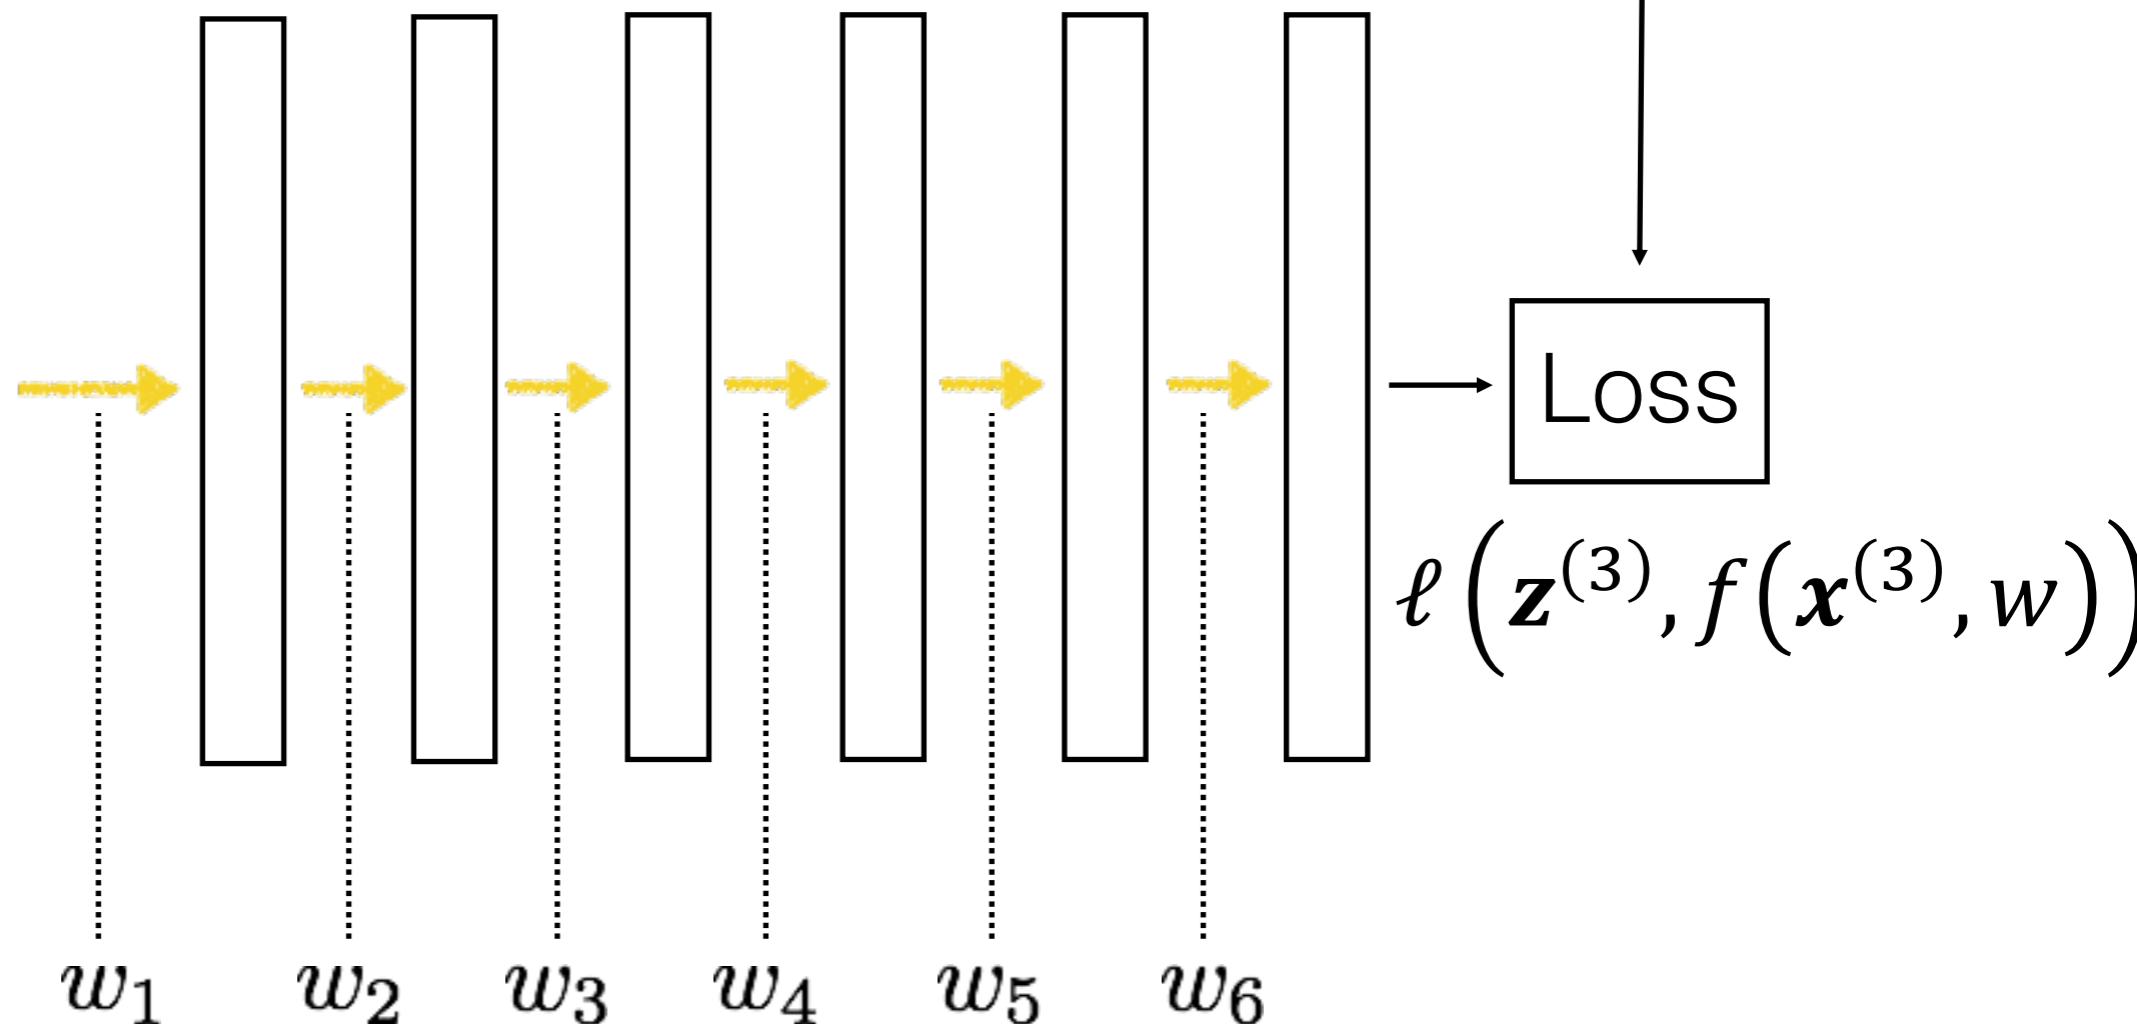
Learning with deep nets

$\mathbf{z}^{(3)}$

“chameleon” —

Learned

$\mathbf{x}^{(3)}$



$$\operatorname{argmin}_{\mathbf{w}} \sum_i \ell(\mathbf{z}^{(i)}, f(\mathbf{x}^{(i)}, \mathbf{w}))$$

Gradient descent

$$\operatorname{argmin}_w \sum_i \ell(z^{(i)}, f(x^{(i)}, w)) = \operatorname{argmin}_w L(w)$$

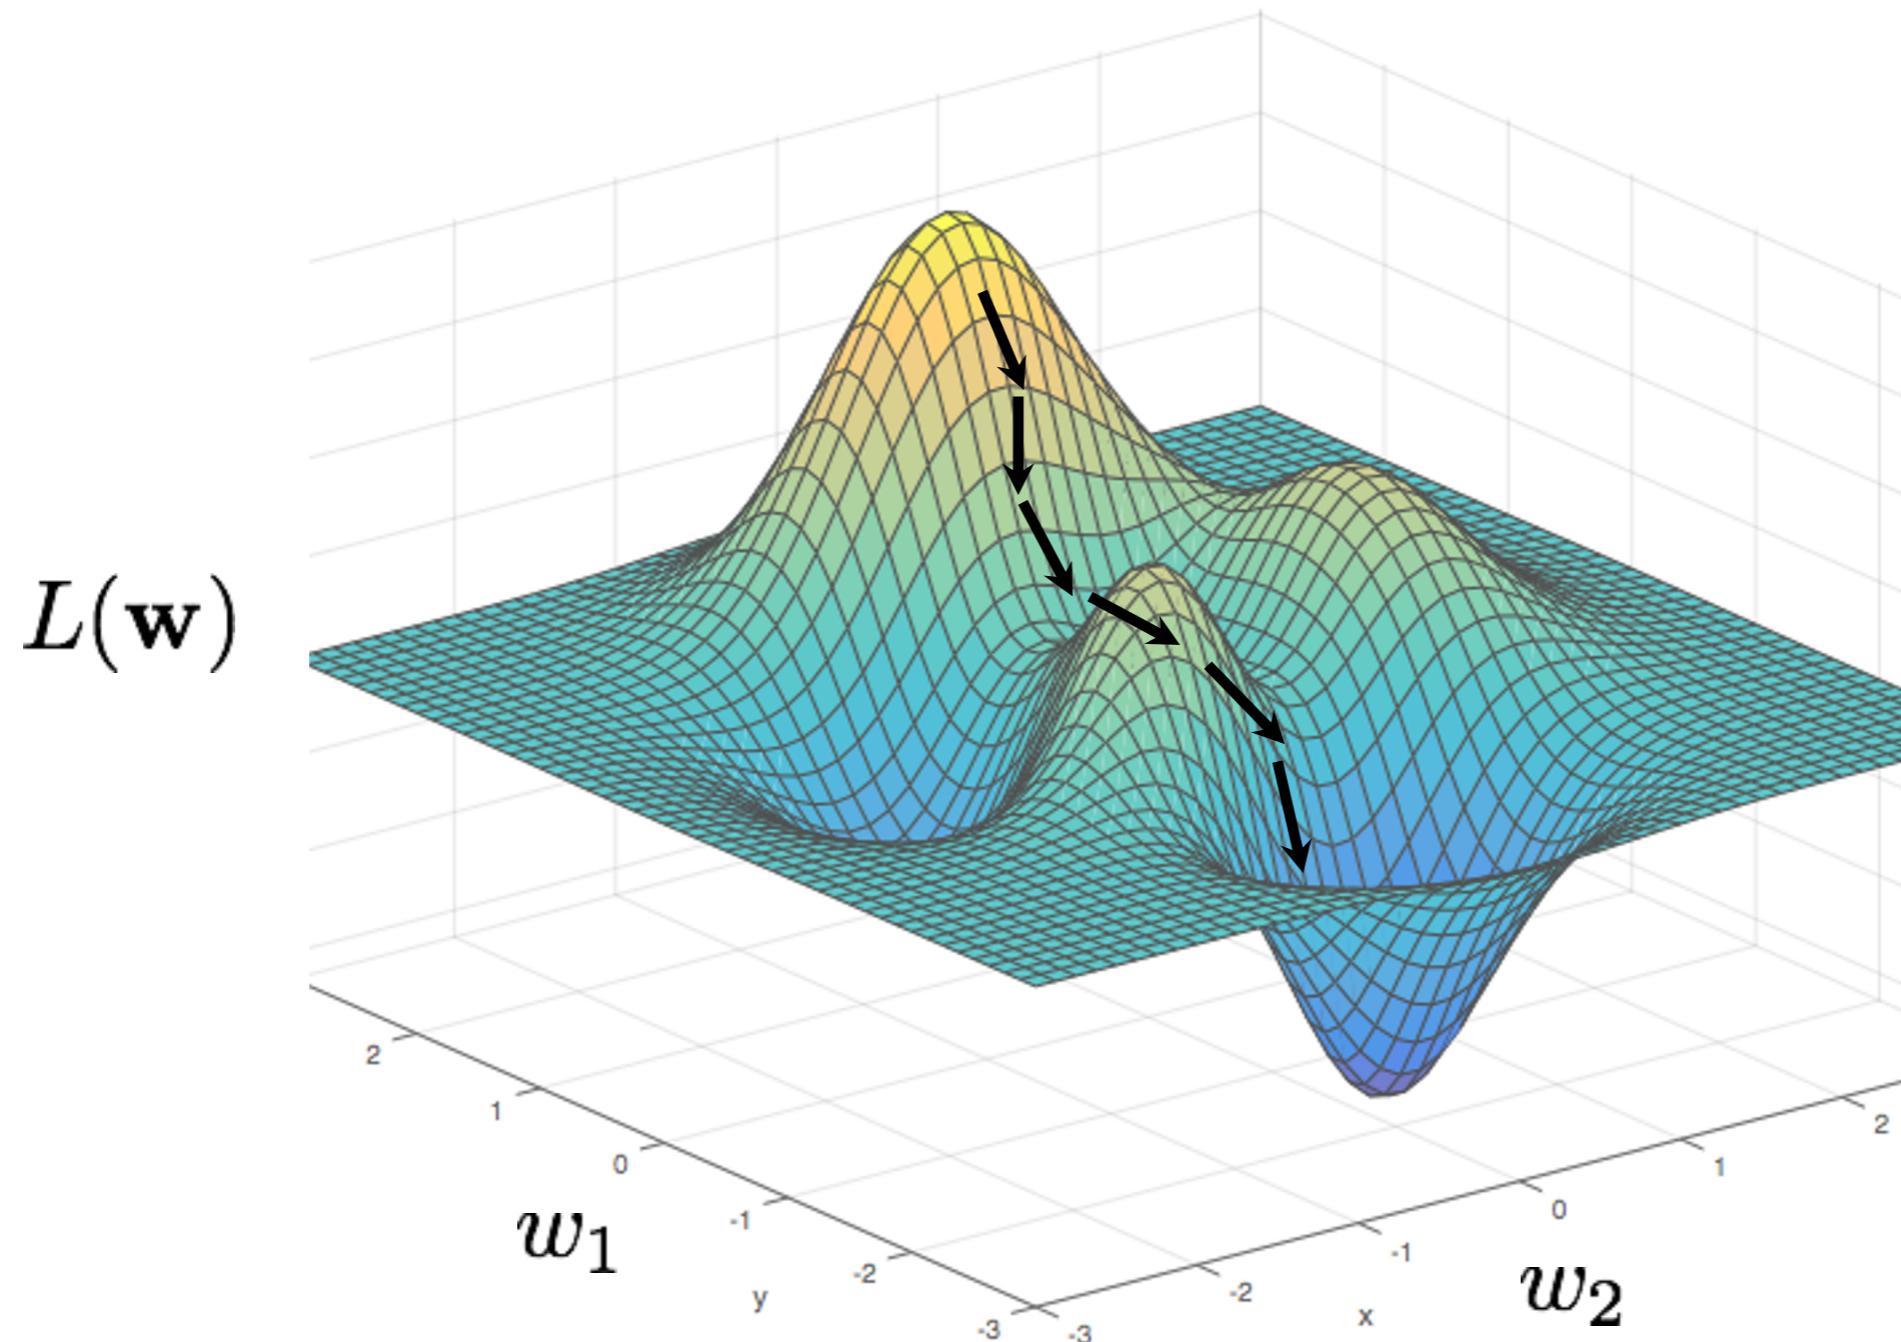
One iteration of gradient descent:

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta_t \frac{\partial L(\mathbf{w}^t)}{\partial \mathbf{w}}$$

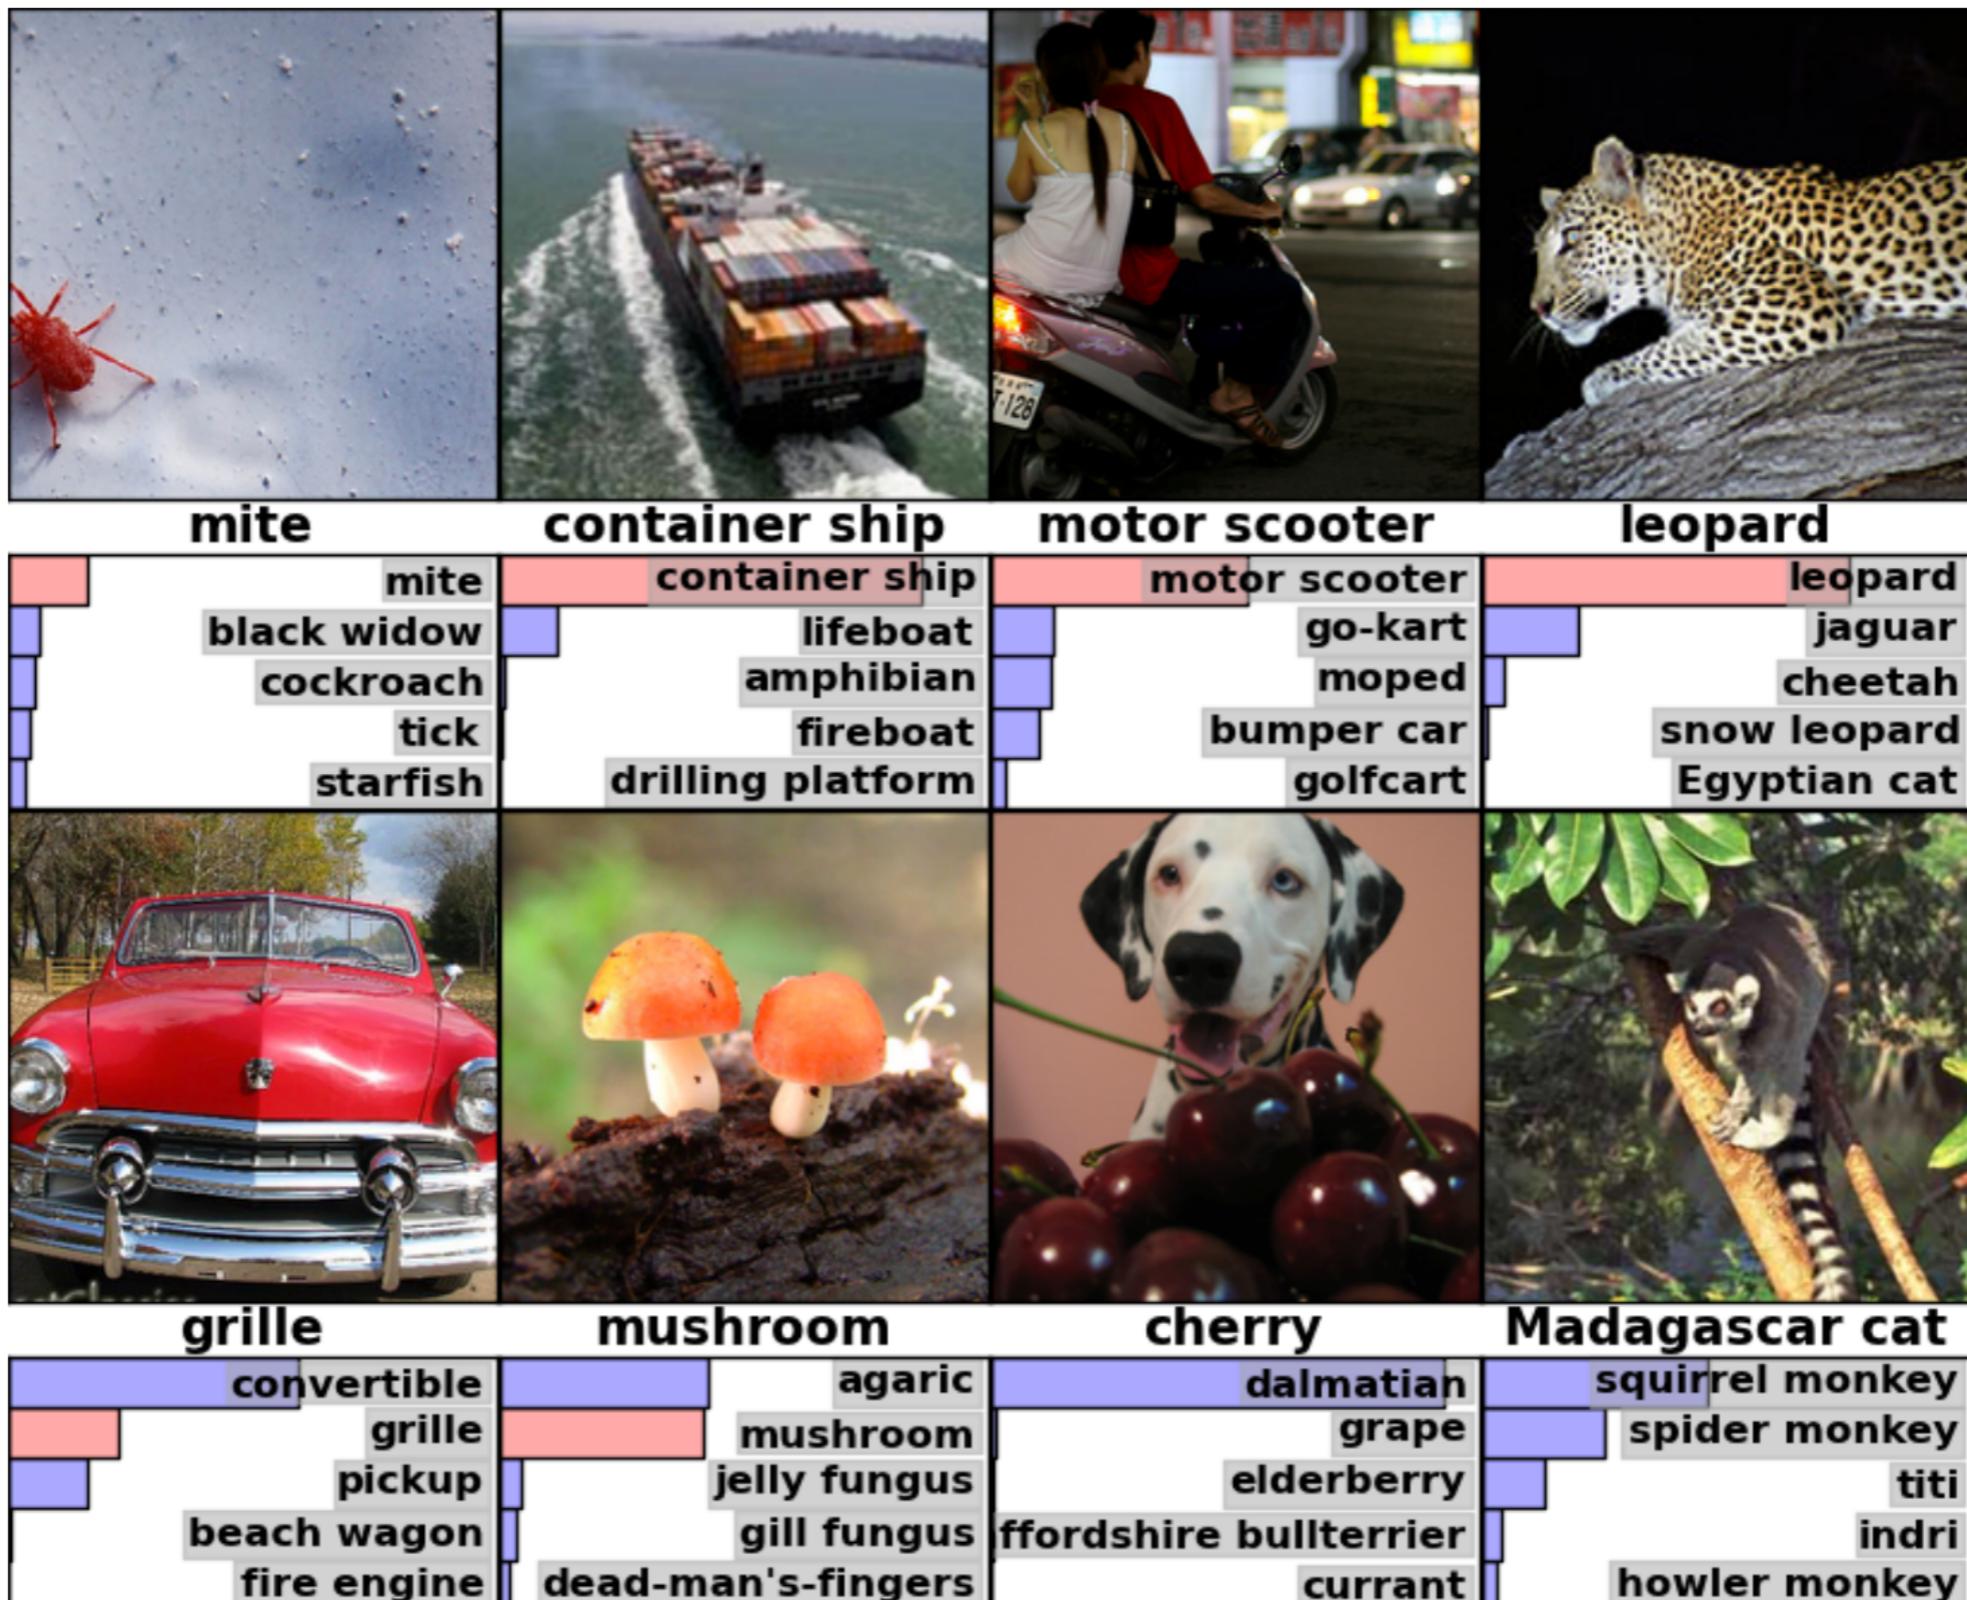


learning rate

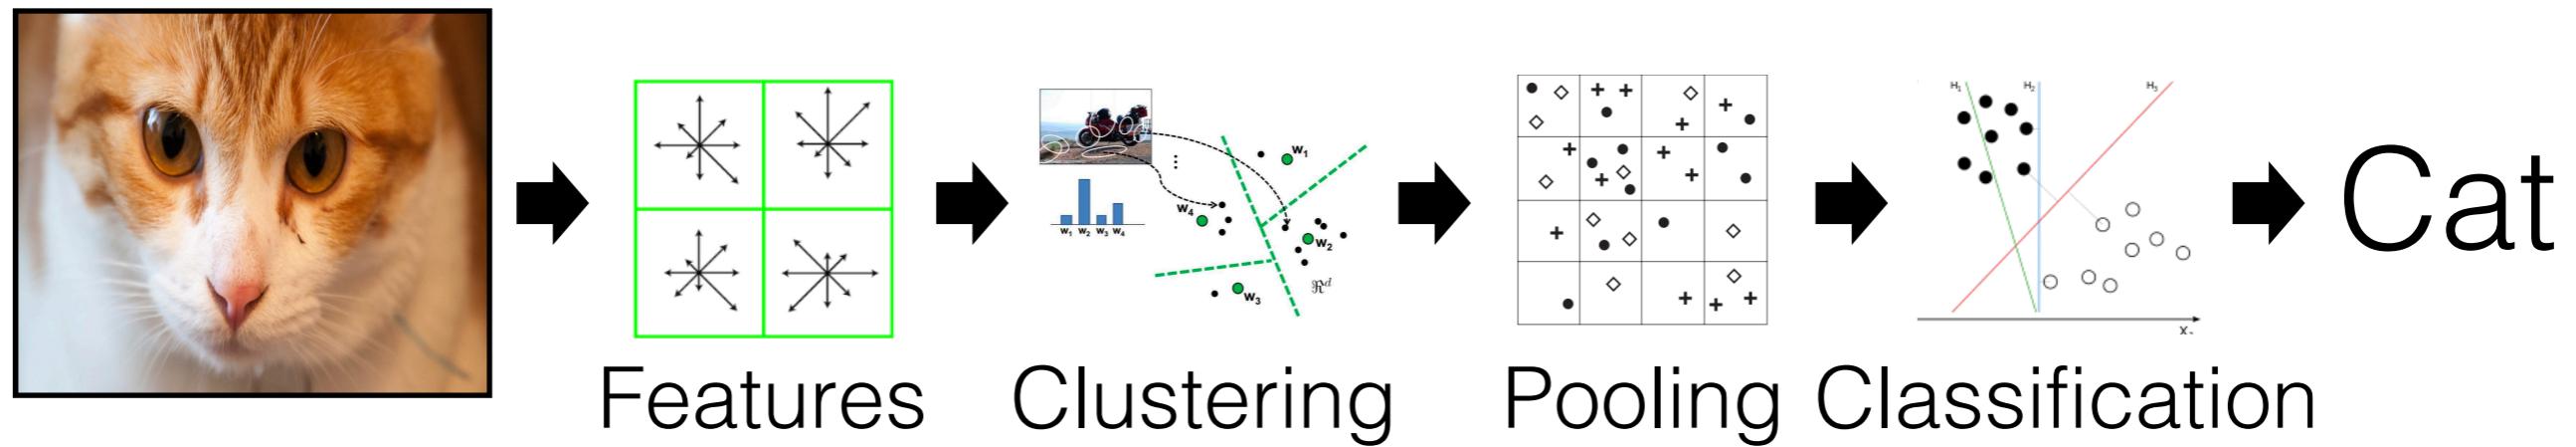
Gradient descent



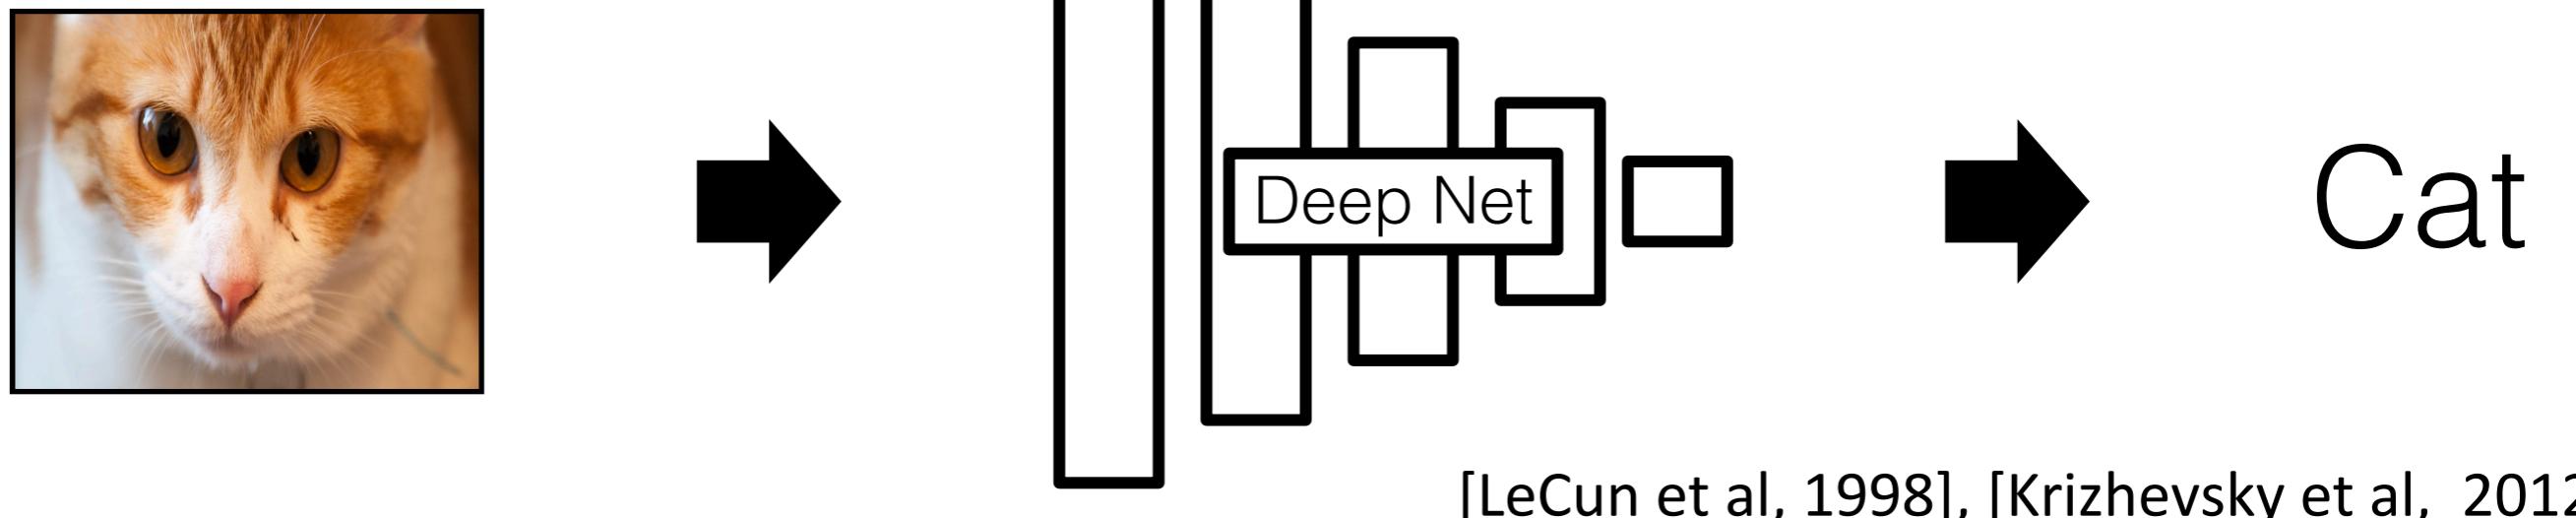
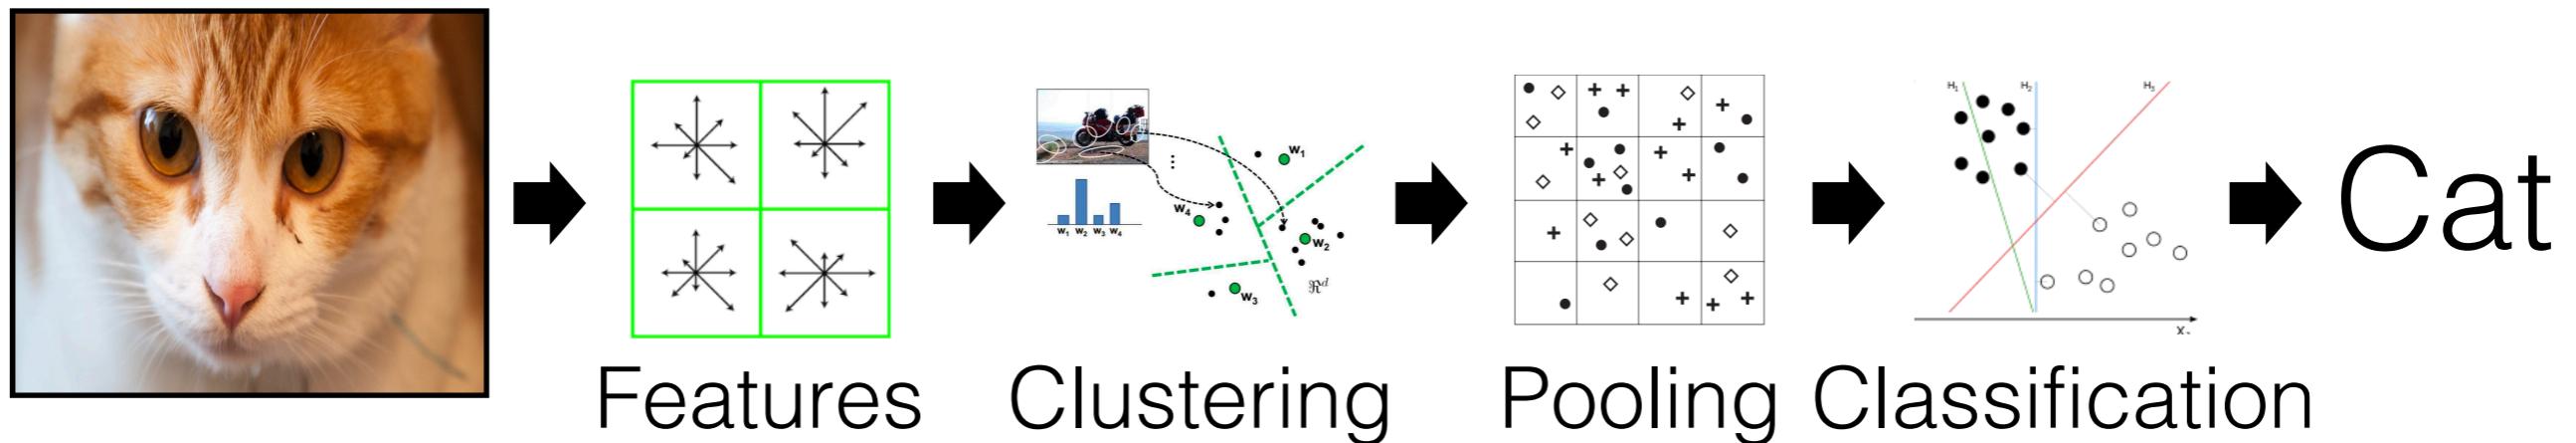
$$p(c|\mathbf{x})$$



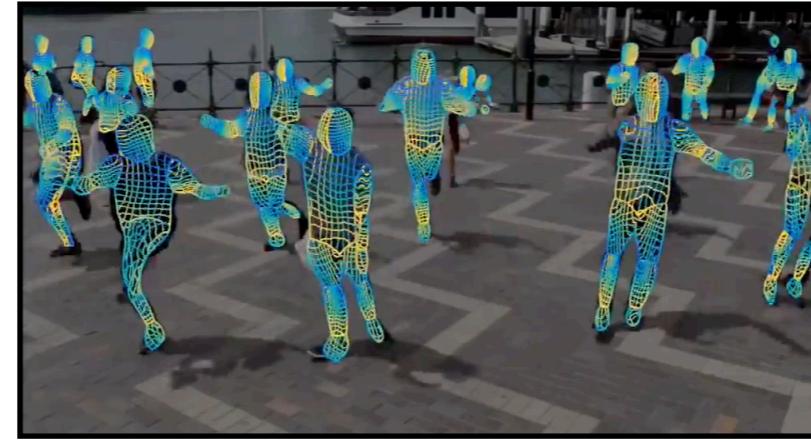
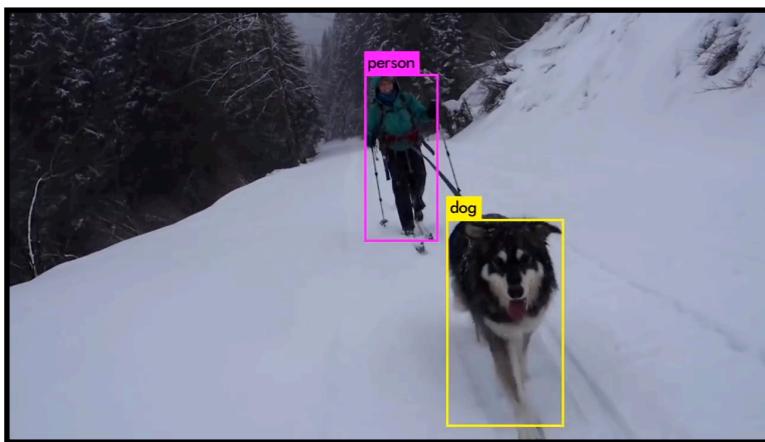
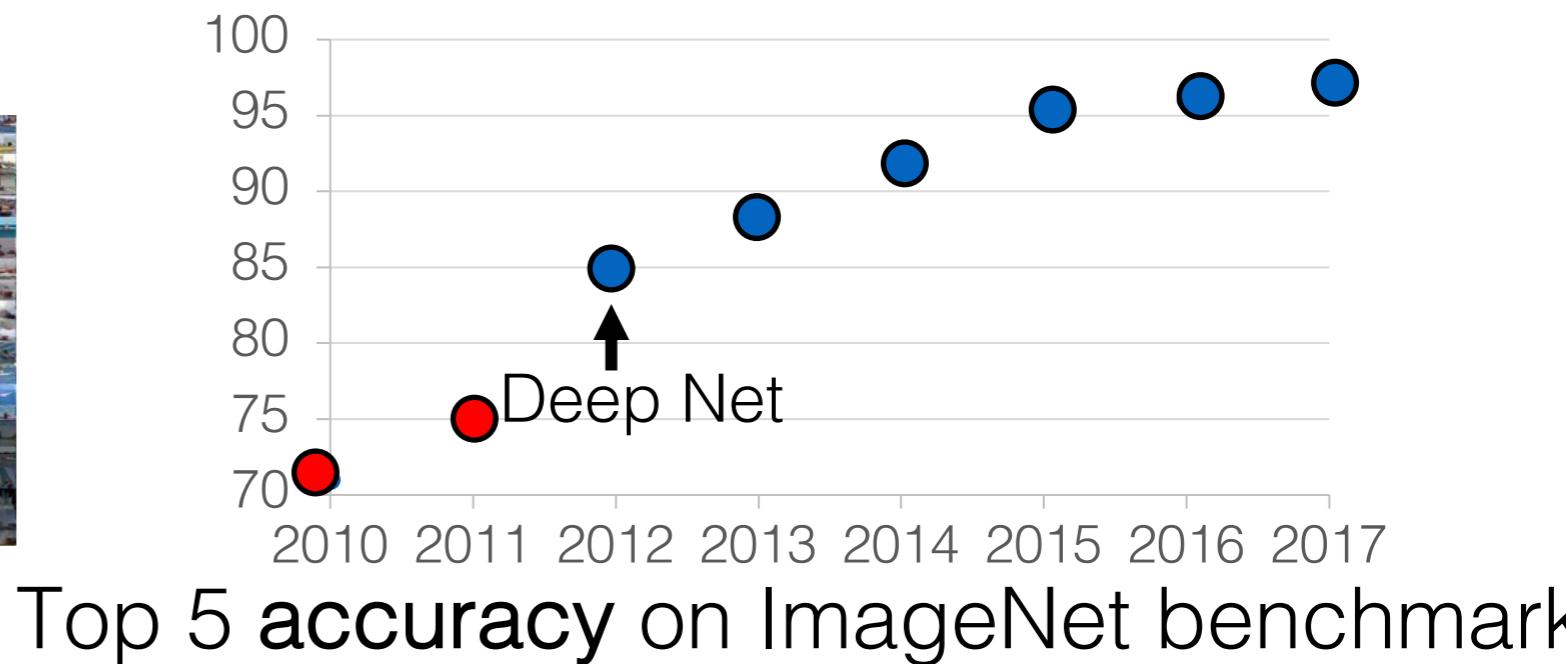
Computer Vision before 2012



Computer Vision Now



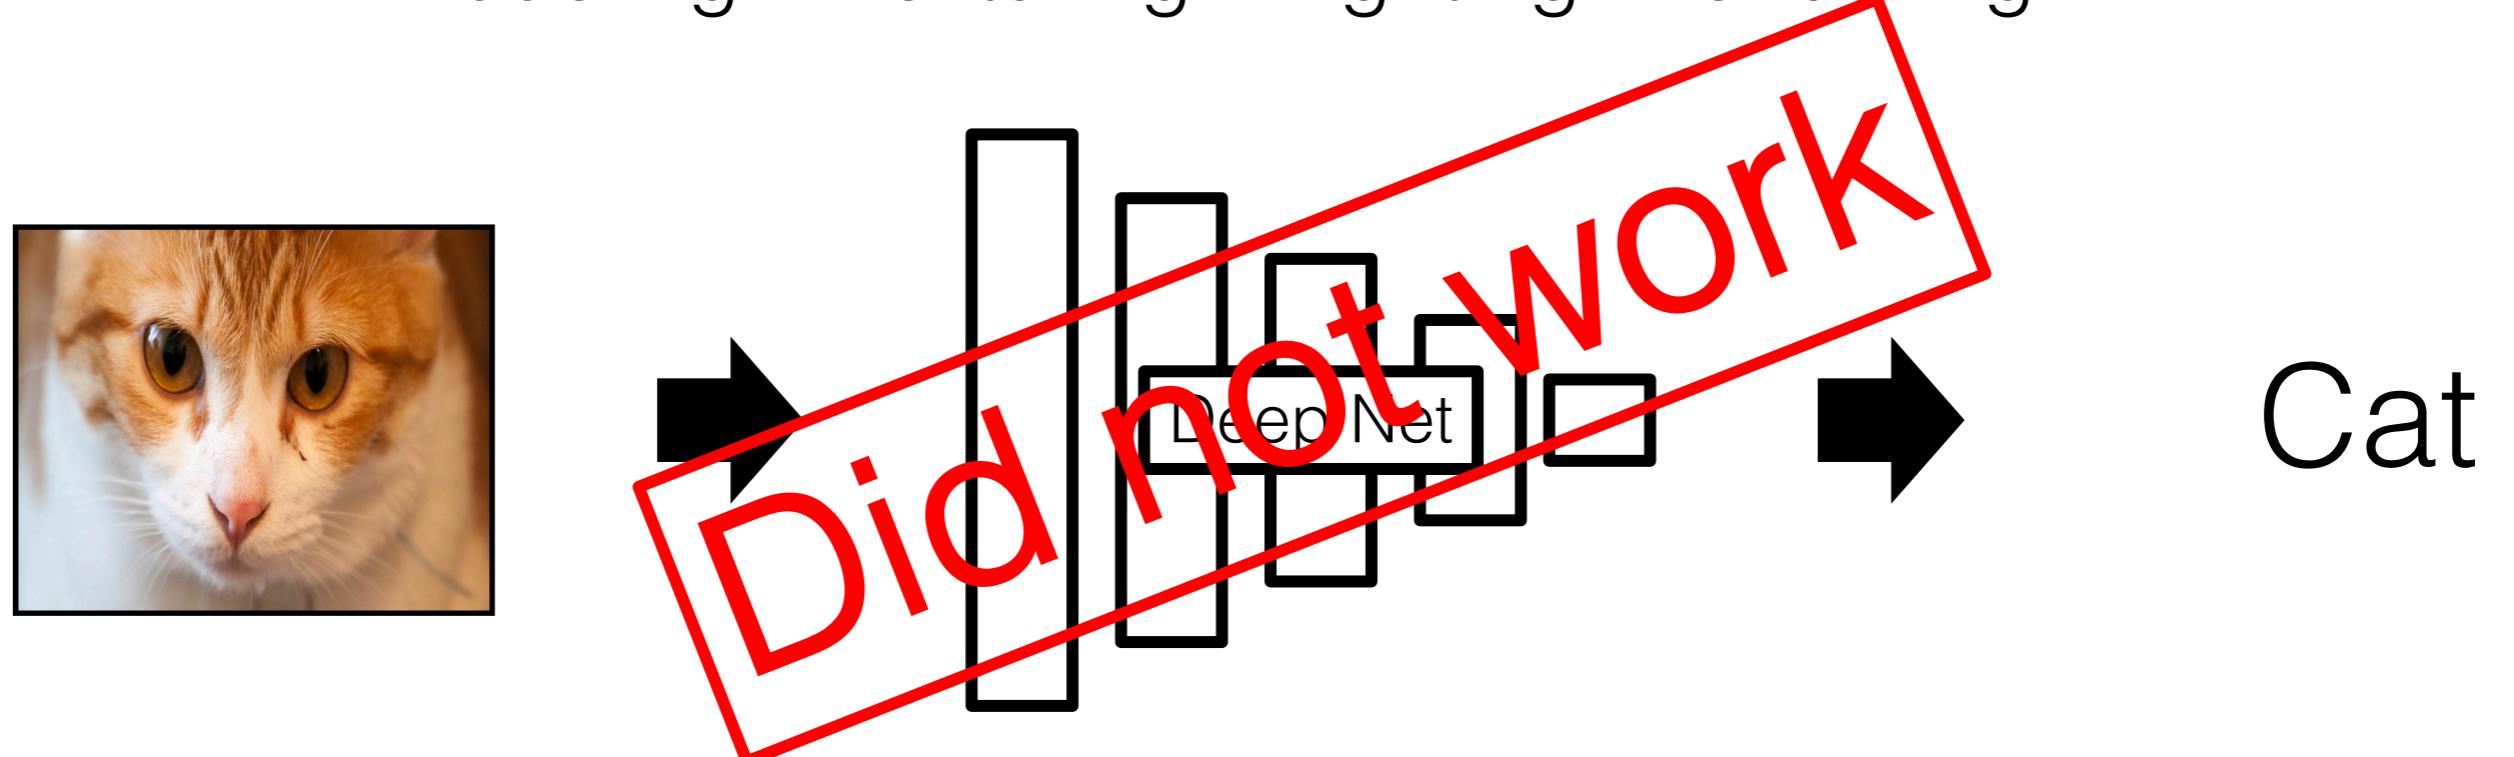
Deep Learning for Computer Vision



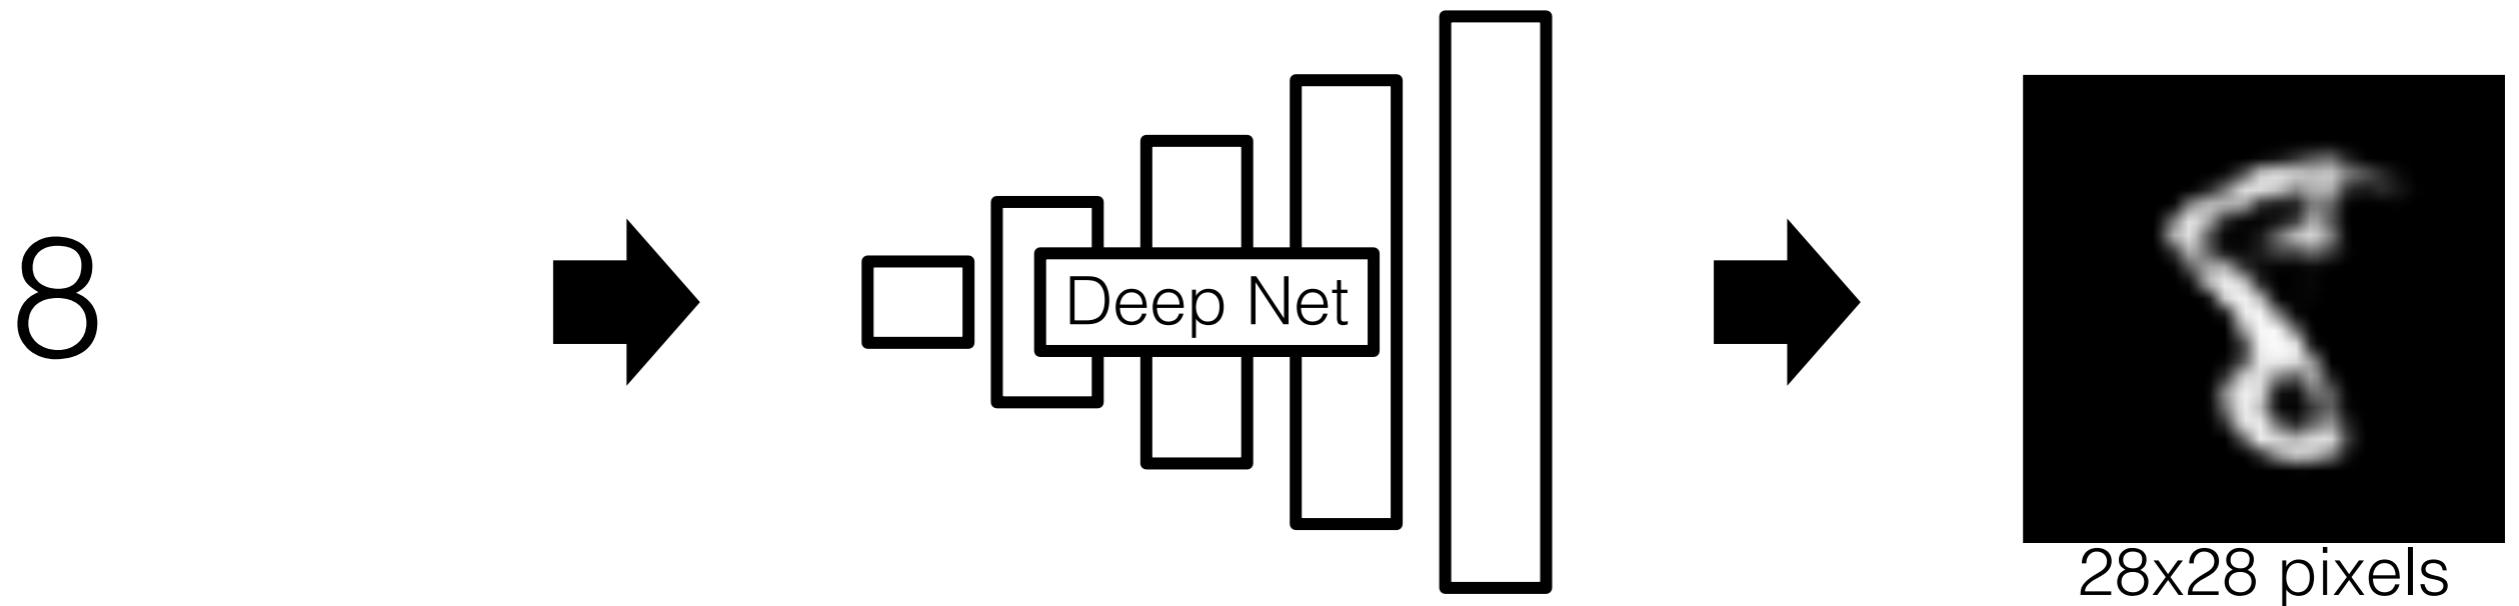
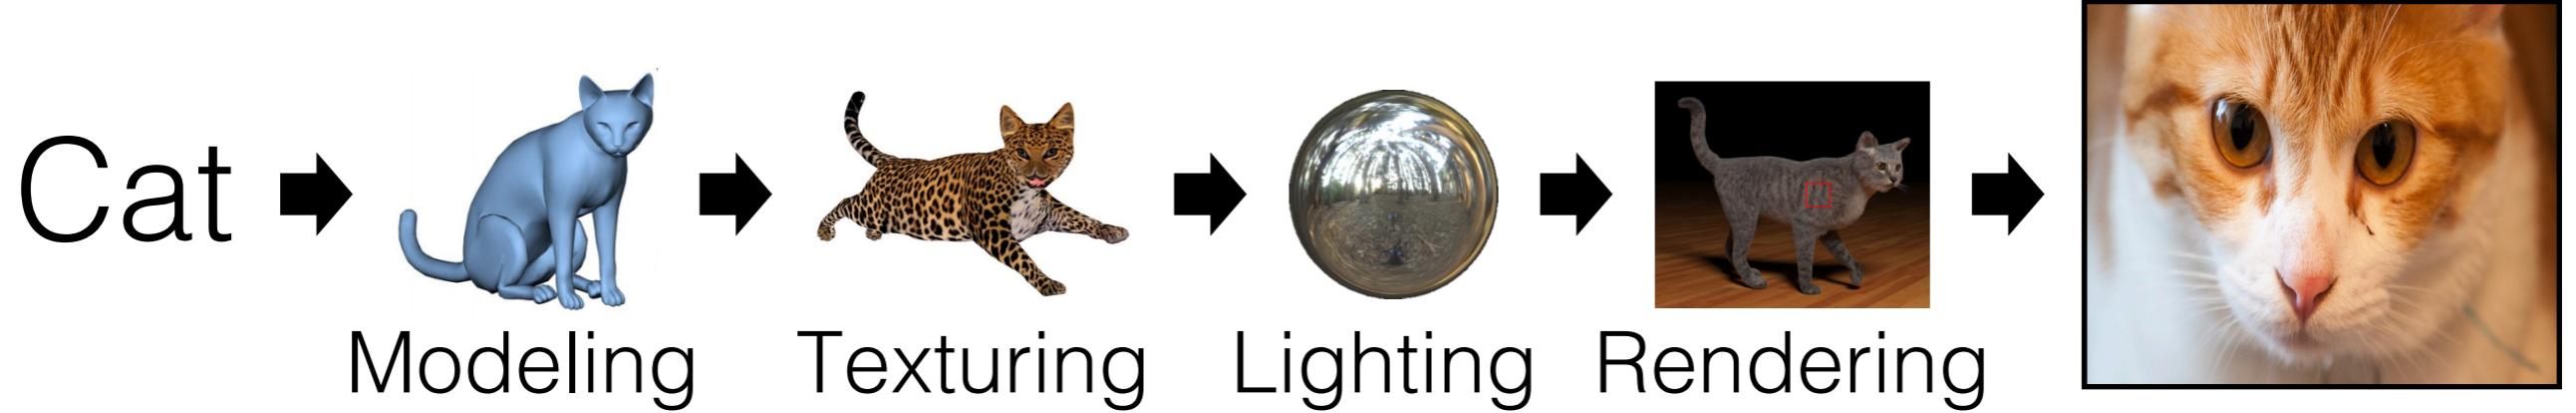
Can Deep Learning Help Graphics?

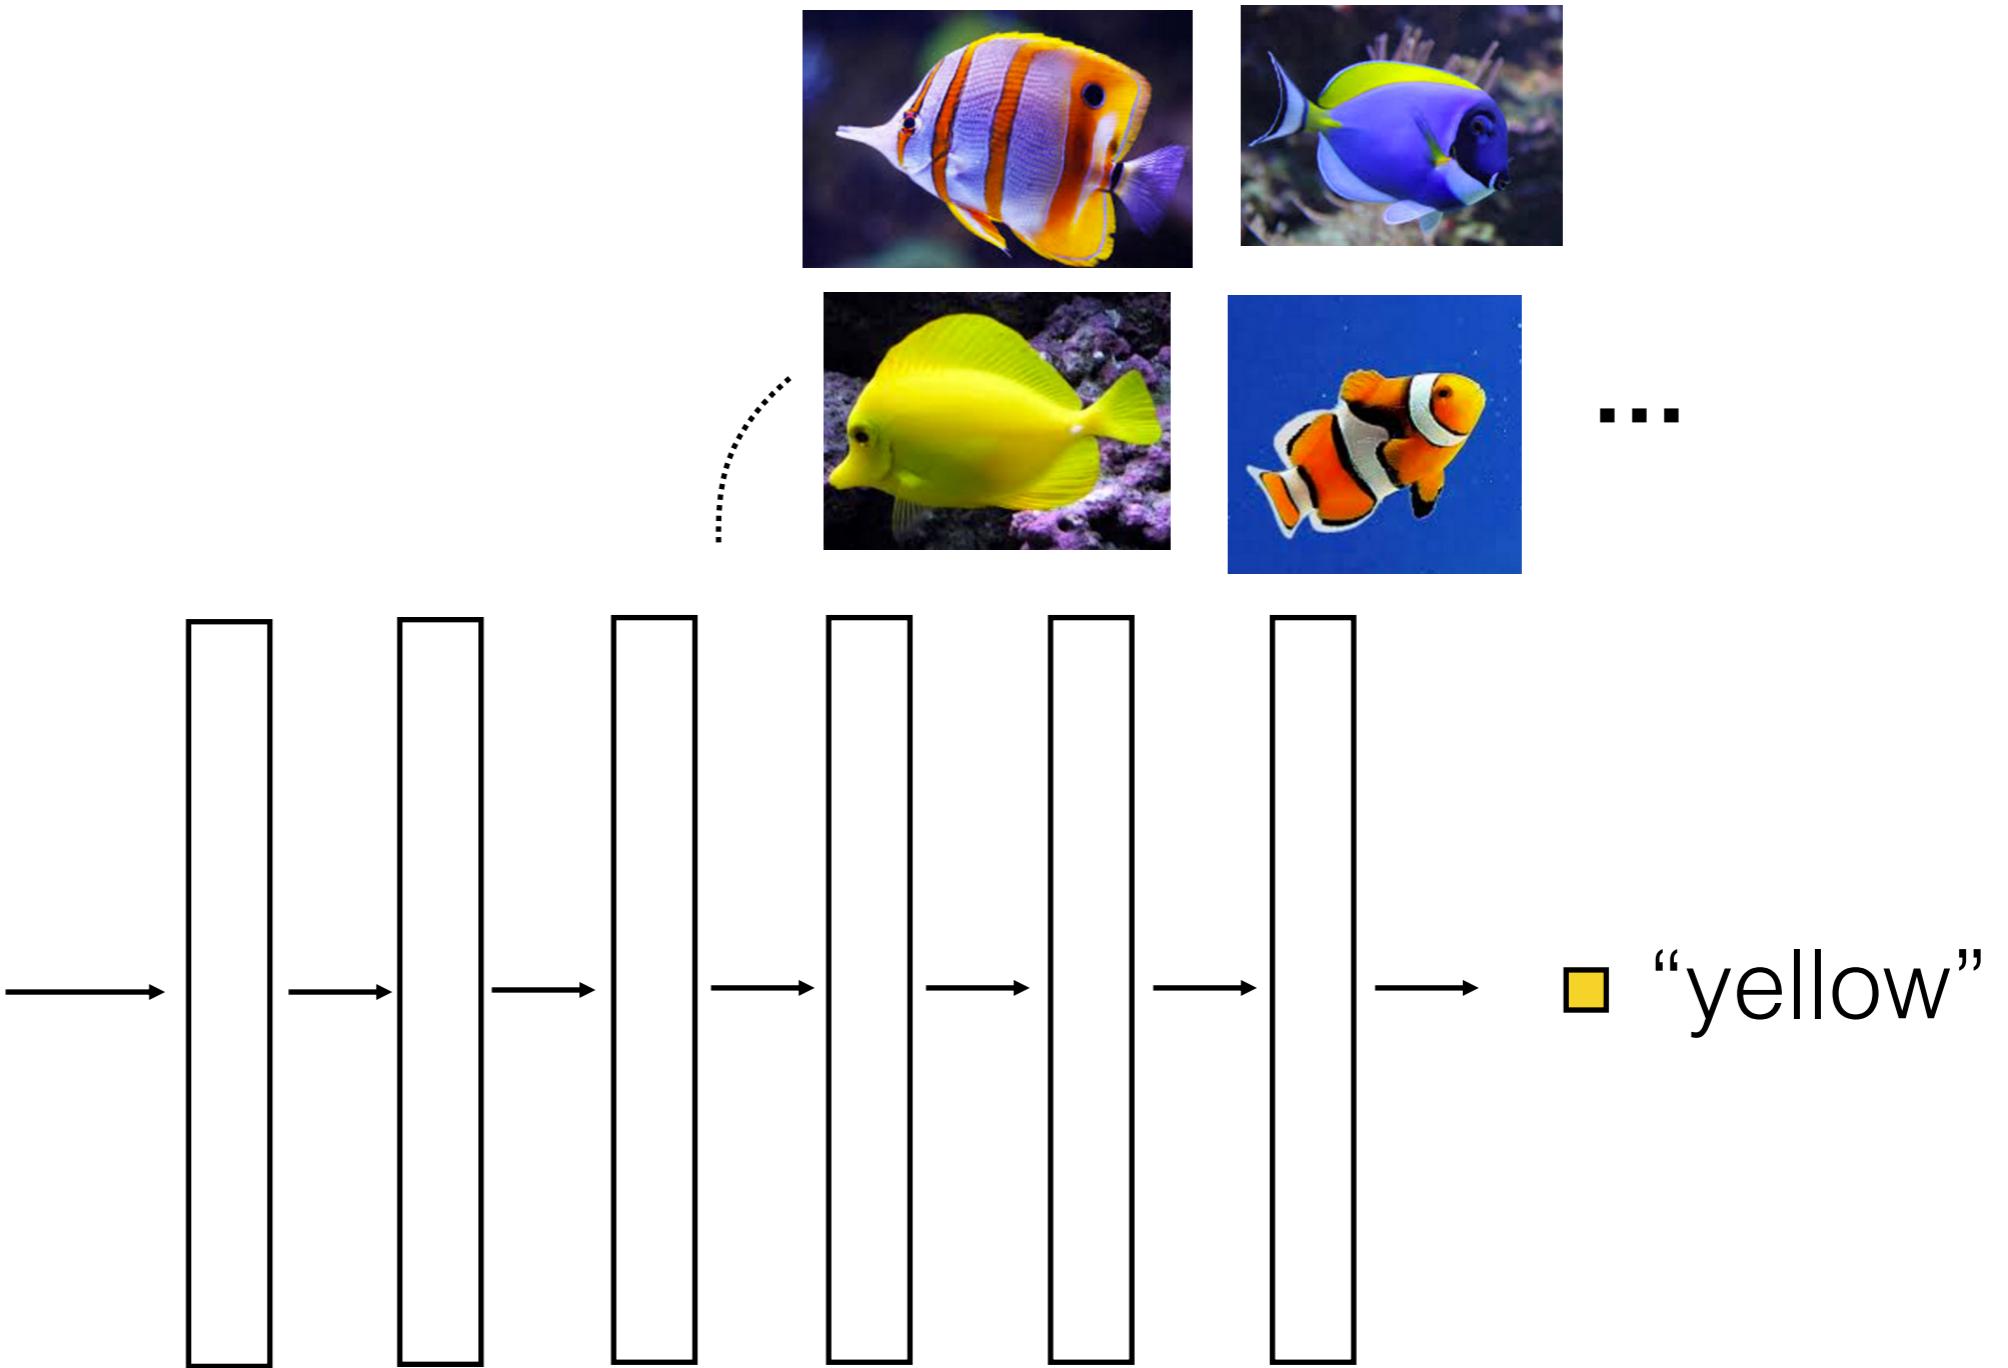


Can Deep Learning Help Graphics?



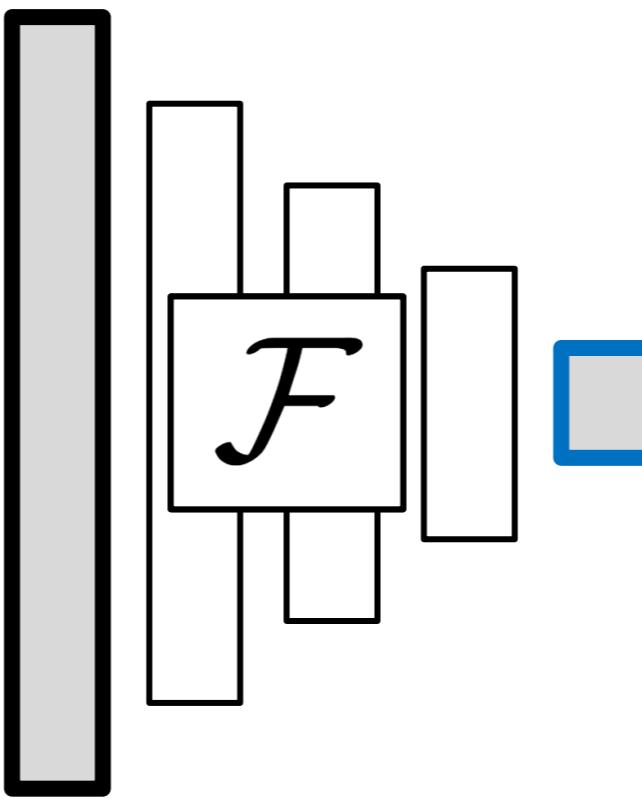
Generating images is hard!



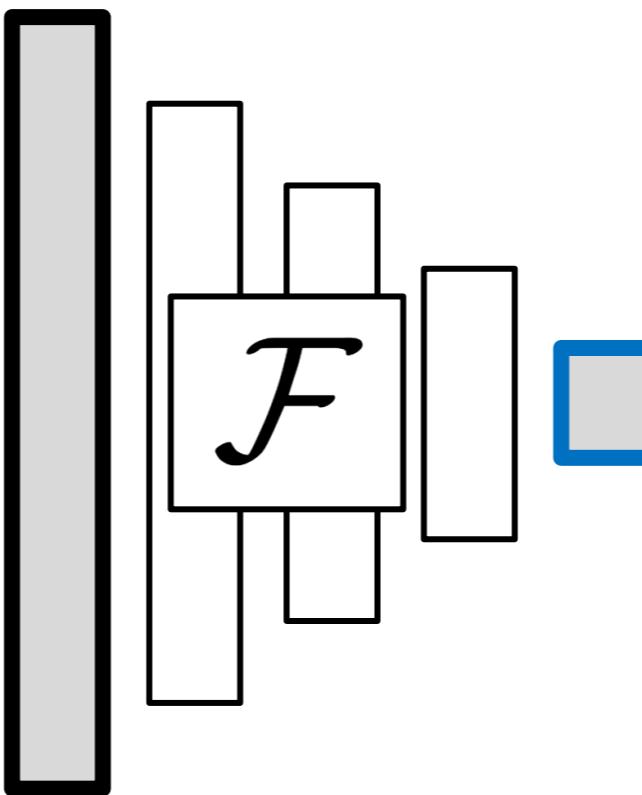


...

Discriminative Deep Networks

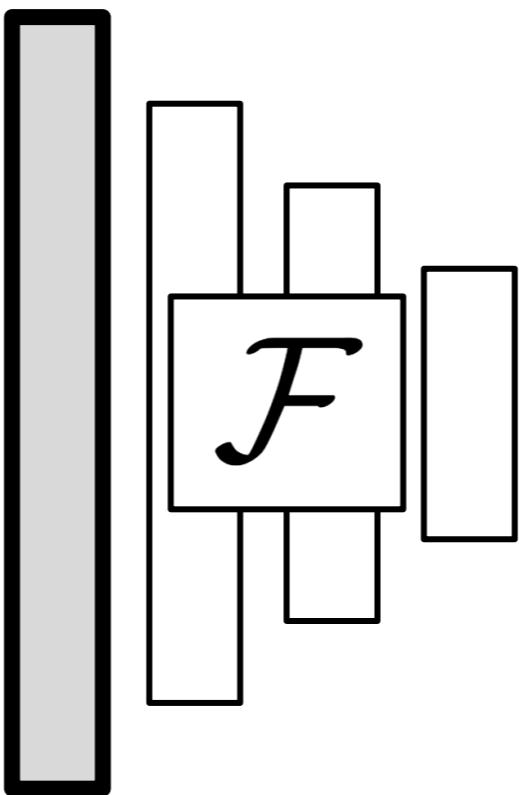
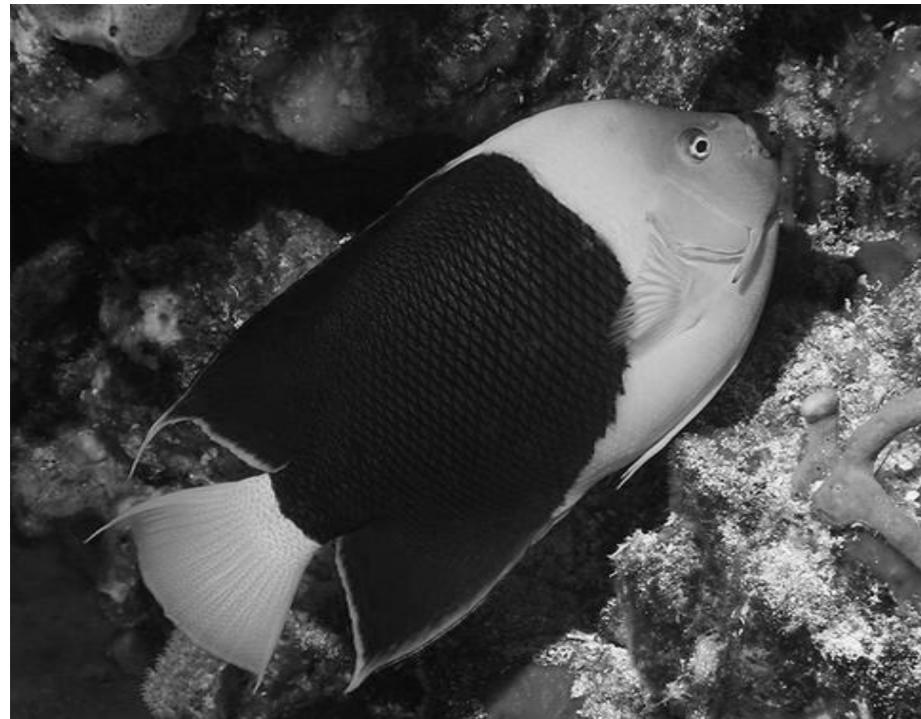


Discriminative Deep Networks



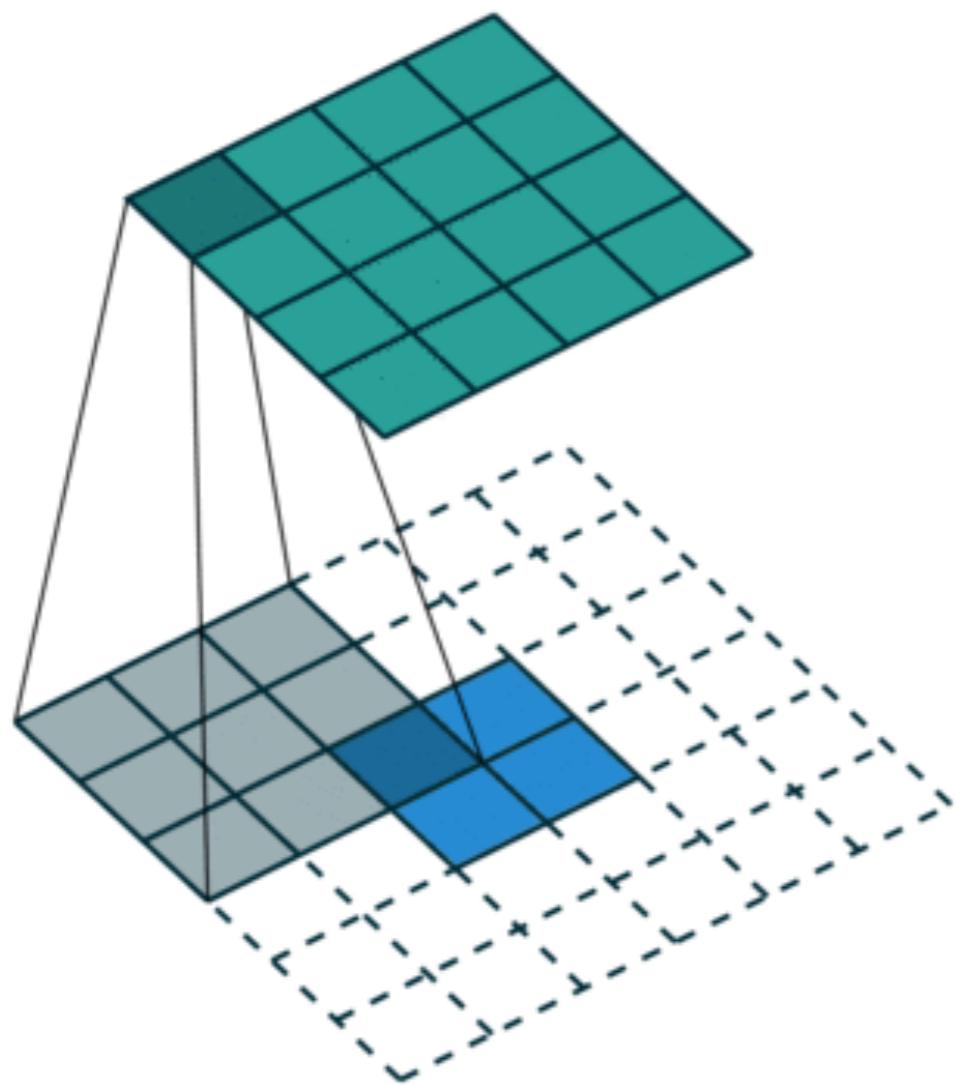
Raw, Unlabeled
Pixels

Generative Deep Networks

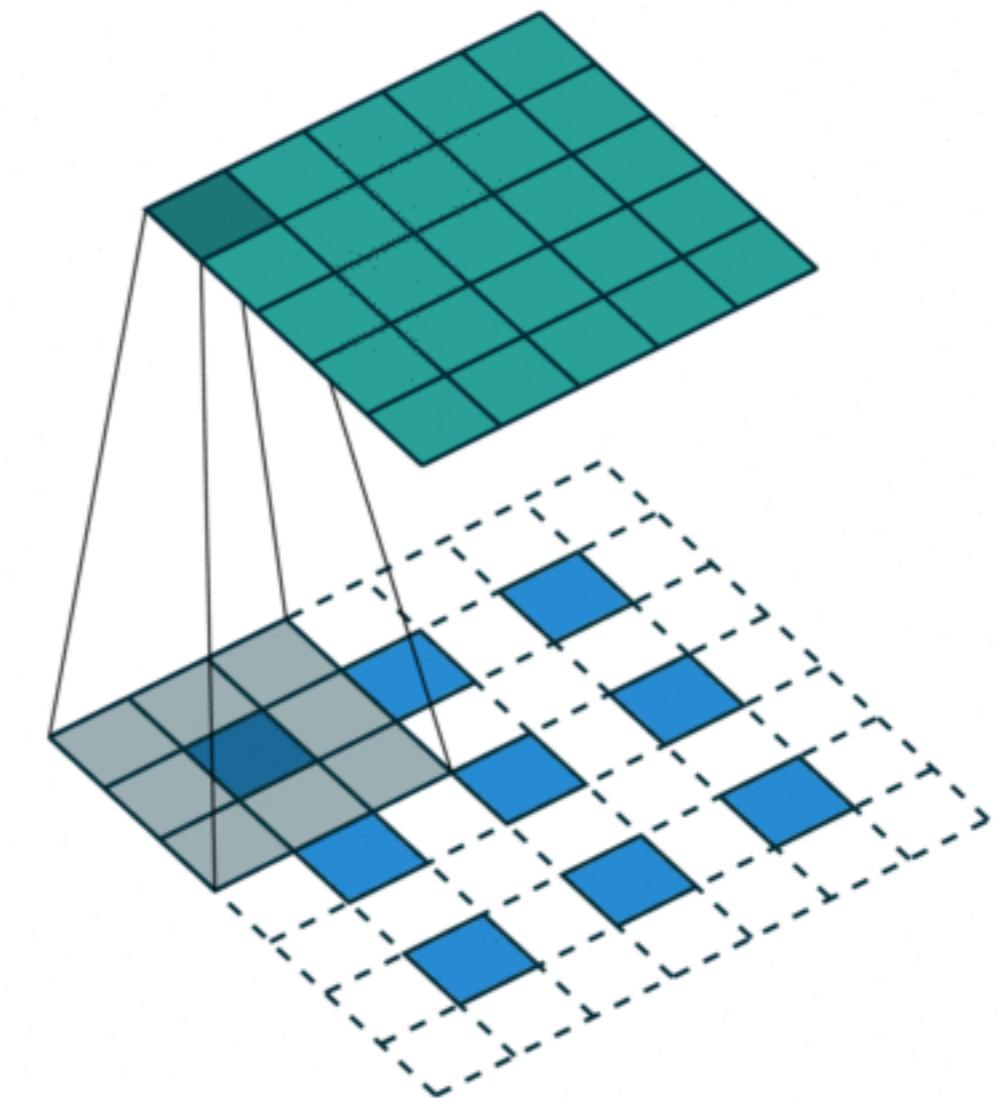


Raw, Unlabeled
Pixels

Fractionally-strided Convolution

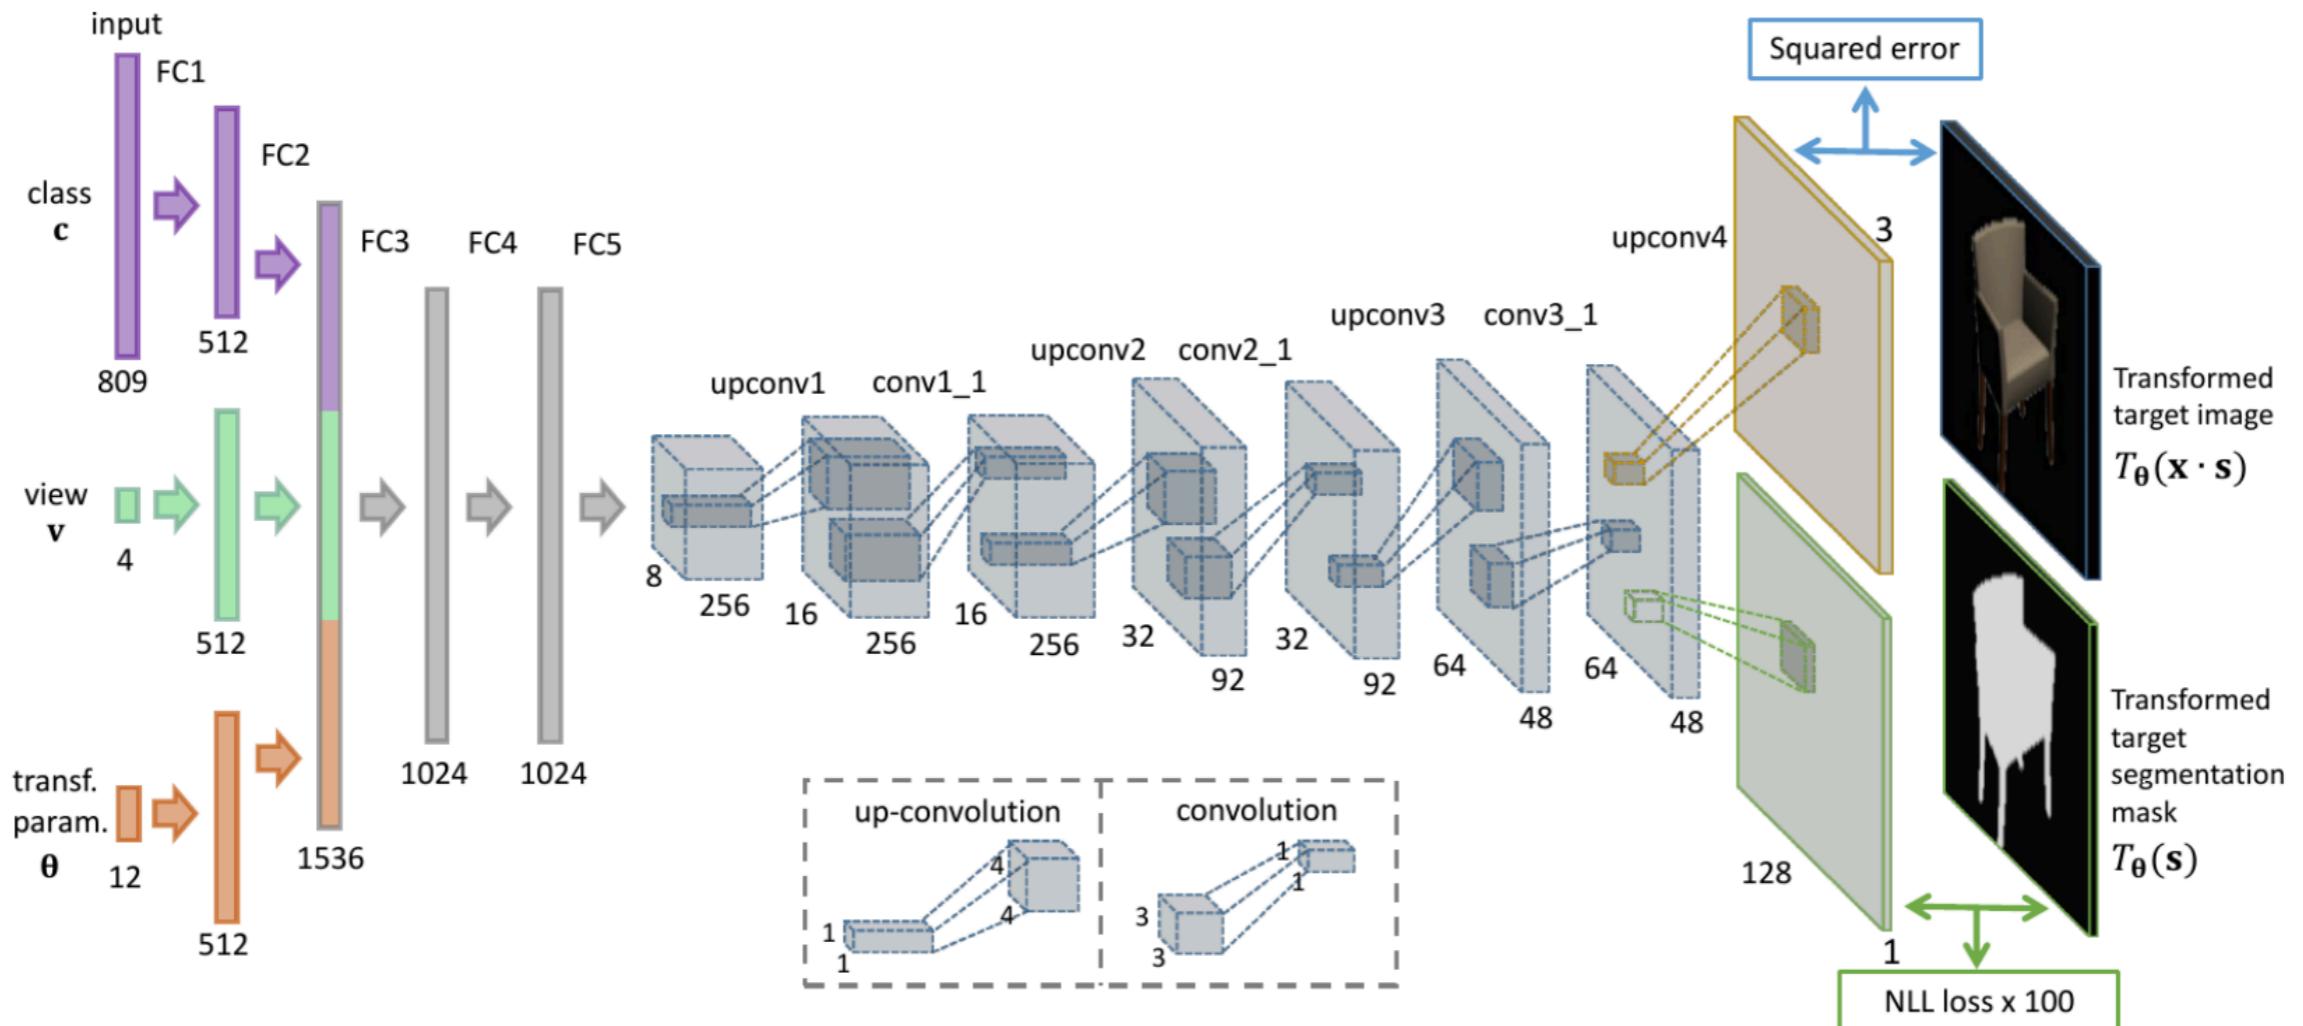


Regular conv



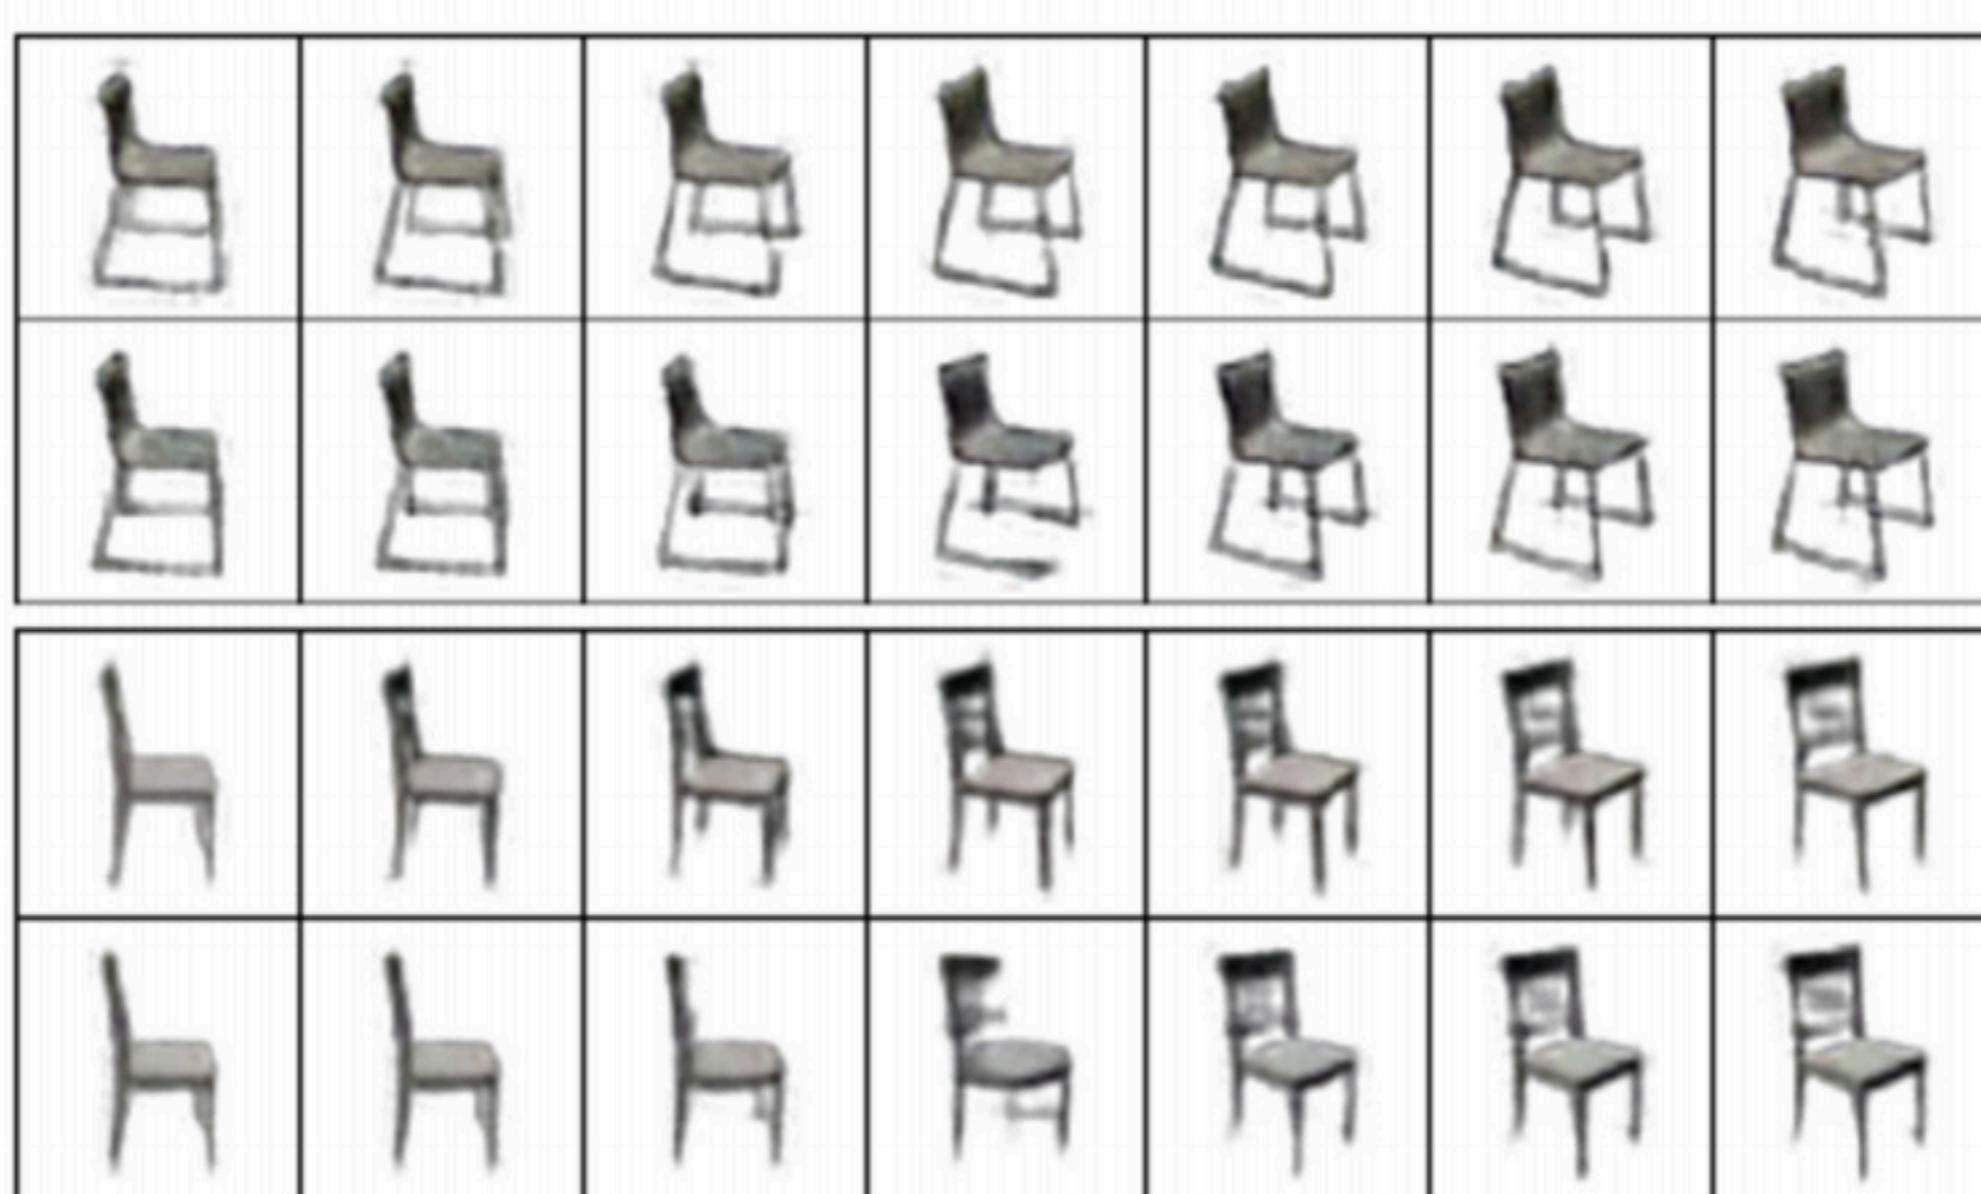
Fractionally-strided conv

Generating chairs given ID, viewpoint, and transformation parameters



Dosovitskiy et al. Learning to Generate Chairs, Tables and Cars with Convolutional Networks
PAMI 2017 (CVPR 2015)

Generating chairs given ID, viewpoint, and transformation parameters

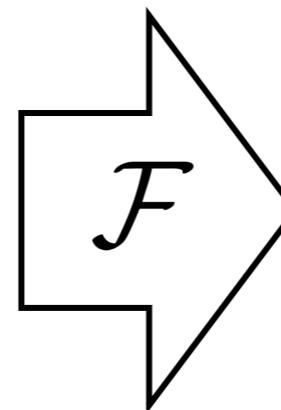
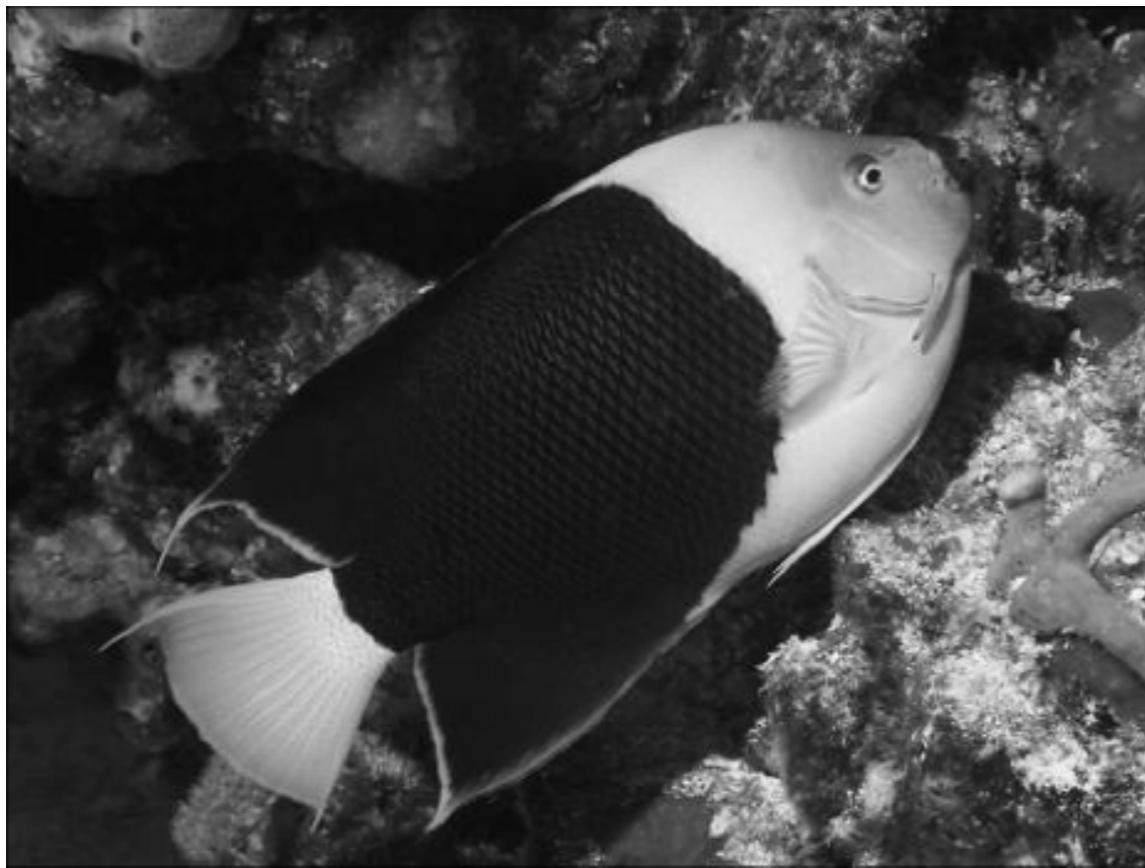


Dosovitskiy et al. Learning to Generate Chairs, Tables and Cars with Convolutional Networks
PAMI 2017 (CVPR 2015)

60

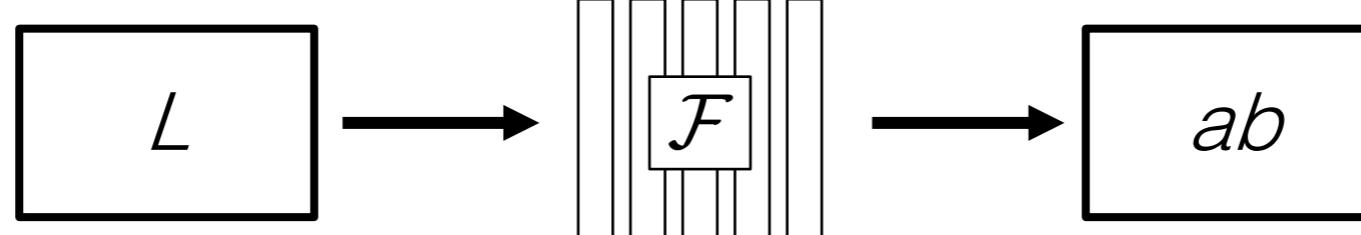


Ansel Adams. *Yosemite Valley Bridge.*



Grayscale image: L channel

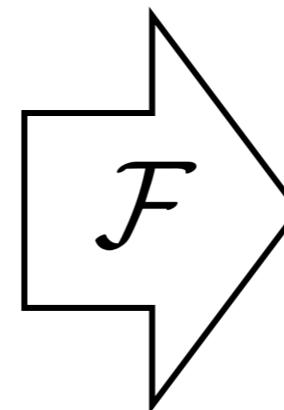
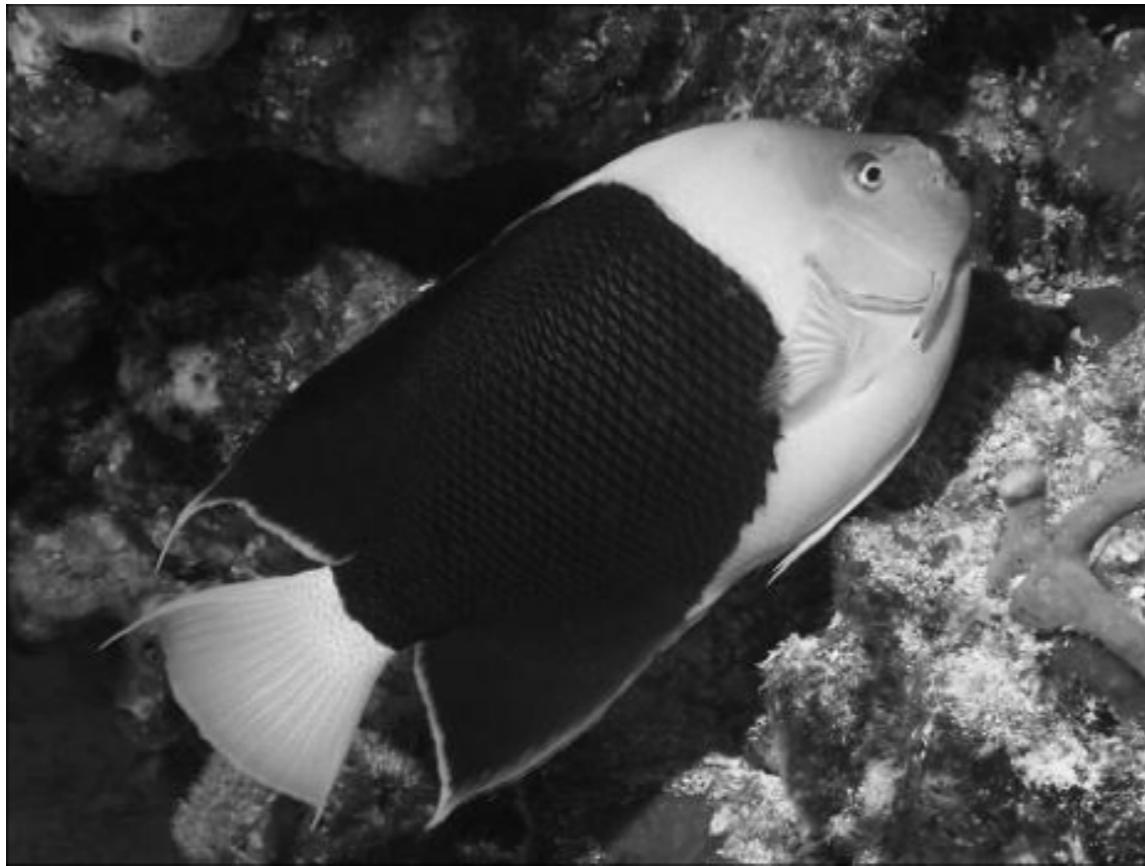
$$\mathbf{X} \in \mathbb{R}^{H \times W \times 1}$$



Color information: ab channels

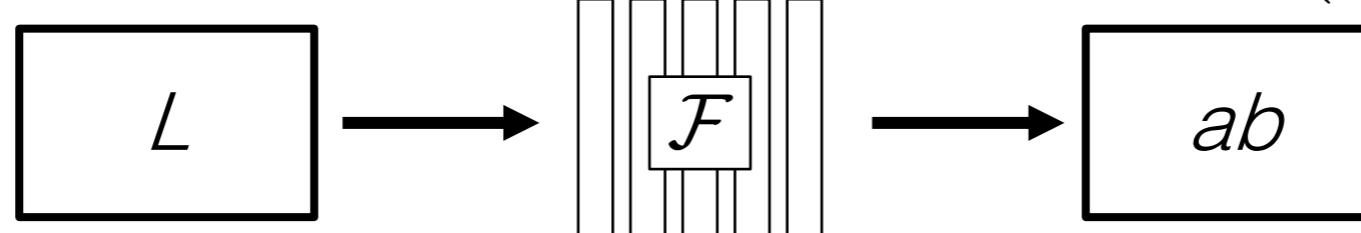
$$\hat{\mathbf{Y}} \in \mathbb{R}^{H \times W \times 61^2}$$

Zhang, Isola, Efros. *Colorful Image Colorization*. In *ECCV*, 2016.



Grayscale image: L channel

$$\mathbf{X} \in \mathbb{R}^{H \times W \times 1}$$



Zhang, Isola, Efros. *Colorful Image Colorization*. In *ECCV*, 2016.

Concatenate (L, ab) channels

$$(\mathbf{X}, \hat{\mathbf{Y}})^{62}$$

Simple L2 regression doesn't work ☹

Input



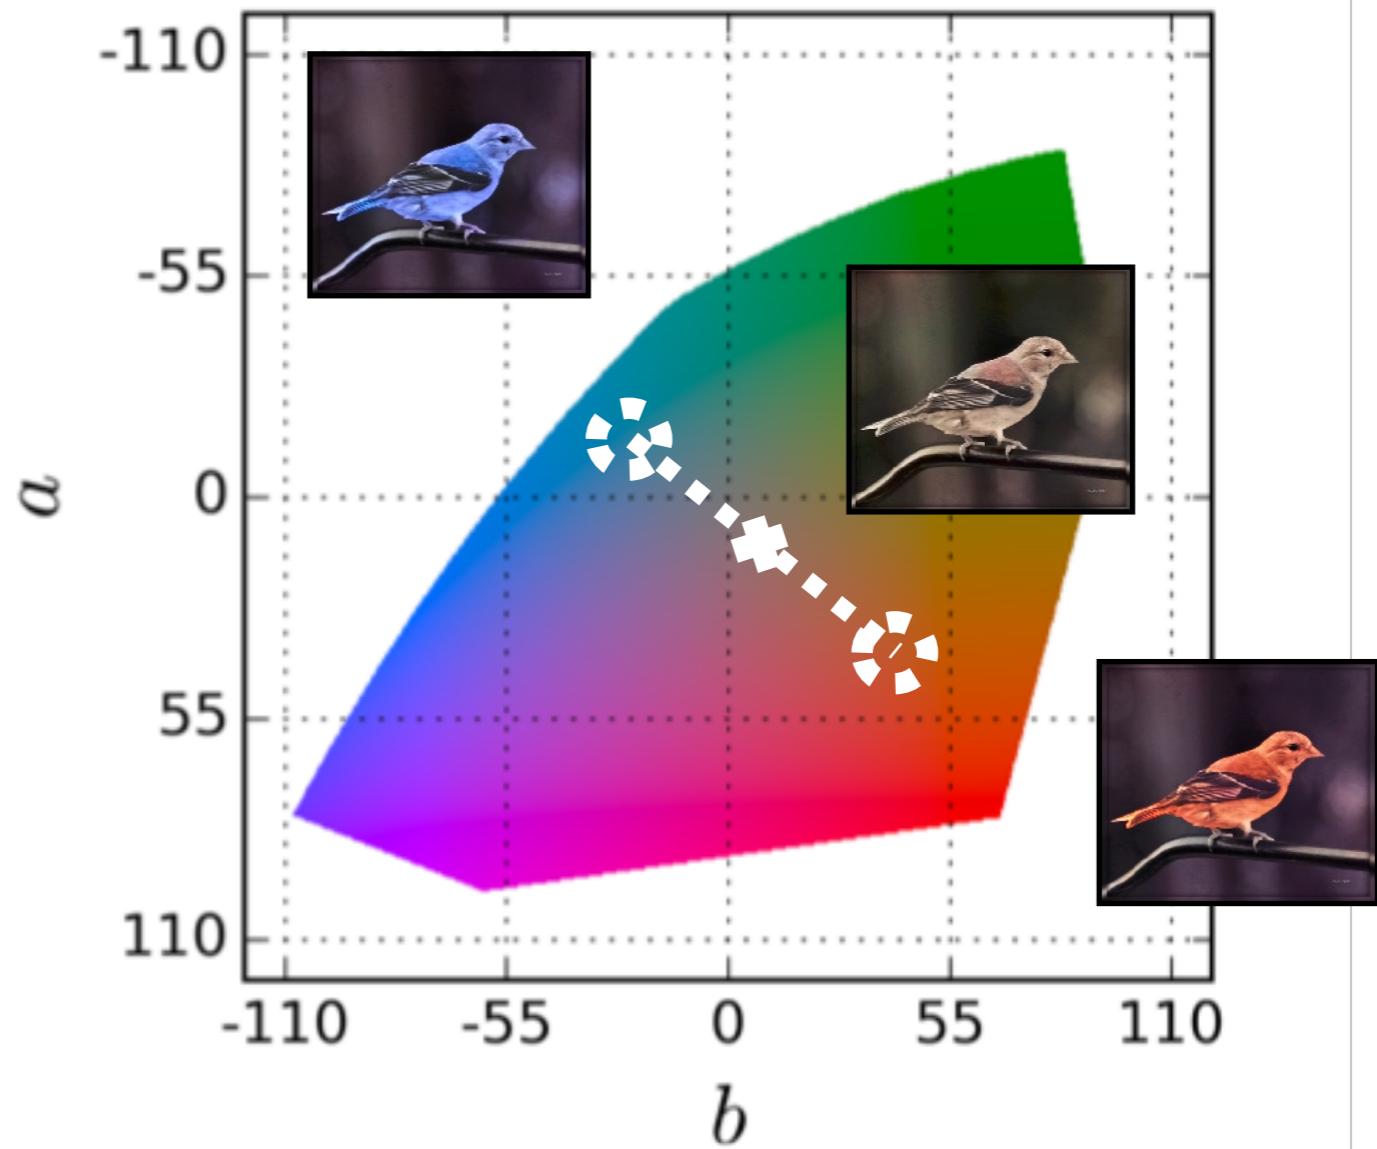
Output



Ground truth



$$L_2(\hat{\mathbf{Y}}, \mathbf{Y}) = \frac{1}{2} \sum_{h,w} \|\mathbf{Y}_{h,w} - \hat{\mathbf{Y}}_{h,w}\|_2^2$$



$$L_2(\hat{\mathbf{Y}}, \mathbf{Y}) = \frac{1}{2} \sum_{h,w} \|\mathbf{Y}_{h,w} - \hat{\mathbf{Y}}_{h,w}\|_2^2$$

Better Loss Function

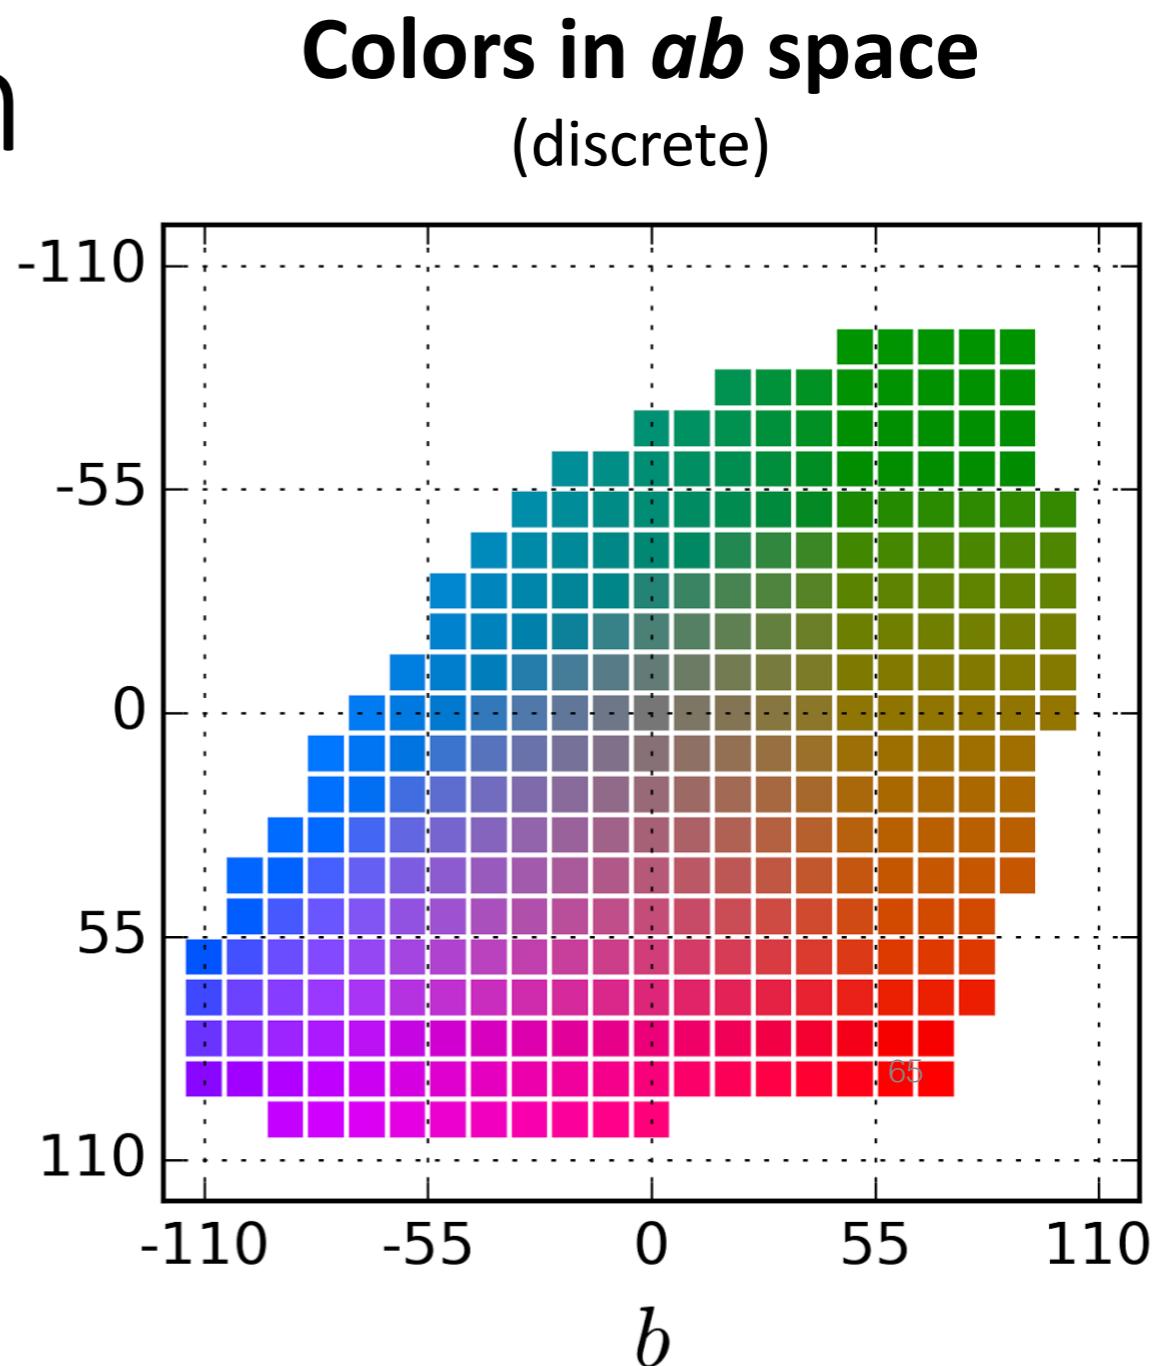
$$\theta^* = \arg \min_{\theta} \ell(\mathcal{F}_{\theta}(\mathbf{X}), \mathbf{Y})$$

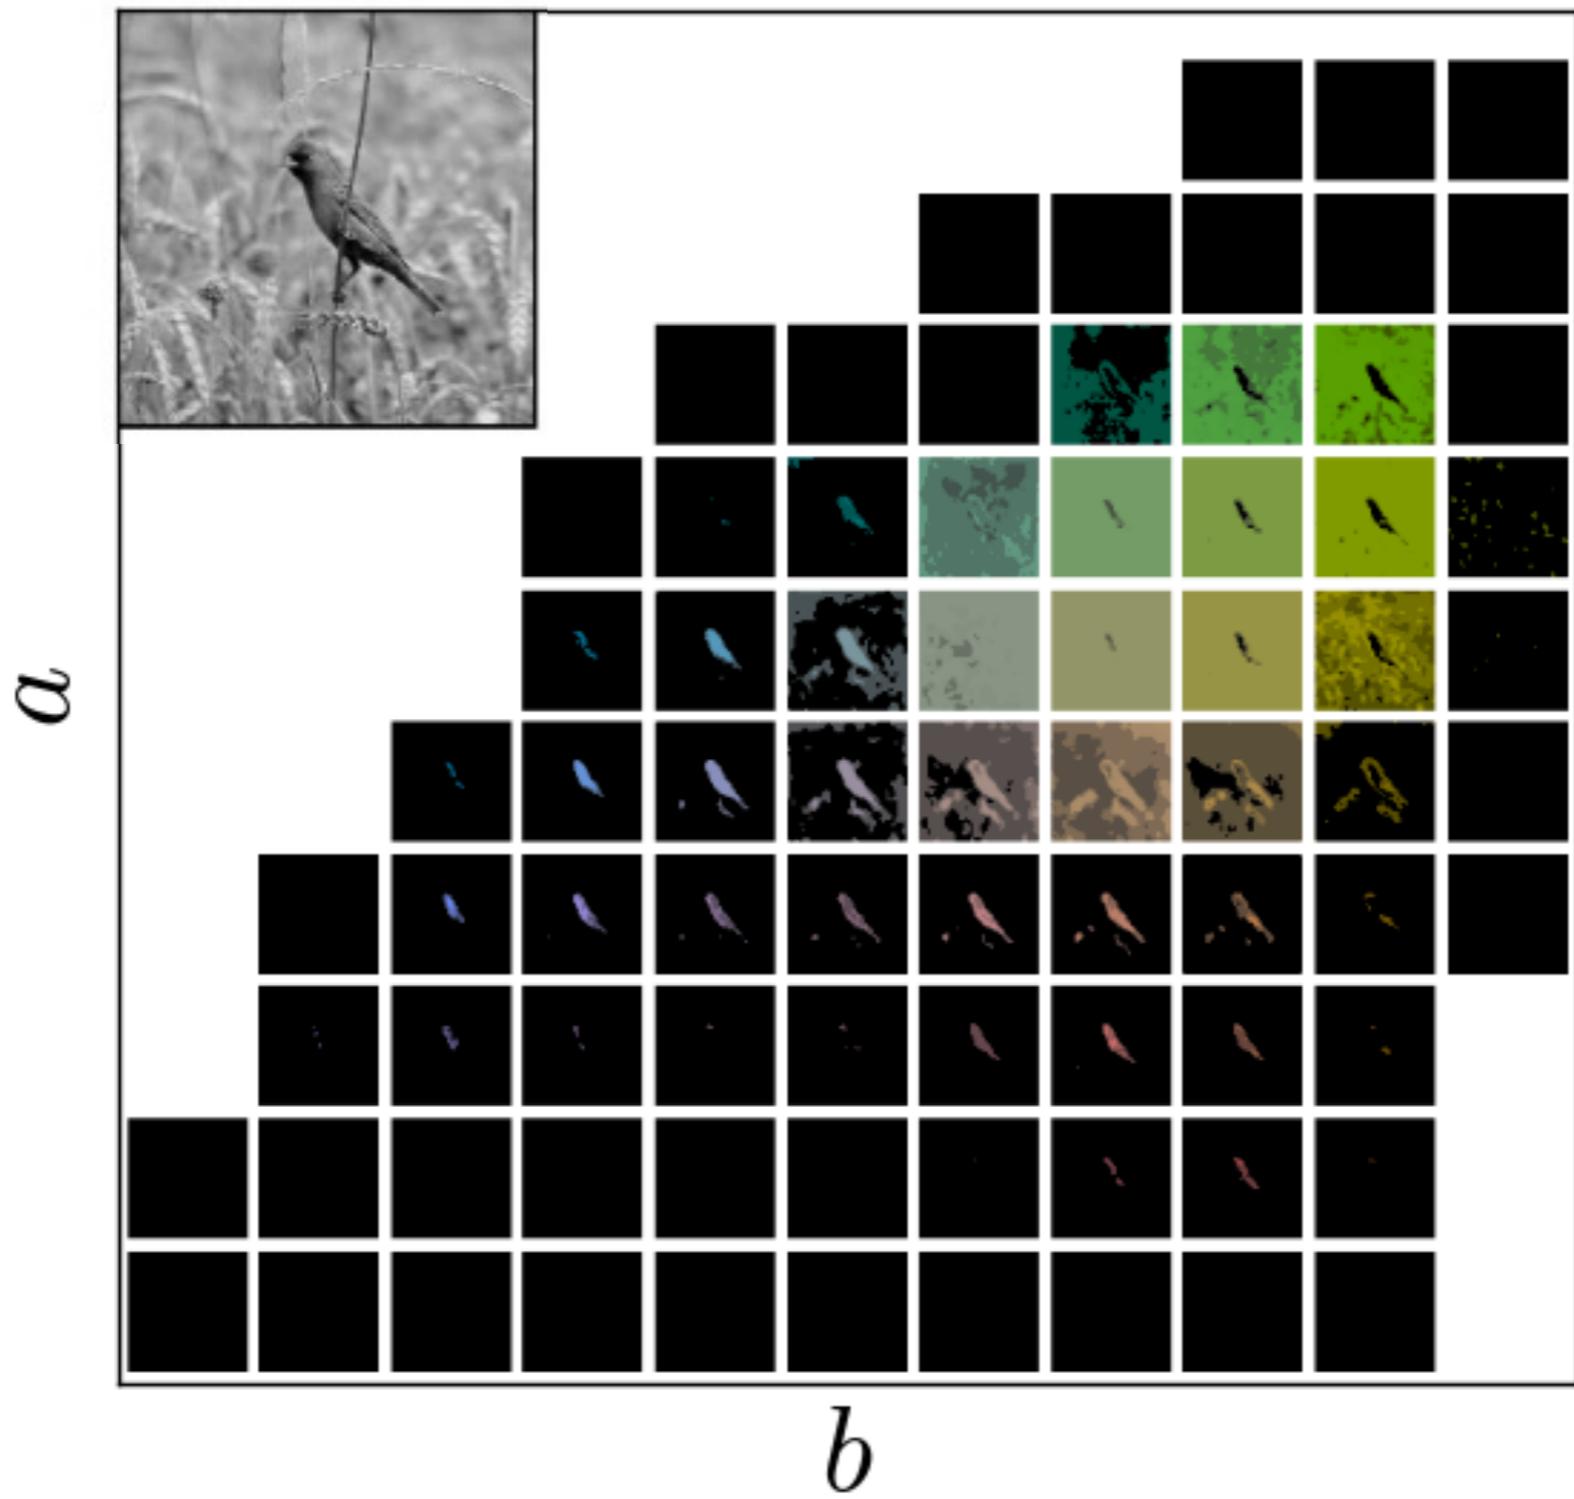
- Regression with L2 loss inadequate

$$L_2(\hat{\mathbf{Y}}, \mathbf{Y}) = \frac{1}{2} \sum_{h,w} \|\mathbf{Y}_{h,w} - \hat{\mathbf{Y}}_{h,w}\|_2^2$$

- Use per-pixel multinomial classification

$$L(\hat{\mathbf{Z}}, \mathbf{Z}) = -\frac{1}{HW} \sum_{h,w} \sum_q \mathbf{Z}_{h,w,q} \log(\hat{\mathbf{Z}}_{h,w,q})$$





Designing loss functions

Input



Zhang et al. 2016



Ground truth



Color distribution cross-entropy loss with colorfulness enhancing term.

[Zhang, Isola, Efros, ECCV 2016]

Thank You!



16-726, Spring 2021

<https://learning-image-synthesis.github.io/>